Enterprise Agents and the Frontier Labs

Investment theses and focus areas for AI agents that finish enterprise work, judged by what the frontier labs will and will not absorb

1. Investment theses: tailwinds, key takeaways and macro

This theme covers AI agents that complete a defined piece of enterprise work end to end, such as resolving a support ticket, working a denied medical claim or clearing a purchase order exception, along with the orchestration, harness and context software that runs them. Our lens throughout is what the frontier labs will and will not absorb from the companies in it. Since January, OpenAI and Anthropic have moved from selling models to selling agent platforms, department agents, solution sets for regulated industries and their own deployment services, so we judge every category on whether a company there still holds up once the labs arrive. Our research across filed expert calls, podcasts, company announcements and industry surveys answers whether the market has tipped, what the labs have absorbed, which layer holds a right to win, where we should focus today and where the main risk sits.

1. **Enterprise agents have tipped on adoption but not on depth, measured value or pricing, because almost every big enterprise now runs agents in production while few run them at scale, and buyers still pay through seats and usage**
    - Breadth has tipped, and agent counts on the biggest enterprise platforms are compounding every quarter. In a July 2026 survey 95.3% of enterprises ran at least one agent workflow in production, about 11 on average (IDC), Microsoft counts over 50M agents in its registry, up more than 30% in a quarter (Microsoft earnings call), and ServiceNow reports agentic AI consumption in production up 13x this year (ServiceNow earnings call).
    - Depth has not tipped, because most agents in production are simple, supervised and confined to a single function. About 1 in 3 enterprises has a cross-functional multi-agent solution and under 50% of proofs of concept reach production (IDC), and the chief executive of Amazon Web Services says the agents most enterprises have built are relatively simple and mostly non-autonomous, with people in the loop (AWS interview).
    - Measured value lags usage, which keeps most programs short of scale across the business. Only 13% of CIOs report significant or transformational value from AI (Gartner), 67% of enterprises overran their agent budgets by more than 10% (IDC), and the chief AI officer of a Salesforce integrator puts the share of its clients in serious production at about 20% (expert call).
    - Buyers pay through seats and usage, and outcome pricing remains a minority of deals on both sides of the market. Only 19% of services buyers and 13% of seller-side service agreements use outcome-based arrangements (Gartner), decision makers rank per-seat pricing first at 42.3% for separately billed AI, ahead of consumption at 36.6% and outcome pricing at 21.1% (Futurum), and the share of AI vendors using outcome pricing rose only from 18% to 23% in six months (ICONIQ).
    - The platforms that sell the most agents see no demand for outcome pricing and keep moving buyers toward consumption. ServiceNow's finance chief says customers are "not asking at all for outcome pricing" (ServiceNow earnings call), Microsoft is "expanding from per seat subscriptions to include consumption" (Microsoft earnings call), and Anthropic says its enterprise pricing stays usage-based as Cowork reaches general availability (Anthropic).
2. **Since January the labs have moved from selling models to selling the agent runtime, department agents, regulated solution sets and their own deployment services, so every layer close to the model now belongs to them**
    - Both labs now sell the runtime, orchestration, memory and grading that agent platforms used to sell, priced at or near the cost of tokens. Anthropic launched Managed Agents at $0.08 per session-hour in April and added multiagent orchestration, an outcomes grader and memory six weeks later (Anthropic), OpenAI

opened its Agents API in September with no fees beyond tokens and tools (OpenAI), and OpenAI launched Frontier in February as an enterprise platform to build, run and govern agents on a shared layer of business context (OpenAI).

- Both labs ship department agents that take a goal to finished work, which absorbs generic sales, finance, support and back-office tools. Anthropic open-sourced 11 Cowork plugins for roles from sales and finance to legal and customer support in January (Anthropic), OpenAI launched ChatGPT Work in July as an agent that can "turn a goal into finished work" (OpenAI), and its September developer conference previewed specialist agents tested on procurement, invoice processing and commercial contracting (OpenAI).
- Both labs now ship solution sets for the richest regulated verticals, built with licensed data partners and named design-partner customers. Anthropic released Claude for Healthcare with a prior authorization review skill in January, financial services agent templates for month-end close, reconciliation and KYC screening in the spring, and 12 legal practice-area plugins a week later (Anthropic), while OpenAI launched a life sciences model used by Amgen and Moderna and ChatGPT for Financial Services with Morgan Stanley and Evercore in September (OpenAI).
- Both labs moved into services this year, which puts them in direct competition with every company that sells deployed, finished work. OpenAI formed a majority-owned deployment company with 19 partners led by TPG and more than $4B of initial investment (OpenAI), Anthropic formed a services firm with Blackstone, Hellman & Friedman and Goldman Sachs aimed at mid-sized companies (Anthropic), and in October Anthropic committed $100M to train 10,000 forward-deployed engineers (Anthropic).
- Public markets priced the threat quickly, and the labs also walk away from layers that do not pay them. Goldman Sachs estimated about $2T was wiped off stocks in a week after the Cowork plugins and Frontier launches (Fortune), shares of Schrodinger fell as much as 8.3% on the day Anthropic launched its Claude Science research workbench (trade press), and OpenAI is winding down the agent builder and evaluations products it launched a year earlier (OpenAI).

3. **Vertical workflows hold the right to win, because the labs and record holders are absorbing orchestration, the harness and generic context, while the labs have not shown they can own proprietary data, write access, regulated liability or messy operational outcomes**
   - Orchestration holds value only as a neutral control plane across models and clouds or when the platform sells the executed work itself, and Microsoft and the labs now contest the control plane. Microsoft sells Agent 365 to govern agents at $15 per user per month (Microsoft) and OpenAI's Frontier governs agents built elsewhere (OpenAI), while 81% of big-company CIOs use three or more model families (a16z CIO survey) and 67% of agent platform users expect the control plane to sit at least partly outside a provider-managed service (VentureBeat survey), and a customer of an execution platform that bills per agent run says it saves 30 to 40 full-time roles (expert call).
   - The harness is turning into a commodity that each enterprise configures for itself, so its value goes to whoever owns the domain logic around it. Two studies cited by a harness engineer found harness choice did not significantly change coding scores (podcast), Amazon Web Services now offers a managed harness inside its AgentCore agent platform (AWS), a global integrator assigns a model to each agent itself because routing fails in production (expert call), and a fintech built its own harness in six weeks (trade press).
   - Generic context is a feature the labs bundle, and context holds value only when the company that collects it keeps it current by doing the work. Business context lifted a data agent's accuracy from about 30% to about 90% before it fell back to 60% two months later (VentureBeat), and a director at a work-context vendor expects customers to stop needing his company once the models learn to map work themselves (expert call).
   - The labs have not shown they can absorb proprietary data and verifiers, write access inside systems of record, regulated liability or outcomes in fragmented markets. An agent framework founder says that if value comes through verification, "the people who win are the domain experts" (podcast), Sapphire Ventures finds that "Labs don't want to handle the operational mess" (Sapphire Ventures), and Andrej Karpathy points founders to domains that are "valuable, verifiable, and undertrained by frontier labs"

(conference summary), and a contract manufacturer's own 18-month machine learning effort on purchasing exceptions capped out around 30% touchless (company claim).

- Vertical companies keep that right to win only when they own the accountable result and the data that improves it, because the labs now ship templates into regulated industries but do not take the liability for the work. Anthropic's services firm adds deployment capacity in healthcare, manufacturing and financial services (Anthropic), Menlo Ventures finds that regulatory certification, human sign-off and compliance infrastructure "slow down competitors" (Menlo Ventures), and Bessemer argues that every delivered unit of verifiable work becomes a training signal for the company that delivers it (Bessemer).

4. **We should focus on Series A and B vertical companies that own an accountable result for named enterprise customers in regulated verticals and in messy operational work, and back horizontal platforms only where they sell executed work priced per run**
   - Our focus follows the absorption test toward regulated or operationally messy markets with many mid-sized buyers, which the labs and horizontal vendors find costly to serve. Sapphire Ventures notes that horizontal vendors "need concentrated, high-value customers to make deployment economics work" (Sapphire Ventures), and Anthropic's services firm aims at mid-sized companies, starting with those owned by private equity (Anthropic).
   - In regulated verticals such as healthcare administration, insurance and lending, vertical companies already run production work for named enterprise customers at Series A and B prices. An insurance agent vendor runs servicing for a multi-line insurer inside the insurer's own cloud (company case study), a loan servicing vendor has processed over $1 billion in transactions for named lenders (investor announcement), and a health insurer's product leader says a claims agent let his teams bring outsourced work back in-house (expert call).
   - In non-regulated operations such as procurement and finance, vertical companies report touchless rates, volume shares and renewals at named customers. A procurement agent runs 45% to 50% of a global retailer's procurement volume after an $18M Series A (expert call, company announcement), a procurement exceptions vendor launched with $56M and says a contract manufacturer reached a 91% touchless rate in four weeks (company claim), and a drug developer that replaced a procurement consultant with an agent vendor says it would "100% renew" (expert call).
   - Horizontal platforms with the same proof are few, so we back them only where they bill per run and show named customers in production. An execution platform with a $32M Series A cut a named retailer's reconciliation from 70 days to 15 (company announcement), while an agent operating system raised a $65M seed without naming any customers (trade press).
   - The category leaders are already priced beyond our stage, so our entry point is the next Series A or B company in each focus area. The leading customer service agent company is valued at $15.8B (trade press), the freight agent leader at $1.2B (trade press), and an AI-native claims administrator raised a $125M Series C led by KKR (trade press).

5. **Outcome pricing is one pricing option among several, and the main risk is that the labs reach vertical work through data partners, services and acquisitions before the companies we back own the result and the data**
   - Outcome pricing works where the vendor can count and own the result, such as a resolved contact, a paid claim or a completed case. Salesforce charges $2 per resolution for its help agent (CX Today), Pega charges a fixed fee per completed case (CIO Dive), and a16z advises pricing at the highest layer of value a vendor can reliably measure, attribute and defend (a16z).
   - Buyers prefer spend they can predict, and outcome contracts have hurt margins where vendors cut prices ahead of automation. In a buyer survey 43% preferred consumption pricing and 27% outcome pricing (Futurum), a fintech buyer wanted "a more transactional pricing without outcome commitment" (expert call), and a former director at an AI-native outsourcer says prices fell 10% to 40% on flat costs (expert call).

- The labs reach vertical work through licensed data and services partners, which narrows the edge a vertical company holds on domain knowledge alone. Anthropic's healthcare connectors cover the federal coverage database, ICD-10 codes and the national provider registry (Anthropic), OpenAI's Frontier Alliances bring in BCG, McKinsey, Accenture and Capgemini to deploy its platform (OpenAI), and Foundation Capital reports the labs are "aggressively training agents inside simulated versions of the top 150 websites and every major system of record" (Foundation Capital).
- The labs and the best-funded AI companies also buy the positions they cannot build, so a neutral or vertical position erodes without owned results. AI companies bought 195 other AI companies through September 29, with OpenAI making 10 of those deals and Anthropic 5 (Forbes), and Anthropic's partner network program put $100M behind integrators such as Accenture, which is training 30,000 staff on Claude (Anthropic).
- Companies that hold a license or take liability for the work keep more of the price, so owning the result matters more than the pricing model. An AI-native contract firm whose lawyers take liability for the reviews its AI produces sells the finished review (company announcement), and Pegasystems' finance chief says pricing on completed work "cannot be disputed" (CIO Dive).

2. Interesting categories: where to focus

We split the theme into horizontal software that runs agents across a company, regulated verticals where a license or rulebook governs the work, and non-regulated verticals where the work is operational. We judge each area first on what the labs have shipped or will absorb there, and we focus where a company that finishes the work owns data, write access or liability that the labs do not. We do not need a position in every category, and we focus only where the labs leave a company room to own the result.

A. Horizontal

A1. Orchestration and Harness Platforms

- **Tailwinds: Enterprises run hundreds of systems and a growing number of agents, and someone has to build those agents, coordinate them and record what each one did.** A founder in the category describes enterprises with "hundreds of applications" and thousands of people acting as glue between them (founder interview).
- **Platforms that sell executed work from the operations budget reach production without a build by the customer's engineers.** An execution platform bills per agent run and its customer says it saves 30 to 40 full-time roles while opening new processes (expert call), and a teammate platform reports more than 1M hours of automated work for named media and technology companies (company claim).
- **Headwinds: At least one industrial buyer learned from a horizontal platform and now plans to build its next agents itself, which caps how far such accounts grow.** An industrial buyer used a platform to skip a 12 to 14 month build and now expects its spending to grow only 5% to 10% (expert call).
- **Custom agent systems built by outside engineers often fail to last, which weighs on platforms sold with heavy services.** Gartner expects 70% of enterprises to abandon agent systems built by vendor engineering teams by 2028 because of maintenance costs (Gartner).
- **The labs and clouds now give away the generic runtime, orchestration and control plane that horizontal platforms sell.** OpenAI's Agents API, Anthropic's Managed Agents and the harness in Amazon's AgentCore give every buyer a managed runtime, orchestration and grading at or near token prices, and Microsoft governs agents from any vendor through Agent 365 (OpenAI, Anthropic, AWS, Microsoft).
- **Why we should not focus here: Orchestration and harness platforms sit in the layer the labs absorb fastest, so we should not make them a core focus.** These platforms charge per seat or per run for a runtime the labs now price near token cost, and an agent framework founder says off-the-shelf harnesses win wherever a task sits inside what the model was trained on (podcast), so we should back a horizontal

platform only when it sells executed work priced per run, grows with each agent the customer adds and names customers in production.

Example companies: Duvo, Sycamore, Ema, Convey

A2. Work Context

- **Tailwinds: Agents need context that no system of record holds, such as how exceptions get handled and which email settles a dispute.** A work-context vendor builds a living knowledge base from emails, logs, tickets and chat histories (trade press), and Cognizant's chief executive says its partnership with a context vendor generated about $200M of incremental pipeline (trade press).
- **The process-mining incumbent argues that every agent needs a map of how processes connect.** Its co-chief executive says about 40% of work in big businesses happens on desktops outside systems (podcast).
- **Headwinds: SAP is buying a work context graph company for its HR suite, and standalone context vendors show uneven proof of paying customers.** SAP agreed to buy a work context graph company for its HR suite (SAP), and while one work-context vendor counts three named enterprise customers after a $30M Series A (trade press), another with a well-known founder names no customers on its site (company site).
- **The labs bundle business context into their platforms and department agents, which leaves little for a standalone context product.** OpenAI's Frontier ships a shared semantic layer that connects warehouses, CRM and ticketing systems for every agent (OpenAI), and Anthropic's Cowork plugins package connectors and company context for each role (Anthropic).
- **Why we should not focus here: Context holds value when a company collects it while doing the work and uses it to finish the next job, which the vertical companies in regulated and non-regulated work already do.** Standalone context vendors sell context as a separate product to buyers that the labs and record holders now court with context bundled into platforms those buyers already pay for, so we should not focus here.

Example companies: Edra, Workfabric AI

B. Regulated Verticals

B1. Healthcare Administration

- **Tailwinds: Prior authorization, coding and billing are rule-bound work that providers and health plans already pay for by the claim or by the dollar recovered.** Fees to health plans for prior authorization can include a charge per automatically approved request (trade press), and at least one coding review vendor bills hospitals only for findings that are approved and billed (vendor pricing pages).
- **Clinical review of care requests has reached production at health plans, starting with fast approvals of cancer care at a regional plan.** A prior authorization vendor that has raised $64M in total says a regional health plan approves cancer care in about 155 seconds (company claim).
- **Headwinds: The dominant electronic health record vendor controls the record at most big health systems and limits outside AI, which slows new vendors in those accounts.** A revenue cycle executive says the vendor does not "allow integrations of AI outside of what they're already working on" (expert call).
- **Providers still review claims before submission because errors carry liability, so the agent does not yet replace the whole job.** A dermatology group's chief information officer asks, "If any of the wrong claims go in, who takes a liability?" (expert call).
- **Anthropic has shipped healthcare connectors and a sample prior authorization skill, which set a free baseline for generic review tools.** Claude for Healthcare connects to the federal coverage database, ICD-10 codes and the national provider registry and includes a sample skill for reviewing prior authorization requests (Anthropic).
- **Why we should focus here: Healthcare administration is paid by the claim, the authorization or the dollar recovered, so an agent that completes the unit replaces work providers and plans already buy from outsourcers.** Payer rules, write-back into the record and liability for errors are what a lab template lacks,

and the record vendor's limits matter least at health plans and independent and specialty practices, so we should back companies that sell into those buyers and show results in the buyer's cash.

Example companies: Anterior, TJM Labs, Procode, Alaffia Health

B2. Insurance Operations

- **Tailwinds: Insurance runs on outsourced servicing, underwriting support and claims work, much of it on old systems that agents can now operate directly.** Agents can now drive old insurance screens reliably enough for production (company claim), an insurer reports its servicing agent resolves more than 90% of service cases without a person (company claim), and an annuity insurer managing over $8B of reserves chose an agent builder for new business transfers and in-force servicing (company announcement).
- **Companies that hold a license or act as the service provider can sell the whole job and keep its economics.** An AI-native commercial brokerage and AI-native claims administrators now sell placement and claims handling, from first notice to payout, as a service (trade press).
- **Claims intake is now completed by agents at volume, and a vendor built by former insurance operators reports most intake conversations resolved without a person.** A vendor that ran as a licensed agency in 47 states reports 70% to 73% autonomous resolution across 20 million claims intake conversations (company claim).
- **Neither lab has shipped an insurance solution set, and their nearest moves are generic tools that insurers reach through integrators.** Anthropic's financial services templates cover reconciliation, month-end close and KYC screening (Anthropic), and OpenAI names an insurer exploring its customer service agent for claims after severe weather (OpenAI).
- **Headwinds: Insurance buyers are pitched constantly and expect cheaper offers built on open models, which caps prices for vendors that sell tools.** An insurance executive says he is pitched "at least two or three a day" and expects a rival offer "at a 10th of the price by using open-source models" (expert call).
- **The insurer keeps the data, so switching vendors is easy and contracts do not compound.** The same executive says "All the data is living within our existing systems" (expert call).
- **At least one core insurance systems vendor now invests in agent companies built on its platform, which gives the record holder a say in which vendors win.** Guidewire joined an underwriting agent company's $30M Series B as a strategic investor (company announcement).
- **Why we should focus here: Insurance operations are bought by the claim, the submission and the policy, and most of the work sits with outsourcers on old systems the labs do not touch.** Companies that hold a license or carry the operating liability keep more of the price than tool vendors and face the least pressure from cheaper open models, so we should back the ones with a named carrier in production across more than one line of work.

Example companies: Notch, Pace, Harper, Strala, Brain Co

B3. Lending and Financial Crime

- **Tailwinds: Loan servicing and collections run on scripted, regulated conversations, and at least one servicing vendor's agents now complete them across voice, text and email.** A servicing agent vendor says it cut lenders' handle times by 60% and monitors both AI and human agents for unfair or deceptive practices (investor announcement).
- **Compliance alerts arrive as countable cases, and backlogs draw fines when hiring lags volume.** A compliance agent vendor says alerts pile into backlogs that regulators fine (company site), and at least one outsourcer now packages anti-money laundering checks as agent solutions under committed pricing (trade press).
- **Headwinds: The biggest card issuers and banks can build their own agent platforms under model risk rules, which narrows the opening at those institutions.** Capital One built its own agent platform with evaluations and controls for regulated use (podcast).

- **At least one fintech routes work across three vendors through a router it built itself, so no single vendor holds that account securely.** A fintech routed tickets across three customer service vendors through a router it built itself (expert call).
- **Both labs now sell financial services packages that set a free baseline for generic compliance and analyst work.** Anthropic's templates cover close, reconciliation and KYC screening for banks and insurers (Anthropic), and OpenAI built ChatGPT for Financial Services with Morgan Stanley and Evercore as design partners (OpenAI).
- **Why we should focus here: Lending servicing and financial compliance are countable, regulated and costly when they go wrong, and the company that monitors its own conduct takes on a share of the compliance risk the lender carries.** Anthropic's KYC screening template screens customers, while a company that runs the whole AML and sanctions investigation for the bank owns the case and answers for its quality. The biggest banks build their own and the lab packages aim at analysts, so we should back companies serving consumer lenders, fintechs and regional banks with named customers and volumes.

Example companies: Salient, Bretton AI, Gradient Labs

B4. Legal Services

- **Tailwinds: Legal agents are paid from the budget that pays outside counsel, often per matter.** Personal injury demand packages, the file of evidence and damages a lawyer sends an insurer to settle a claim, are now priced per package (investor research), and an AI-native consumer law firm reports revenue above $35M (trade press).
- **Firms that hold a license and take liability for the result can price the finished work, and a German appeals court has already held a company liable for what its chatbot said.** An AI-native contract firm says it has no billable hours and reinvests profits in the workflow to raise margin (podcast), and a German appeals court held a cosmetic clinic responsible for specialist qualifications its website chatbot invented (law firm commentary).
- **Headwinds: Fixed and outcome pricing for legal work takes time to get right, and at least one AI-native firm lost money while learning it.** The co-founder of an AI-native contract firm says it lost a lot of money on pricing before getting it right (podcast).
- **Anthropic now ships legal practice-area plugins and contract connectors, which makes generic drafting and review tools cheap.** This spring it released 12 legal plugins and more than 20 connectors to contract, deal and document systems (Anthropic).
- **Why we should focus here: As drafting tools get cheap, we should back firms that sell finished legal work under their own license, because they keep the price of the outside counsel they replace and the lab plugins do not carry liability.** The plaintiff legal leader already raised at a valuation above $2B (company announcement), so we should back AI-native firms in commercial contracts and consumer law at Series A and B.

Example companies: Crosby, Lawhive, EvenUp

B5. Life Sciences Research

- **Tailwinds: Drug companies are starting to buy research agents, beginning with target discovery, which gives vendors a paying customer for a clear unit of work.** A research agent vendor says one run equals about 6 months of a scientist's work (company claim), and Pfizer agreed to deploy a research agent across its research and development with forward-deployed scientists (company announcement).
- **At least one research agent vendor trains on each drug company's own data, which the labs avoid touching.** A research agent founder says his company trains on each customer's data (founder interview).
- **Headwinds: Research outcomes take years to prove, so the agent's value is hard to price per result.** A research agent priced at $200 per run ended its founding subscriptions seven months after launch (company announcements).

- **Incyte keeps research decisions with its own teams while deploying a research agent, which slows the move from pilot to paid runs.** Incyte's deployment covers research teams that still own every program decision themselves (trade press).
- **OpenAI launched a life sciences model with Amgen, Moderna and Thermo Fisher, and Anthropic launched a research workbench for scientists in June.** OpenAI's Life Sciences plugin connects the model to more than 50 scientific tools (OpenAI), and Anthropic's Claude Science combines scientific databases, code and compute in one desktop workbench (Anthropic).
- **Why we should focus here: Life sciences research pays for discovery work in a regulated industry, and a research agent that trains on a drug company's own data does work a lab workbench does not.** Results take years to prove, so we should back companies that already show repeat paid runs at the same customer and price each run against the program decision it informs.

Example companies: Edison Scientific

C. Non-Regulated Verticals

C1. Procurement, Logistics and Distribution

- **Tailwinds: Procurement savings are measured against total spend, which ties an agent's value to the money it saves on purchases as well as the labor it replaces.** For a manufacturer spending $10B a year, a 50% more efficient team saves $50M, while a 1% to 2% cut in purchase prices is worth $100M to $200M (founder interview).
- **At a contract manufacturer, exception work in purchasing and payables defeated a long internal automation effort, which opens the door to vendors that finish it.** A contract manufacturer's internal machine learning effort ran for 18 months and capped out around 30% touchless (company claim).
- **Freight operators and distributors are buying agents for quoting, dispatch, collections and service, and the leaders show scale.** A freight agent vendor reports more than 150 enterprise customers and 5x growth since its Series B (trade press), and a distribution agent vendor raised a $60M Series B with a home improvement retailer among its investors (trade press).
- **Headwinds: Procurement agents that run beside the buyer's record are easy to replace, and at least one global retailer puts a replacement at seven to eight months.** A global retailer that pays on a consumption license puts a replacement of its procurement agent vendor at seven to eight months (expert call).
- **Buyers keep a person on the final step, so the full price of the job moves to the agent slowly.** The chief executive of a procurement outsourcer says "I don't have any client that lets it go on its own" (expert call).
- **OpenAI is testing its own specialist agents for procurement, invoice processing and commercial contracting, which sets a lab baseline for the simplest purchasing work.** Its September developer conference previewed specialist agents with their own identity for those roles and a marketplace where buyers spend their OpenAI commitment on 32 partners (OpenAI).
- **Why we should focus here: Procurement, logistics and distribution work runs across suppliers, carriers, phone calls and email outside any single record, which makes it hard for an ERP vendor or a lab agent to absorb.** The agents sit on the buyer's ERP, so results are counted in a record the buyer trusts, and value grows with the spend the agent handles. We should back companies that clear exceptions end to end and show touchless rates at named manufacturers, freight operators and distributors, at Series A or B prices since the freight leader already trades as a Series C company.

Example companies: Fragment, Magentic, Didero, InstaLILY, BackOps, HappyRobot

C2. Finance and Accounting Operations

- **Tailwinds: Finance teams repeat the same work every month, and each invoice, collection or reconciliation has a clear count and cost.** Vendors in payables and collections report volumes above 100M invoices a year and overdue balances cut by 40% (company claims).

- **Accounting and payables agents now sell the finished job and take responsibility for it.** An accounting agent vendor says it never had a seat or token price for its core product, and a payables vendor says its agents decide whether to pay each invoice and "take accountability for the outcome" (company interviews).
- **At least one accounting agent vendor built for a single trade has won customers at scale, which shows room for agents built around one industry's books.** That vendor raised a $30M Series B and counts more than 750 construction companies as customers (company announcement).
- **Headwinds: Oracle is building agent orchestration into its Fusion ERP, which limits how far a new vendor can reach inside Fusion finance teams.** Oracle sells its agents through paid AI units, with a rules layer validating every action the model proposes (trade press).
- **Agents cost more to run than the scripted bots many finance teams already use, so vendors must show savings beyond the old automation.** An integrator's automation lead says a task that costs $200 or $300 a month on bots costs $1,000 or $1,200 a month with AI (expert call).
- **Anthropic's financial services templates now cover the generic monthly close directly, which removes the opening for generic close tools.** Its spring release includes a month-end closer, a general ledger reconciler and a statement auditor alongside KYC screening (Anthropic).
- **Why we should focus here: Finance and accounting work recurs with the monthly close and is checked in the buyer's own books, and the companies that take accountability for each invoice or filing charge for the finished job.** ERP vendors and lab templates cover the generic close, so we should back companies that serve accounting firms, global payroll or a specific industry's books, where rules and liability sit outside the ERP. The best-known accounting agent company already raised a $100M Series B at a $1.15B valuation (company interview), so our entry point is the next company serving one industry's books or global payroll.

Example companies: Basis, Freehand, Niural

C3. Customer Service and Sales

- **Tailwinds: Customer service is the first job agents complete end to end at scale, because a resolved ticket is easy to define and check and the labor it replaces is easy to price.** Software companies spend $20 to $30 per ticket resolved by people, which gives buyers a clear benchmark for an agent that resolves the ticket itself (founder interview).
- **Voice is opening a second wave of spending in contact centers that still run mostly on people.** The US still employs 2,666,000 customer service representatives whose work voice agents can now take on (government data).
- **Headwinds: Sales agents churn heavily, because agents that send cold email burn through prospects and fail to show value inside short pilots.** An advisor to sales agent companies puts churn at 80% to 90%, the worst of any category he advises (expert call).
- **Salesforce bundles its own sales agents to replace point tools, which squeezes standalone vendors in its accounts.** Salesforce says its Agentforce 1 Edition at $550 replaces roughly $1,300 of point-solution spend per US sales rep (Salesforce conference remarks).
- **OpenAI now sells its own customer service agent deployed by its own engineers, and Salesforce already sells a help agent priced per resolution.** Salesforce charges $2 per resolution for its help agent (CX Today), and OpenAI launched Presence in July with written procedures, guardrails and automated scoring, and says it resolves 75% of inbound issues on its own phone line without human assistance (OpenAI).
- **Why we should not focus here: Customer service proves that agents can finish a whole job at scale, but OpenAI and Salesforce now sell that job themselves.** Vendors here charge per conversation or per resolution, Salesforce already prices each resolution at $2 (CX Today), and the leaders already raised at prices that leave a new entrant little room. Sales agents are paid on meetings and leads that are easy to game, and Salesforce bundles its own. We should use customer service resolution rates and prices as the benchmark for other functions and make no new bets in the category.

Example companies: Sierra, Decagon