Enterprise Agents and the Frontier Labs

Investment theses and focus areas for AI agents that finish enterprise work, judged by the budget they draw on and by what the frontier labs will and will not absorb

1. Investment theses: key takeaways and macro

This theme covers AI agents that complete a defined piece of enterprise work end to end, such as resolving a support ticket, working a denied medical claim or clearing a purchase order exception, along with the orchestration, harness and context software that runs them. We judge every company in the theme with two questions about money and about the labs. The first asks which budget pays for the work, because an agent that finishes a job is paid from the money a company spends on people and outsourcers, which is several times what it spends on software. The second asks what the frontier labs will absorb, because since January OpenAI and Anthropic have moved from selling models to selling agent platforms, department agents, regulated solution sets and their own deployment services.

Key takeaways

- **Adoption of enterprise agents has tipped, while the depth of use and the measured value have not.** Buyers still judge each agent by the work it completes and want the value to show up in their own records before they scale it.
- **The labor budget is the prize, because agents that finish work replace spending on people and outsourcers.** Labor, outsourcing and services budgets are several times the software budget, and spending on AI is moving from the IT department to the business owners who hold those budgets.
- **The labs now own every layer close to the model and sell it into the IT budget, and their new services arms are their first move toward the labor budget.** Runtime, orchestration, generic context, department agents and starter kits for regulated industries are sold by the labs at or near the cost of tokens and seats.
- **Vertical companies that own the result of the work hold the right to win against the labs.** The labs have not shown they can own proprietary data, write access to systems of record, regulated liability or messy operations, and each of these keeps a company inside the labor budget when the labs arrive.
- **We focus on vertical work and back horizontal platforms selectively where they own the work.** We back companies with enterprise customers in production, paid from the labor or outsourcing budget, and horizontal platforms that own executed work and stay neutral across models, while the main risk is the labs reaching vertical work through partners and acquisitions.

1. **Enterprise agents have tipped on adoption but not on depth or measured value, because almost every big enterprise now runs agents in production while few run them across functions or can show a return**
   - Breadth has tipped, and agent counts on the biggest enterprise platforms are compounding every quarter. In a July 2026 survey 95.3% of enterprises ran at least one agent workflow in production, about 11 on average (IDC), Microsoft counts over 50M agents in its registry, up more than 30% in a quarter (Microsoft earnings call), and ServiceNow reports agentic AI consumption in production up 13x this year (ServiceNow earnings call).
   - Depth has not tipped, because most agents in production are simple, supervised and confined to a single function. About 1 in 3 enterprises has a cross-functional multi-agent solution and under 50% of proofs of concept reach production (IDC), and the chief executive of Amazon Web Services says the agents most enterprises have built are relatively simple and mostly non-autonomous, with people in the loop (AWS interview).
   - Measured value lags usage, which keeps most programs short of scale across the business. Only 13% of CIOs report significant or transformational value from AI (Gartner), only 22% of organizations have scaled

AI across multiple business units (Gartner), and 67% of enterprises overran their agent budgets by more than 10% (IDC).

- Buyers still pay mostly through seats and usage, and ServiceNow says its customers are not asking for outcome pricing. Only 19% of services buyers use outcome-based arrangements (Gartner), decision makers rank per-seat pricing first at 42.3% for separately billed AI, ahead of consumption at 36.6% and outcome pricing at 21.1% (Futurum), and ServiceNow's finance chief says customers are "not asking at all for outcome pricing" (ServiceNow earnings call).

2. **Agents that finish work are paid from the labor budget, which is several times the software budget and is moving to business owners outside IT, while the labs still sell mostly tokens and seats into IT and are now building services arms to reach it**
   - The money for finished work sits in labor and services budgets, which investors estimate at several times software spend and more. Sequoia estimates that "For every dollar spent on software, six are spent on services" (Sequoia), Bessemer puts professional services at about 13% of US GDP, 10x the size of the software industry (Bessemer), Felicis finds labor budgets run 35 times software budgets (Felicis), and IT services alone already outspend software at $1,557B against $1,464B in Gartner's 2026 forecast (Gartner).
   - The US wage base for the office work agents now complete runs to millions of people in each job family. Government data counts 17,753,430 office and administrative support workers at a mean wage of $51,560, 10,537,240 business and financial operations workers at $95,230 and 2,595,750 customer service representatives at $46,590 (BLS), and Foundation Capital sizes the global salary and services spend open to agents at "a $4.6 trillion opportunity" (Foundation Capital).
   - The outsourcing pool is already repricing, which shows the labor budget moving from people to software in real contracts. HFS projects the outsourcing market at $725B in 2035 against $786B on its conventional path, with agent-delivered services inside it becoming "a segment worth $422 billion" (HFS Research), a contact center outsourcing executive says the standard rate his firm can charge fell 20% to 40% in two years and that billing is moving from per FTE to consumption (expert call), and Concentrix took a $1,050 million goodwill impairment in its third quarter after its share price fell (Concentrix).
   - Spending on AI is moving out of the IT budget and into the business units that own labor costs. BCG finds that "more than 80% of AI spending now sits outside enterprise IT" (BCG) and that companies spend 2.8% of revenue on AI funded outside the IT budget on top of 3.3% on enterprise IT (BCG), Gartner finds functional leaders put 12% of their budgets into AI in 2025 with 85% planning increases (Gartner), and among over-budget organizations that fund AI by reallocating within IT, 60.9% cut external contractors and consultants, while 23.1% draw on a non-IT business unit budget (Futurum).
   - A labor budget owner compares the agent with the outsourcer's invoice or the team's payroll, which favors the company that owns the finished work over a lab that sells tokens. Sequoia notes that replacing an outsourcing contract with an AI provider "is a vendor swap" while replacing headcount "is a reorg" (Sequoia), and a Cognizant automation head says 60% to 70% of outsourced business process work can be automated (expert call).

3. **Since January the labs have moved from selling models to selling the agent runtime, department agents, regulated solution sets and their own deployment services, so every layer close to the model now belongs to them**
   - Both labs now sell the runtime, orchestration, memory and grading that agent platforms used to sell, priced at or near the cost of tokens. Anthropic launched Managed Agents at $0.08 per session-hour in April and added multiagent orchestration, an outcomes grader and memory six weeks later (Anthropic), OpenAI opened its Agents API in September with no fees beyond tokens and tools (OpenAI), and OpenAI launched Frontier in February as an enterprise platform to build, run and govern agents on a shared layer of business context (OpenAI).
   - Both labs ship department agents that take a goal to finished work, which absorbs generic sales, finance, support and back-office tools. Anthropic open-sourced 11 Cowork plugins for roles from sales and finance to legal and customer support in January (Anthropic), OpenAI launched ChatGPT Work in July as an agent

that can "turn a goal into finished work" (OpenAI), and its September developer conference previewed specialist agents tested on procurement, invoice processing and commercial contracting (OpenAI).

- Both labs now ship solution sets for the richest regulated verticals, built with licensed data partners and named design-partner customers. Anthropic released Claude for Healthcare with a prior authorization review skill in January, financial services agent templates for month-end close, reconciliation and KYC screening in the spring, and 12 legal practice-area plugins a week later (Anthropic), while OpenAI launched a life sciences model used by Amgen and Moderna and ChatGPT for Financial Services with Morgan Stanley and Evercore in September (OpenAI).
- Both labs moved into services this year, which puts them in direct competition with every company that sells deployed, finished work. OpenAI formed a majority-owned deployment company with 19 partners led by TPG and more than $4B of initial investment (OpenAI), Anthropic formed a services firm with Blackstone, Hellman & Friedman and Goldman Sachs aimed at mid-sized companies (Anthropic), and in October Anthropic committed $100M to train 10,000 forward-deployed engineers (Anthropic).
- Public markets priced the threat quickly, and OpenAI has already walked away from one layer that did not pay. Goldman Sachs estimated about $2T was wiped off stocks in a week after the Cowork plugins and Frontier launches (Fortune), and OpenAI is winding down the agent builder and evaluations products it launched a year earlier (OpenAI).

4. **Vertical workflows hold the right to win, because the labs and record holders are absorbing the runtime, the harness and generic context, while the labs have not shown they can own proprietary data, write access, regulated liability or messy operational outcomes**
   - Orchestration holds value as a neutral control plane across models and clouds or when the platform sells the executed work itself, and Microsoft and the labs now contest the generic control plane. Microsoft sells Agent 365 to govern agents at $15 per user per month (Microsoft) and OpenAI's Frontier governs agents built elsewhere (OpenAI), while 81% of big-company CIOs use three or more model families (a16z CIO survey) and 67% of agent platform users expect the control plane to sit at least partly outside a provider-managed service (VentureBeat survey).
   - The harness is turning into a commodity that each enterprise configures for itself, so its value goes to whoever owns the domain logic around it. Two studies cited on a practitioner podcast found harness choice did not significantly change coding scores (podcast), Amazon Web Services now offers a managed harness inside its AgentCore agent platform (AWS), and a global integrator assigns a model to each agent itself because routing fails in production (expert call).
   - Generic business context is a feature the labs and record holders bundle, and context keeps its value only when it is collected and kept current by doing the work. OpenAI's Frontier ships a shared semantic layer across warehouses, CRM and ticketing systems for every agent (OpenAI), and SAP agreed to buy a work context graph company for its HR suite (SAP).
   - The labs have not shown they can absorb proprietary data and verifiers, write access inside systems of record, regulated liability or outcomes in fragmented markets. Sapphire Ventures finds that "Labs don't want to handle the operational mess" (Sapphire Ventures), Andrej Karpathy points builders to domains that are "valuable, verifiable, and undertrained by frontier labs" (conference summary), and Menlo Ventures finds that regulatory certification, human sign-off and compliance infrastructure "slow down competitors" (Menlo Ventures).
   - Vertical companies keep that right to win when they own the accountable result and the data that improves it, because the labs ship templates into regulated industries without taking the liability for the work. Bessemer argues that every delivered unit of verifiable work becomes a training signal for the company that delivers it (Bessemer), and nearly three-quarters of the companies McKinsey counts as AI high performers fundamentally redesigned their workflows, against one-quarter of the rest (McKinsey).
5. **We should focus on vertical companies that own an accountable result paid from the labor budget and back horizontal platforms selectively, while the main risk is the labs reaching vertical work through data partners, services and acquisitions**

- Our focus follows the budget and the absorption test toward regulated or operationally messy markets with many mid-sized buyers, which the labs and horizontal vendors find costly to serve. Sapphire Ventures notes that horizontal vendors "need concentrated, high-value customers to make deployment economics work" (Sapphire Ventures).
- We back companies that show enterprise traction in the buyer's own records, meaning named enterprise customers in production, results the buyer counts and spend that comes from the labor or outsourcing budget. Incumbents treat bold savings claims with suspicion, as a revenue cycle executive whose own gains run to a "low-double digit" says of anything higher, "you can pretty much scrub it" (expert call), and Menlo Ventures notes that "The companies competing for labor budgets get you into the work" (Menlo Ventures).
- We back horizontal platforms selectively, where they own executed work paid from the labor budget and stay neutral across models and clouds, because the labs and Microsoft now sell the generic runtime and control plane themselves. More than 80% of AI spending now sits outside enterprise IT (BCG) and 81% of big-company CIOs use three or more model families (a16z CIO survey), while Gartner expects 70% of enterprises to abandon agent systems built by vendor engineering teams by 2028 because of maintenance costs (Gartner), which rules out platforms that depend on custom builds.
- Outcome pricing is one way to sell finished work among several, and the labor budget pays for a finished job whether it is priced per outcome, per case or per run. Salesforce charges $2 per resolution for its help agent (CX Today), Pegasystems' finance chief says pricing on completed work "cannot be disputed" (CIO Dive), buyers in a survey preferred consumption pricing at 43% over outcome pricing at 27% (Futurum), and a16z advises pricing at the highest layer of value a vendor can reliably measure, attribute and defend (a16z).
- The main risk is the labs reaching vertical work through licensed data, services partners and acquisitions before the companies we back own the result and the data. Anthropic's healthcare connectors cover the federal coverage database, ICD-10 codes and the national provider registry (Anthropic), OpenAI's Frontier Alliances bring in BCG, McKinsey, Accenture and Capgemini to deploy its platform (OpenAI), Foundation Capital reports the labs are "aggressively training agents inside simulated versions of the top 150 websites and every major system of record" (Foundation Capital), and AI companies bought 195 other AI companies through September 29, with OpenAI making 10 of those deals and Anthropic 5 (Forbes).

2. Interesting categories: where to focus

We split the theme into horizontal software that runs agents across a company, regulated verticals where a license or rulebook governs the work, and non-regulated verticals where the work is operational. We judge each area on three things, which are the size of the labor and outsourcing budget it draws on, what the labs have shipped or will absorb there, and whether a company that finishes the work owns data, write access or liability that the labs do not.

A. Horizontal

A1. Orchestration and Execution Platforms

- **Tailwinds: Enterprises now run many agents on many models and want a control layer they own across all of them.** In a VentureBeat survey, 67% of agent platform users expect the control plane to sit at least partly outside a provider-managed service (VentureBeat survey).
- **Agents are gaining real authority over business decisions faster than companies are building the controls to govern them.** In a BCG survey of 1,330 executives, 42% expect agents to hold real decision authority by 2030 while only 5% have the controls in place (BCG).
- **Functional leaders now hold their own AI budgets and aim them first at productivity, which is the budget a platform that executes cross-system work sells into.** Functional leaders put 12% of their budgets into AI in 2025, productivity is the target for 75% of them and it takes about 30% of their AI spend (Gartner).

- **Headwinds: Custom agent systems built by outside engineers often fail to last, which weighs on platforms sold with heavy services.** Gartner expects 70% of enterprises to abandon agent systems built by vendor engineering teams by 2028 because of maintenance costs (Gartner).
- **Many agent projects stall before they pay back, which lengthens sales cycles for any platform sold ahead of a proven use case.** Gartner expects over 40% of agentic AI projects to be canceled by the end of 2027 (Gartner).
- **The labs and clouds now give away the generic runtime, orchestration and control plane that horizontal platforms sell.** OpenAI's Agents API, Anthropic's Managed Agents and the harness in Amazon's AgentCore give every buyer a managed runtime, orchestration and grading at or near token prices, and Microsoft governs agents from any vendor through Agent 365 (OpenAI, Anthropic, AWS, Microsoft).
- **Why we should focus here selectively: Horizontal platforms deserve a place in our focus where they sell finished work into the labor budget instead of a runtime into the IT budget.** We should back a horizontal platform that owns the executed work across many processes, is paid from the operations budget and stays neutral across models and clouds, with process context built up by running the work, and we should pass on platforms that resell the generic runtime or control plane that the labs and Microsoft now price near token or seat cost.

Example companies: Duvo, Sycamore, Ema, Convey

A2. Work Context

- **Tailwinds: Enterprises expect to tailor agents to the way their own business works, which requires context that no off-the-shelf model carries.** In a Deloitte survey of 3,235 leaders, 85% expect to customize agents to fit the unique needs of their business, while only 21% have mature agent governance (Deloitte).
- **SAP agreed to buy a work context graph company for its HR suite, which shows a record holder paying for context about how work gets done.** SAP says the deal brings the target's context graph, AI models and applied research team into SAP and establishes work intelligence as a strategic SAP asset (SAP).
- **Headwinds: Record holders can buy or build the context layer themselves, which leaves a standalone context vendor competing with the suite its buyer already runs.** The same SAP deal puts a context graph for work inside the HR suite that SAP customers already license (SAP).
- **The labs bundle business context into their platforms and department agents, which leaves little for a standalone context product.** OpenAI's Frontier ships a shared semantic layer that connects warehouses, CRM and ticketing systems for every agent (OpenAI), and Anthropic's Cowork plugins package connectors and company context for each role (Anthropic).
- **Why we should focus here selectively: Context holds value when a company collects it while doing the work and uses it to finish the next job, so we should back context companies selectively.** We should back a context layer that is captured and kept current by executing work in specific processes, that sits close to the labor budget of the teams it serves and that the buyer sees as neutral across models, and we should pass on context sold as a separate catalog to buyers whose platforms already bundle it.

Example companies: Edra, Workfabric AI

B. Regulated Verticals

B1. Healthcare Administration

- **Tailwinds: Federal rules now put deadlines and public reporting on prior authorization, which turns slow manual review into a measurable cost for health plans.** A CMS rule requires most affected payers to decide within 72 hours for urgent requests and seven calendar days for standard ones from January 1, 2026, give a specific reason for each denial and publish their metrics (CMS).
- **The volume of authorization work is high on both sides, and practices staff for it as a full-time job.** Practices complete 39 prior authorization requests per physician per week, 40% of physicians have staff

who work only on prior authorizations, and Medicare Advantage insurers received more than 50 million requests in 2023 (AMA).

- **Revenue cycle work is already bought from outsourcers on contracts that tie price to labor or cash collected.** A revenue cycle services executive says "if it's FTE-based, then we take the savings" and that contracts based on collections earn more when clients collect more (expert call).
- **Headwinds: The dominant electronic health record vendor controls the record at most big health systems and limits outside AI, which slows any new agent in those accounts.** A revenue cycle executive says the vendor does not "allow integrations of AI outside of what they're already working on" (expert call).
- **Prior authorization decisions are often reversed on appeal, so providers and plans keep a person on contested cases until a vendor takes liability for the result.** Among appealed Medicare Advantage denials, 81.7% saw the initial denial fully or partially overturned (AMA).
- **Anthropic has shipped healthcare connectors and a sample prior authorization skill, which set a free baseline for generic review tools.** Claude for Healthcare connects to the federal coverage database, ICD-10 codes and the national provider registry and includes a sample skill for reviewing prior authorization requests (Anthropic).
- **Why we should focus here: Healthcare administration is paid by the claim, the authorization or the dollar recovered, so an agent that completes the unit replaces labor that providers and plans already buy from outsourcers.** Payer rules, write-back into the record and liability for errors are what a lab template lacks, and the record vendor's limits matter least at health plans and independent and specialty practices, so we should back companies that sell into those buyers and show results in the buyer's cash.

Example companies: Anterior, TJM Labs, Procode, Alaffia Health

B2. Insurance Operations

- **Tailwinds: Insurance runs on big claims and policy teams whose work is defined by forms, rules and deadlines that agents can follow.** Government data counts 324,230 claims adjusters, examiners and investigators at a mean wage of $80,470 and 214,260 insurance claims and policy processing clerks (BLS).
- **Traditional outsourcing, where much insurance servicing sits today, is losing share of delivery to models built on automation.** HFS expects traditional outsourcing models to "decline from 55% to 37% of delivery" within two years (HFS Research).
- **Neither lab has shipped an insurance solution set, and their nearest moves are generic tools that insurers reach through integrators.** Anthropic's financial services templates cover reconciliation, month-end close and KYC screening (Anthropic), and OpenAI names an insurer exploring its customer service agent for claims after severe weather (OpenAI).
- **Headwinds: Insurance buyers are pitched constantly and expect cheaper offers built on open models, which caps prices for vendors that sell tools.** An insurance executive says he is pitched "at least two or three a day" and expects a rival offer "at a 10th of the price by using open-source models" (expert call).
- **The insurer keeps the data, so switching vendors is easy and contracts do not compound on their own.** The same executive says "All the data is living within our existing systems" (expert call).
- **Why we should focus here: Insurance operations are bought by the claim, the submission and the policy, and most of the work sits with in-house teams and outsourcers on old systems the labs do not touch.** Companies that hold a license or carry the operating liability keep more of the price than tool vendors and face the least pressure from cheaper open models, so we should back the ones with a named carrier in production across more than one line of work.

Example companies: Notch, Pace, Harper, Strala, Brain Co

B3. Lending and Financial Crime

- **Tailwinds: Financial crime compliance is a big, rising cost whose biggest line is people.** Financial crime compliance costs institutions in the US and Canada $61 billion a year, costs rose for 99% of institutions and "Labor costs rank highest in financial crime compliance spending" (LexisNexis and Forrester).

- **Alert volumes keep climbing and institutions want to cut the cost of clearing them, which favors work sold per case.** Screening alerts rose at 83% of mid-sized and big institutions, and 70% of institutions prioritize cutting compliance costs (LexisNexis and Forrester).
- **Headwinds: Capital One built its own agent platform under model risk rules, which shows the biggest card issuers and banks can build their own and narrows the opening at those institutions.** Its platform includes the evaluations and controls the bank built for agents in regulated use (podcast).
- **Both labs now sell financial services packages that set a free baseline for generic compliance and analyst work.** Anthropic's templates cover close, reconciliation and KYC screening for banks and insurers (Anthropic), and OpenAI built ChatGPT for Financial Services with Morgan Stanley and Evercore as design partners (OpenAI).
- **Why we should focus here: Lending servicing and financial compliance are countable, regulated and costly when they go wrong, and the labor in them is the biggest line of spend.** A lab screening template screens customers, while a company that runs the whole AML and sanctions investigation for the bank owns the case and answers for its quality. The biggest banks build their own and the lab packages aim at analysts, so we should back companies serving consumer lenders, fintechs and regional banks with named customers and volumes.

Example companies: Salient, Bretton AI, Gradient Labs

B4. Legal Services

- **Tailwinds: Legal work draws on a big services budget that buyers already pay per matter or per hour to outside firms.** Speakers at a Sequoia event put legal services alone at about $400B in the US (Sequoia event transcript), and government data counts 392,880 paralegals and legal assistants at a mean wage of $69,700 (BLS).
- **Legal buyers already pay more for agentic tools, which shows a budget for finished legal work beyond research seats.** Thomson Reuters says 32% of its annualized contract value comes from GenAI-enabled products and that "the agentic offering is a premium tier pricing mechanism for us" (Thomson Reuters earnings call).
- **Headwinds: A German appeals court has already held a company liable for what its chatbot said, which raises the bar for any vendor whose agent speaks to clients.** A German appeals court held a cosmetic clinic responsible for specialist qualifications its website chatbot invented (law firm commentary).
- **Thomson Reuters bundles agentic features into the research tools firms already license, which squeezes standalone drafting tools.** Thomson Reuters prices its agentic offering as a premium tier on top of its existing products (Thomson Reuters earnings call).
- **Anthropic now ships legal practice-area plugins and contract connectors, which makes generic drafting and review tools cheap.** This spring it released 12 legal plugins and more than 20 connectors to contract, deal and document systems (Anthropic).
- **Why we should focus here: As drafting tools get cheap, we should back firms that sell finished legal work under their own license, because they keep the price of the outside counsel they replace and the lab plugins do not carry liability.** Commercial contracts and consumer law are the areas where buyers pay per matter today and where a licensed firm takes the liability a lab plugin avoids.

Example companies: Crosby, Lawhive, EvenUp

B5. Life Sciences Research

- **Tailwinds: Drug research is one of the biggest professional budgets in any industry, and it keeps growing.** R&D spending at the biggest pharma companies rose to $190Bn in 2024 from $163Bn in 2023 and reached 25.2% of sales (IQVIA Institute).
- **Clinical trials are taking longer to fill, and the labs already name drug companies such as Novo Nordisk and Sanofi as customers.** Median clinical trial enrollment time rose from 13 months to 16 months between

2021 and 2023 before stabilizing (IQVIA Institute), and the labs now name drug companies such as Novo Nordisk and Sanofi as customers (Anthropic).

- **Headwinds: Public research software makers lose value as soon as a lab launches a research tool, which shows how quickly the labs can reprice this work.** Shares of Schrodinger fell as much as 8.3% on the day Anthropic launched its Claude Science research workbench (trade press).
- **OpenAI launched a life sciences model with Amgen, Moderna and Thermo Fisher, and Anthropic launched a research workbench for scientists in June.** OpenAI's Life Sciences plugin connects the model to more than 50 scientific tools (OpenAI), and Anthropic's Claude Science combines scientific databases, code and compute in one desktop workbench (Anthropic).
- **Why we should focus here: Life sciences research pays for discovery work in a regulated industry, and an agent that trains on a drug company's own data does work a lab workbench does not.** Results take years to prove, so we should back companies that already show repeat paid runs at the same customer and price each run against the program decision it informs.

Example companies: Edison Scientific

C. Non-Regulated Verticals

C1. Procurement, Logistics and Distribution

- **Tailwinds: Purchasing and freight coordination employ big teams whose work runs on email, portals and phone calls.** Government data counts 491,430 buyers and purchasing agents, 202,810 dispatchers outside emergency services and 97,670 cargo and freight agents (BLS).
- **Procurement buyers want help with turning commercial strategy and demand changes into sourcing plans, which today runs as a manual and disconnected process.** A Honeywell sourcing manager says his biggest pain is the manual and disconnected process of translating demand changes into sourcing forecasts (expert call).
- **Headwinds: At least one big buyer, Honeywell, builds its own procurement agent on top of the models, which caps how much a vendor sells into accounts like it.** The same Honeywell manager says his team runs its own internal AI agent built on internal data and a ChatGPT model (expert call).
- **ERP and procurement suites now ship agents inside the record, so any vendor has to work alongside the system the buyer already trusts.** Oracle launched 22 Fusion Agentic Applications that act inside its business processes (Oracle), and a PayPal procurement executive says "We weren't trying to replace our system of record" (expert call).
- **OpenAI is testing its own specialist agents for procurement, invoice processing and commercial contracting, which sets a lab baseline for the simplest purchasing work.** Its September developer conference previewed specialist agents with their own identity for those roles and a marketplace where buyers spend their OpenAI commitment on 32 partners (OpenAI).
- **Why we should focus here: Procurement, logistics and distribution work runs across suppliers, carriers, phone calls and email outside any single record, which makes it hard for an ERP vendor or a lab agent to absorb.** The agents sit on the buyer's ERP, so results are counted in a record the buyer trusts, and value grows with the spend the agent handles. We should back companies that clear exceptions end to end and show touchless rates at named manufacturers, freight operators and distributors.

Example companies: Fragment, Magentic, Didero, InstaLILY, BackOps, HappyRobot

C2. Finance and Accounting Operations

- **Tailwinds: Finance teams repeat the same work every month with big clerical staffs, and each invoice or reconciliation has a clear count and cost.** Government data counts 1,373,680 bookkeeping, accounting and auditing clerks at a mean wage of $53,560 and 404,060 billing and posting clerks (BLS).

- **Payment networks now let agents complete business payments, which lets an agent finish the payables job instead of handing it back.** Visa runs a live pilot that "enables trusted AI agents to initiate and complete approved business invoice payments through supplier portals" (AIM Media House).
- **Headwinds: Oracle is building agent orchestration into its Fusion ERP, which limits how far a new vendor can reach inside Fusion finance teams.** Oracle sells its agents through paid AI units, with a rules layer validating every action the model proposes (trade press).
- **Agents cost more to run than the scripted bots many finance teams already use, so vendors must show savings beyond the old automation.** An integrator's automation lead says a task that costs $200 or $300 a month on bots costs $1,000 or $1,200 a month with AI (expert call).
- **Finance processes amplify errors, so speeding one step without fixing the next creates more work downstream.** Deloitte warns that "Agents can amplify broken or inefficient processes" and that faster upstream steps can overwhelm manual approval downstream (Deloitte).
- **Anthropic's financial services templates now cover the generic monthly close directly, which removes the opening for generic close tools.** Its spring release includes a month-end closer, a general ledger reconciler and a statement auditor alongside KYC screening (Anthropic).
- **Why we should focus here: Finance and accounting work recurs with the monthly close and is checked in the buyer's own books, and it is staffed by clerks and outsourcers paid from the labor budget.** ERP vendors and lab templates cover the generic close, so we should back companies that serve accounting firms, global payroll or a specific industry's books, where rules and liability sit outside the ERP and the company takes accountability for each invoice or filing.

Example companies: Basis, Freehand, Niural

C3. Customer Service and Sales

- **Tailwinds: Customer service is the first job agents complete end to end at scale, because a resolved ticket is easy to define and check and the labor it replaces is easy to price.** BLS projects US customer service employment to decline 5% from 2025 to 2035, to 2,524,100, and cites automation and self-service as the cause (BLS).
- **The big contact center outsourcers are rebuilding around AI, which shows the labor budget moving in this function first.** Concentrix says "50% of our revenue is coming from business we have won and deployed within the last 3 years since the introduction of AI" (Concentrix), and TP's core services revenue declined like for like in the second quarter (TP).
- **Headwinds: Salesforce bundles its own sales agents to replace point tools, which squeezes standalone vendors in its accounts.** Salesforce says its Agentforce 1 Edition at $550 replaces roughly $1,300 of point-solution spend per US sales rep (Salesforce conference remarks).
- **OpenAI now sells its own customer service agent deployed by its own engineers, and Salesforce already sells a help agent priced per resolution.** Salesforce charges $2 per resolution for its help agent (CX Today), and OpenAI launched Presence in July with written procedures, guardrails and automated scoring, and says it resolves 75% of inbound issues on its own phone line without human assistance (OpenAI).
- **Why we should not focus here: Customer service proves that agents can finish a whole job at scale, but OpenAI and Salesforce now sell that job themselves at a published price per resolution.** Sales agents are paid on meetings and leads that are easy to game, and Salesforce bundles its own. We should use customer service resolution rates and prices as the benchmark for other functions and make no new bets in the category.

Example companies: Sierra, Decagon