tegus by AlphaSense | Expert Insights

COMPANY DEEP-DIVE

Former Chief Product Officer at *Chime* Believes AI Tools Like *Decagon* and Sierra Are Transforming Customer Support Dynamics

DATE PUBLISHED	INTERVIEW DATE	EXPERT PERSPECTIVE	ANALYST PERSPECTIVE	PRIMARY COMPANY
16 Aug 2026	18 May 2026	Customer	Investor-Led (Buy-Side)	DECAGON.AI

SMART SUMMARY

Summary

The interview covers the expert's approach to evaluating and implementing AI-driven customer support tools, comparing vendors like Decagon and Sierra, and discussing the operational, technical, and strategic considerations involved. Topics include vendor selection, integration, cost optimization, and the evolving landscape of AI in customer support and related functions.

EXPERT BACKGROUND

- Former Chief Product Officer at a leading US neobank, with extensive experience evaluating and implementing AI solutions for customer support, including direct work with vendors such as Decagon and Sierra, and prior product leadership roles at other technology-driven companies.

EXPERT SENTIMENT

- The expert believes the AI customer support tool landscape is highly dynamic, with no clear long-term winner, and emphasizes the importance of continuous benchmarking and flexibility in vendor selection.
- **They see significant potential for cost reduction and efficiency gains**, but stress that the real value comes from optimizing customer satisfaction and resolution rates, not just cost savings.
- **The expert is optimistic about future margin expansion opportunities for vendors** as models improve and as solutions expand into adjacent operational areas beyond customer support.

TAKEAWAYS

- **The primary driver for adopting AI tools like Decagon and Sierra was to improve resolution rates and customer satisfaction, rather than immediate cost savings**.
- A router system was developed to dynamically assign cases to different vendors (Decagon, Sierra, others) based on performance, case type, and ongoing benchmarking, with frequent reevaluation as technology evolves.
- **Both Decagon and Sierra have distinct strengths: Decagon excels in speed and transactional efficiency, while Sierra is stronger in complex, multi-turn conversations and customer loyalty focus**.
- Integration with internal systems is managed by the company's own engineering team, ensuring all vendors have equal access to necessary data and tools, and making switching or adding vendors relatively straightforward.
- The expert anticipates that as AI models improve, more routine or low-risk cases may be handled by internally built solutions, while premium vendors will need to differentiate on complex, high-value cases and technical capabilities.

AlphaSense Assistant is AI Generated

Table of Contents

Introduction and Expert Background

Initial Evaluation of Tools

Comparison of Decagon and Sierra

Routing and Integration Strategies

Future Trends and Optimization

Technical Challenges and Market Dynamics

Expansion Opportunities and Margins

Voice Channel Considerations

Expert Bio

Former Chief Product Officer at Chime, leaving in April 2026. Former Chief Product Officer at Chime, leaving in April 2026. Reporting to Co Founder, the expert is responsible for product management, product design, user research and product operations teams at Chime. Prior to

©2026, AlphaSense, Inc. All Rights Reserved. AlphaSense is a service mark of AlphaSense, Inc. All other trademarks mentioned belong to their respective owners.

Chime, the expert was the Chief Product Officer at Notion Labs, leaving {May, 2023. Reporting to the Founder, the expert was responsible for the product management, product design, research and product partnerships teams at Notion. The expert {can speak to managing large teams of over 100 people within the product, partnership, and research space. Has over 15 years of experience in building big named products in industry like mobile, web, fintech, design, etc.

Interview Transcript

Client 00:00:00

Really appreciate your time today. As we get into things, really want to talk about the competitive set, sounds like you have a bunch of knowledge about Decagon as well as Sierra, and just understanding the other players in the market. Maybe just to kick things off, it'd be great to hear a little bit about your background and just your experience evaluating the tools broadly.

Expert 00:00:21

Sure. I recently was the chief product officer at Chime, the largest neobank in the United States. We were also on the very forward edge of adopting, evaluating, and rolling out AI solutions across the company. Before that, I led product at Notion, which is a B2B SaaS productivity product, but notably, very AI-native.

My background is mostly in product, and over the last, let's call it two and a half to three years, I've been quite close to teams across OpenAI, Anthropic, and now their vendors and partners, including Decagon, Sierra, etc. Specifically with these, I would just say, experience has been over the last, call it, 12-18 months of evaluating them and applying them across text-based customer support, but then also voice-based as well.

Client 00:01:12

That's helpful. Would love to dive into the differences, but just to start, it would be helpful to understand what was the process like at ***Chime*** before you implemented ***Decagon*** and Sierra, and what was the catalyst to pivot and decide to either outsource or change from a BPO? What drove that decision to look at these?

Expert 00:01:34

Totally. This is very similar for all companies, so I'm not going to speak directly to Chime. We're a regulated fintech. That means you have customer support that's being offered by humans. For the most part, these are pretty complicated cases, so you've already built tools to support those agents. Tools that include very simple things like how do you track SOPs and make sure that agents can keep track of the steps they're taking, all the way to more complicated tools, things that recommend the next step or automatically pull up data that they're going to need during the conversation.

Most companies are in the same spot, which is that they've moved operations to some centralized centers. I'm happy to chat about where Chime was, which is, we had BPOs largely in India and the Philippines, and we'd built a lot of custom tools for them. That said, we also had some deterministic tree-style automations that we already had in place, and then we were looking to drive up the efficiency and effectiveness of those and drive down the cost.

Client 00:02:40

Got it. As you started to evaluate the tools, could you talk about what the key ROI points you were focused on, and to the extent that you've got any insight into how important pricing and how you thought about margins were, versus these human BPO call centers, versus the tools that you're implementing with Sierra and Decagon?

Expert 00:03:01

Totally. I think the first thing that we were trying to understand is how real is real. This started end of 2024 into 2025 and 2026. In 2024, there was **quite a lot of bluster, but maybe there wasn't quite as much proof** just yet that these systems could not only handle very simple things like, "What is my balance?" but could actually handle huge complex interactions and cases. That was the beginning, was just to try to understand what's possible.

To that end, I'm not sure cost was the major concern. There was a belief that, over time, if you could move away from more human-centered service delivery and you could move it to a hybrid model, that would probably be cost-effective, but I wouldn't say that there was a huge driver there. We just knew that **the cost would come down** pretty dramatically.

The third thing I would call out is just applicability overall. Not every service could be rolled out to every type of customer support problem in 2024, so we were looking for people who had already had the ability to start with the easy things. By that, I mean simple queries that could have been answered by a website, before moving our way into the more complicated things.

©2026, AlphaSense, Inc. All Rights Reserved. AlphaSense is a service mark of AlphaSense, Inc. All other trademarks mentioned belong to their respective owners.

Client 00:04:18

Got it. Not even just specifically to Chime, but would be helpful for my background and understanding, with these human BPOs, are you typically paying for a 1099 contract for headcount? Are you paying by conversation? Could you talk a little bit about that pricing structure, and if it's comparable for a Decagon or a Sierra?

Expert 00:04:41

It totally varies. I've been in places where you have full-time employees because from a regulated perspective you need them. Before this, I led product at Robinhood. Robinhood is a brokerage. Brokerage regulations do require a level of oversight that requires those people to be in the U.S. and to have full oversight, and then there are places like Chime where you have BPOs, so you're paying differently. You might not be bringing those people onto staff. You're doing a contract directly with the BPO that's usually renegotiated over total number of hours and cases, that kind of stuff.

Client 00:05:19

Got it. How often are you reevaluating which tools you're using? Is it every year when you go through contract renewals? Are you out testing different products, or do you feel like Sierra and Decagon are leading the pack and they'll continue to be the top tools for a while moving forward, so you haven't had to do RFPs?

Expert 00:05:41

Yeah. I'd separate these two things, which are, how often do we evaluate tools, and how often do we evaluate tools in this space? The reason I say that is because we know that this space is moving relatively quickly. Our COO might spend a lot of time looking at tools on a two-year, three-year basis just to renegotiate contracts, understand new volume, try to find discounts.

From the product and engineering lens, anything that says, "Hey, look, we can drive down a tremendous amount of spend, we have this new capability," especially AI-driven things, we're almost always open to look at them. I would say, right now, **we are probably in a very frothy environment** where weekly, monthly, some new AI for disputes, AI for fraud, AI for customer support tool will come up. Unfortunately, we're probably eyes and ears open on a lot of that, just given the promise.

Client 00:06:41

That makes sense. Would love to dive into the competitive piece, but maybe just to start with Sierra and ***Decagon***. As I understand, you use both of those at ***Chime***, so it would be great to understand what it was about these two products separately, and then as a comparison against each other, what made you select the tools over the other competitors in the space?

Expert 00:07:02

Totally. Like I said in the beginning, our main goal was not to save costs. It was to understand what tool can do what, how well can it do, what percentage of our total support volume it could take. That started with an RFP-style canvassing. These companies are relatively small, so a lot of it did come from word of mouth. Again, this is 2024 into 2025, so it's not what we see today.

The beginning of that canvassing got us to probably four or five companies. That group would include Maven AGI, Google had some older technology that was doing decision trees, and then we had our own internal tools that could also help answer some questions. You should assume that all of those things were in the competitive set.

What we were looking for is to get an understanding, first focused entirely on chat, of three things. One is, for any given case type, how high could the resolution rate on that case be? Number two is, what is the satisfaction of the customer beyond what that resolution rate might imply? Third is the cost to handle that case. A lot of this was just for baselining purposes to try to get a sense of what is good, and what I'd say is that Sierra and Decagon both spiked in very different directions.

Sierra is quite good at these longer conversational multi-turn types of cases and tickets. **Decagon was very, very good at accurate, quick, fast responses across the whole corpus**. There were cases where you might want to spend a little bit more money. Assume that the customer has been with you for a longer time or the churn likelihood on that specific case is very high, you might be okay with paying a little bit more to get a higher resolution rate, higher satisfaction. In other cases, you're just trying to accomplish a task, let's say you're closing an account, speed, accuracy, and cost would matter more. Those are the kinds of ways that we would look at these problems.

Client 00:09:11

Makes sense. How did you decide which set of problems you were going to have Decagon solve versus Sierra? Is it just a 50/50 baseline?

©2026, AlphaSense, Inc. All Rights Reserved. AlphaSense is a service mark of AlphaSense, Inc. All other trademarks mentioned belong to their respective owners.

Expert 00:09:22
Yeah. At first, we built basically a router system, which allowed us to have control over all of the data and then all of the assignment. When we started this, we had a ticket come in, we would assign it to a human agent. Pretty simple, straightforward. The ops teams can change rules, no problem.

Imagine a modification on top of that, which is, for certain ticket types, ones we felt comfortable testing after a while and had regulatory approval to test, we would say, "Okay, instead of handing this to a human agent, let's hand this to Decagon," or "Let's hand this one to Sierra," or "Let's hand this one to Maven." That was done randomly in a controlled way, so that we could monitor and baseline the results, cut against our own data.

Client 00:10:10
Got it. Today, are there different types of problems that always go to Decagon versus Sierra, or is it based on capacity? How do they get routed?

Expert 00:10:21
Yes. There's still a router, and there are lots of things that go into the router and the router's decision-making. This is where I would say, very forward-looking technology companies will operate differently than everybody else, and I would put us on that front edge of this, given that I know a ton of other customers who don't do this.

When we built this router out, our goal was to understand the baseline, just to understand who's actually works versus who doesn't, what's the cost of that, what types of problems work, what types of customers like these things, but you can imagine the variables start to move. For example, as soon as we make a tool available to close an account, we might want to try closing accounts. As soon as voice is available, we might want to try voice. With every one of these permutations, you get a level of complexity of, "Well, actually, this tool is better than that tool," so you want to push them.

The second part is you want to continue to maintain the measurement because they're evolving very quickly. Every new model that comes out, every new feature that **they add, does actually have a material impact** on how well they perform on case types that you might have seen a week ago or a month ago. We were not cost optimizing, we were visibility optimizing, and that means even today across text channels and voice channels, we still employ multiple vendors, so that we can understand what's performing well, what's not performing well, where to send majority of flow, and again, that changes on a week-to-week, month-to-month basis.

Client 00:11:54
Got it. When you talk about, for example, the tool that closes the account that you've built, are you building that inside of Decagon and Sierra, or is that being built inside of different internal systems that you no longer have to route to a Decagon or a Sierra?

Expert 00:12:11
Great question. The way to think about the world is, again, not sure everyone will think this way, but we certainly did, there's a core set of, let's call it, fuel that any of these agents is going to need. It's going to need data where it's just like, "Who is this customer, what's their history with us?" etc. They're going to need a bunch of context that the company offers, so our SOPs, our policies, how our products work, and then they're going to need some tools. "Hey, can I pull up a balance? Can I close an account? Can I send an email? Can I send a notification?"

These are all generic, so for us, we kept these all inside of Chime systems. We're very happy to connect any provider to that corpus. Let's say tomorrow you decided you wanted to start a voice agent company and **you came knocking on our door**. We'd give you the access to the same corpus, but we were not going to put anything unique into any of the vendors, outside of just wiring it all up and saying, "Okay, we'll do a little bit of custom work to make sure you can plug into this." Does that make sense?

Client 00:13:23
Got it.

Expert 00:13:24
The tools we're talking about here, these could be APIs that we make available internally, these could be buttons that show to a tool somewhere that they can access, but broadly speaking, these are just capabilities that we're offering any of the vendors.

©2026, AlphaSense, Inc. All Rights Reserved. AlphaSense is a service mark of AlphaSense, Inc. All other trademarks mentioned belong to their respective owners.

Client 00:13:39

Got it. Is it your internal engineering team that's creating these tools, or are you working with Decagon's forward engineers, Sierra, etc. ?

Expert 00:13:51

No, it's ours. Decagon's forward engineers are mostly working on integrating into that tool or in times telling us, "Hey, look, we can't solve a problem unless you expose something to us." We might take that as input and say, "Okay, great. Let's go build a tool that enables that," but as soon as we build the tool, we're making it generic so that Sierra can plug in, Vapi can plug in, whoever.

Client 00:14:17

Makes sense. Do you have a perspective on whether Decagon or Sierra or a different system have been better at plugging into your tools?

Expert 00:14:28

Not really. They all have to get good, and they all forward deploy anyway because if they can't use our contacts or they can't get to the data, then they can't solve cases. For the most part, both parties are incentivized to make this pretty easy to use, and it's not a differentiator.

Client 00:14:46

That makes sense. Apologies, I'm not a technical background, do **you have a sense for how difficult it is** to develop this set of tools? If a new entrant comes into the market, how far behind do you think that they are, to be able to catch up to a Decagon or Sierra and be able to have the same success?

Expert 00:15:04

Not that bad. Let's take a big step backwards. Let's say we were still in the world where there was no AI. You still had to make these things available to your human agents. Otherwise, your human agent can't do anything. Let's say, tomorrow, you want to refund customers, give them $20 of credit to your shop if they had a problem, you have to find some way of enabling that. That enablement will almost always come through some engineering team, whether it's yours or a vendor's, and then you just expose that to Sierra or Decagon. These are not things that Sierra and Decagon are going to build.

They're going to instead put a forward-deployed person against it and say, "Okay, how do we connect our pipes to these pipes?" and that doesn't tend to be too bad. They're good and quick, and like I said, everyone's incentive is the same because if they can't plug in, then they can't solve problems.

Client 00:16:03

Got it. I know we've talked a lot about chat. I'd love to get into voice in a minute, but just for chat, if you could only stick with Sierra or Decagon, do you have a perspective on which you would keep?

Expert 00:16:17

No. Again, the technology is constantly moving. If we had had this conversation one month after we started the test, we would've told you, "Oh my goodness, Decagon is doing an amazing job. It's really fast and it's really accurate, and it's kicking everyone's butt on cost." If you were to check in two months later, we'd say, "Whoa, the new models are out. Sierra's models are doing really great. We're seeing a change."

I would liken this to, we're a horse race. There are a lot of horses. They're all doing slightly different things. Maybe NASCAR is a better analogy. I don't know if you're a sports person, but everyone's supposed to use the same kinds of engines, but one of these cars is going faster than the other. Does that mean that it won't change? No. It's absolutely going to change. We know that.

I would say the skills of these two companies seem different. Sierra definitely is a little bit more differentiator on the go-to-market side because Bret is well-known, he's the chairman of the board at OpenAI, he has worked at Salesforce, so he has this ability to get into companies quickly. They have a great, very talented team.

Decagon, on the other hand, is a little bit more upstart, Silicon Valley scrappy, work hard, burn the midnight oil, and make sure the teams that you work with are really happy. Their approach is a little bit more innovative for the time when we first started this. They focus on different things. I would not bet on either.

In fact, I'll add to your chaos a little bit. One of the other reasons we keep the benchmarking open and the router open is because we can build and deploy our own internally as well. As the local models get better and better, there's probably a subset of cases and tickets that could be taken care of by an internally-built model as well.

©2026, AlphaSense, Inc. All Rights Reserved. AlphaSense is a service mark of AlphaSense, Inc. All other trademarks mentioned belong to their respective owners.

Client 00:18:21
Got it. Do you have a sense for how close we would be to the point where you'd start using internal models to take care of, for example, closing accounts? Sounds like that's pretty repetitive. Is that in the next six months? Is that years away?

Expert 00:18:39
Yeah. Probably. Let me try to explain maybe the variables, so that this doesn't come off as being doomsday-y. In the beginning, you can only really address, I don't know, 5%, 10%, 15% of your total tickets because you don't have all the SOPs, you don't have all the data, you haven't exposed everything, you're not in every channel, not everything's been approved by banks or regulators, etc. Over time, that denominator is going to grow. Because you can handle more and more case types, you'll set up more and more things on your own, you'll have more and more capabilities. As those things are approved, you'll start to see differentiation inside of that.

The denominator is almost constantly moving, in that we're exposing more and more things that used to never get to us. For example, if you put a chatbot inside the app, you're going to get a lot more outreach just naturally, than if you didn't have a button for a chatbot in the app.

I would just say this is always a game of optimization, and I would imagine what the future is going to look like is, at least for forward-looking companies, a pretty stratified delivery pattern. Some things that are super-low-risk, super-functional will probably be controlled by models that we build. Some of the less complicated things will go to a low-cost provider. Some of the really high-stakes, high-complexity, high-value things will go to a more premium provider or premium product.

Client 00:20:12
Got it. How do you become the winner or the premium product? Are there specific either things that you think won't be easily in-housed, or is it speed? What are the things that help you continue to charge a premium as opposed to getting more commoditized and being able to take in-house?

Expert 00:20:30
I think it's, at the very macro level, two or three things. The first is, are you resolving problems, do customers like it at a higher rate, and is the cost competitive with putting a human on the task? If the answer to those three things continues to go up and to the right, then we feel great. **It's just ROI-positive as a business** to continue to give to that company, even if it costs more because the LTV of the customer might be really, really high. Paying under $25 a case, which is much more than you would pay any human, might make sense in certain situations when you have a super-high-value customer or a new customer that's going to be very engaged in the future.

I think it really comes down to those things, at least for us. We're looking at resolution rates, so that tickets don't reopen, satisfaction, so customers are actually happy and stay, and then cost. What that means under the hood, one more step, because I think you're implying it a little bit, is, these companies have to make decisions. They have to make serious technical and product decisions as to what they want to prioritize. For example, they can make their agent really, really rich with context.

It can know everything about the customer who called and be very, very dexterous in using their name and the last 10 times that they contacted them, but that comes at a cost. Is that actually going to yield better outcomes? We don't know. Likewise, someone can optimize for speed and supporting every language on the face of the earth and doing it in a very direct tone to save time. Those are choices that these companies are making in their products, and that'll enable them to be more or less competitive over time.

Client 00:22:20
Got it. On that piece where you talked about understanding the client and being able to make decisions for a higher-value client that might be a little bit more complex, is that stuff that you're seeing already in Decagon and Sierra, or is that a future roadmap? How idealistic is that versus something that's in place right now?

Expert 00:22:41
Right. I want to separate their tools from our abilities, so the best way to think about them is to just keep thinking about them almost like people. I know it's counterintuitive, but you have an agent that's provided to you by Decagon. The Decagon agent is X amount of smart, Y amount of capable, Z amount of fast, and then the one presented to you by Sierra is the same. Of course, you are giving those agents directions, you're telling them what the SOPs are, etc., but the agent themself isn't providing you things like routing, not at least as of yet.

Do I think that they will build their own routing? Absolutely, just for the pure play cost optimization reasons that we've seen large LLM providers like Anthropic and OpenAI do, which is that not every query needs to go to the most expensive model. In fact, if it does, people complain, "Oh my God, look at the cost of this. Why does this cost so much?"

The second part is they'll get undercut. If you look at Gemini, for example, in the LLM space, it's much cheaper to use on a daily basis than some of the other models, and if they don't bring their prices down, they'll get out-competed. I think that same dynamic will make its way into

©2026, AlphaSense, Inc. All Rights Reserved. AlphaSense is a service mark of AlphaSense, Inc. All other trademarks mentioned belong to their respective owners.

the customer support space. Does it exist today? No, not yet, because I don't think the industry's matured. For us, people who are building on the front edge, we're already optimizing under the hood. Why would we just use one-blanket approach? We have the technical talent to do it.

Client 00:24:21

That makes sense. As you talk about that routing piece, do you see products that already exist, like a Zendesk, or I don't know if it would go through a HubSpot, but companies like that creating these solutions?

Expert 00:24:35

Not really. I think e-commerce is probably the closest, where you have abandoned cart-types of companies. You have retention loyalty-types of companies. Those places are a little bit more forward-looking, because they've had to deal with this problem before. For the most part, large, let's call them fintechs, banks, etc., have not really had to think in this way because bringing the data together across all these disparate systems is quite hard. It's coming for two or three reasons. One, because all the data will get unlocked. Two, it's coming because otherwise **it becomes cost-prohibitive really, really fast**.

That I'm sure of. I don't know how many weeks or months, but every company will get to a place where they figure out some distribution of how to route to the higher cost versus the lower cost. Actually, let me take a step back. They're doing it today. If you call Chase, Chase has already tried to stop you with a help center. They're probably trying to get you to use a telephone tree before they give you to a human. It's that same logic again, just at the LLM layer this time.

Client 00:25:52

That makes sense. Would it be fair to say that the moat right now for a Decagon is its ability to bring all of the data together?

Expert 00:26:02

No. Unfortunately, that's the company's problem. If **a company has a scattered mess**, no amount of Decagon prayer or OpenAI seances will help you. You have to bring that yourself. I think the moat for these companies right now is twofold. Decagon's is, you can get started very, very quickly. I think the first time we saw a demo, they did it in under 24 hours, when we first gave them any access to any data. That demo was just a very simple ask questions, get responses, type of chatbot. Worked exceptionally well, was correct, accurate, found issues in our own documentation, etc.

I think Sierra's moat is a little bit more technical and conversational. They want to push the brand out to the customer a little bit more. They think a lot about what the interactions do for lifetime loyalty. If you talk to Bret and team, they'll tell you a lot about how they think about customer support and marketing being two sides of the same coin.

I think the moats are really technical moats. **It's very difficult** for some people to build what they're building. It sounds like it would be easy, like, "Hey, just plug in OpenAI." Turns out it's not that easy. The second part is they're very focused on delivery in slightly different ways. I think Decagon may be a little bit more transactional, "Hey, we can help right now, **we can bring your costs down**, we can scale." Sierra is maybe a little bit more artisan.

Client 00:27:38

That makes sense. This is probably going to be a dumb question, I'm not sure you can answer it in three minutes, but when you say it's not that easy to build on the back end, is that governance? What makes it difficult?

Expert 00:27:52

Exactly. In this market, OpenAI is entering it directly. They're working directly, I believe, with T-Mobile. ElevenLabs, a voice provider, is working directly with Revolut. It's not that the model companies won't try to do this, but it's work for them too, and it's exactly what you pointed out.

It's, "How do I make sure that I don't give the wrong information? How do I make sure that I can update an SOP and pass it through to an agent immediately? How do I make sure that I stay on the brand voice? How do I make sure there's no hallucinating?" etc. In a customer support setting, some of those things are much harder than in a chat setting. "How do I do this with low-latency voice communication?" Just turns out it's hard. The trade-offs you have to make there are very different than the ones you might make on a watch.

I think there's a lot of customization. There's a lot of hand-holding. There's a lot of reporting that you have to do. Maybe there's a regulator, maybe there's a state, maybe there's another department. Plugging into the existing operations teams and then ensuring a specific level of support is pretty important. Business continuity planning. There are a bunch of other things.

©2026, AlphaSense, Inc. All Rights Reserved. AlphaSense is a service mark of AlphaSense, Inc. All other trademarks mentioned belong to their respective owners.

Client 00:29:13

Got it. You mentioned that Decagon's moat is speed. If Sierra were to get faster and be able to answer more quickly, would **you replace Decagon, or is there a benefit to having two tools** still?

Expert 00:29:26

Exactly. I think we see the world very differently. I think we're in the first inning of a nine-inning game here. For the companies **that really want to see a cost reduction**, they can do it today. They can just go in and pick either one of these things. **They'll see a giant reduction** compared to their humans.

I think if you're trying to play a longer game and you're trying to understand, how do I give my customers the best support I possibly can, how do you understand this technology to figure out, "Hey, do I need a vendor forever or can I build parts of this?" then you wouldn't cost optimize today.

Now, we're very fortunate. **We're in the position** where we're a technology company and we want to invest for the future. If I were a legacy business and I really needed the margin tomorrow, I would just pick one of these companies, have them implement it, and then move on, but again, I think that comes down to who you are as a customer.

Client 00:30:24

That makes sense. Maybe one more question if you've got time for it. **It sounds like there are huge cost savings from the human perspective**. Is there a way for Decagon to drive greater margins, or is this only going to become compressed over time? What would we have to see for that to happen?

Expert 00:30:44

Yeah. I think it's expansion in two ways. The first is into other processes that are similar to customer support, but themselves are not customer support. I was mentioning this earlier, fraud investigations. That's very similar in that it's SOP'd, it has a lot of oversight, it has some non-deterministic problems that require judgment.

Could they make their way away from just customer support into fraud investigations or into disputes, where a customer tells you, "Hey, I didn't actually get a haircut. You should give me the money back"? Various businesses will have various versions of this. Internal help centers, I think fintech is easier for me to speculate on, that's huge. Those are sometimes as big, if not larger, cost items than customer support is.

I think the second is, get out of this function specifically and migrate into other functions. By that I mean, right now, you're solving an operations problem. Once you build trust that you can use this data, you know how to do things, you might find yourself in a finance team's needs, like doing some treasury workflows, or a legal team's needs in doing some sort of cleanup on all active caseload that you're managing. It's too early to tell where these companies will go. Do I think there's more margin? I definitely do, because the models will get cheaper and they'll get better at this.

Maybe **the last thing, probably the most important question to ask is, is there a benefit to these companies in getting the horizontal learning across multiple companies** who are serving customers? If it's the case that when you do the 20th company or 25th company, you unlock something special where you learn the system better than they can learn the system, whether **that's patterns of fraud or very specific things that are counterintuitive**, then great. These folks will be able to provide even better service than the company could figure out on their own. **Those are, again, margin expansion opportunities**.

Client 00:32:53

That's really helpful, appreciate it. If you have time for one more question, maybe on the voice side, are you using both Decagon and Sierra on voice?

Expert 00:33:03

Yep. Plus a couple other companies, it's also a router.

Client 00:33:08

Got you. How do you think about the opportunity to continue to drive that piece of the business? Is it very similar to chat? Are you trying to push more calls into chat versus voice, just given the cost? How do you think about the way that those two channels work?

©2026, AlphaSense, Inc. All Rights Reserved. AlphaSense is a service mark of AlphaSense, Inc. All other trademarks mentioned belong to their respective owners.

Expert 00:33:25

It's cost optimization. On one end is the cost of sending a human being to the customer's house and having them solve a problem, on the other end is a pop-up on a website, and then everything in between that you can imagine, chat, voice, etc., is all there. We want to match the service to the capacity we have to deliver that service. We can't spend $10 fixing a problem that's going to bring us back $2. **We can spend $10 solving a problem** that's going to bring us back $1,000. I just think of it as one more tool in the tool belt, and yes, we are going to try to push customers to the right place.

You really don't want people calling you on the phone and asking you every single day what their balance on their account is. There's no way to provide that service in a way that that's cost-positive. You could just think about everything I said about chat is pretty much the same on voice. The second part of that is, customers have expectations, and we have to meet them too. If they call and they need something, we can't just say, "Oh, we don't offer this on our AI voice. Sorry." That's quite bad.

Client 00:34:36

Got it. I really appreciate the time. This has been incredibly helpful. Hope you have a great rest of your day.

©2026, AlphaSense, Inc. All Rights Reserved. AlphaSense is a service mark of AlphaSense, Inc. All other trademarks mentioned belong to their respective owners.