AlphaSense Highlighted Document

Demystifying Agentic AI with Agent Teardown; Agents Represent Major Unlock for Enterprise AI

Tue Apr 28 2026 ASR

User Highlights & Comments

Equity Research Report

April 28, 2026

TMT: Application Software and *Agentic AI*
TMT: Infrastructure Software
TMT: Security Software
TMT: Artificial Intelligence, Digital Platforms and IT Services
TMT: Internet

Terry Tillman
404-926-5362
Terry.Tillman@truist.com

Miller Jump
212-401-4427
Miller.Jump@truist.com

Junaid Siddiqui
212-326-6141
Junaid.Siddiqui@truist.com

Youssef Squali
212-303-4180
Youssef.Squali@truist.com

Arvind Ramnani
214-969-5873
Arvind.Ramnani@truist.com

Connor Passarella
212-326-6153
Connor.Passarella@truist.com

111 Page Document

Reasons for this report

- ✓ Deep Dive Report
- ✓ Sector Update
- ✓ Sector Read-Throughs
- ✓ ***Agentic Teardown***

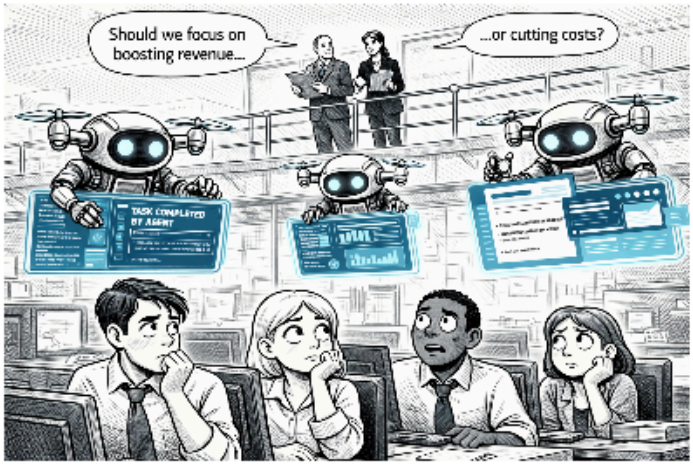
Demystifying Agentic AI with Agent Teardown*; *Agent*s Represent Major Unlock for Enterprise *AI

Agentic AI stands to transform existing software markets across apps, infrastructure & security while creating substantial new monetization opportunities across broader ***AI*** ecosystem & economy. Autonomous work agents could unlock trillions of dollars of value. Agentic AI opens up much larger budget dollars, well beyond IT & innovation budgets to vast dollars associated with $60+ trillion in estimated global labor spend & other operational sources to fund agentic transformation. Along w/ this report aiming to provide a guide to define agents, we share clear prescriptions on potential AI winners & picks and shovel ideas across the broader Truist Securities TMT coverage.

In this report, we profile the two largest AI/agentic pure plays, OpenAI (Private) and Anthropic (Private) along with the most consequential agentic picks and shovel ideas as we see it across the broader Truist Securities TMT research coverage. This includes Microsoft (MSFT, Buy), Salesforce (CRM, Buy), Snowflake (NOW, Buy), ServiceNow (NOW, Buy), Databricks (Private), CrowdStrike (CRWD, Buy), Cloudflare (NET, Buy), Palo Alto Networks (PANW, Buy), Zscaler (ZS, Buy), Alphabet (GOOGL, Buy), Amazon's AWS (AMZN, Buy), Meta (META, Buy), Palantir Technologies (PLTR, Buy) and Accenture (ACN, Buy).

Agentic AI Explained. Building on generative AI and powerful advances in LLMs, AI agents represent the next major unlock in enterprise AI adoption, in our opinion. Software agents are akin to digital workers and autonomously plan and execute a wide range of work, from simple tasks to advanced multi-step business processes and intelligence that offer significant business transformation potential. In this report we define agents and agentic architectures, share implications of LLMs and agent innovation/disruption on System of Record (SoR) and SaaS companies along with development, data and security software markets, & size the agentic market opportunity.

Enterprise Agents = Best Opportunity to Recoup Trillions in Data Center & Infrastructure Spend. NVIDIA's CEO Jensen Huang has suggested data center buildout for AI to scale to $3 trillion - $4 trillion by 2030. McKinsey had estimated $7 trillion in total capex needed by 2030 for AI. Agents represent a direct monetization opportunity related to over $60 trillion in potential global labor spend (labor typically 50%+ of total GDP). We conservatively estimate at least 10% of global labor spend could represent source of investment for enterprise agents. This could translate into $6 trillion+ agent monetization and would start to pay for the substantial anticipated AI data center and infrastructure buildout in rather quick and dramatic fashion, in our opinion.



Source: Image created by ChatGPT

Table of Contents

Executive Summary 3
What is an Agent? 5
Agent Architecture 6
Proliferation of Agents = Need for Interoperability, Management and Observability 8
Protocols: MCP and A2A; Important for Enterprise Agent Success 9
Agentic Application Design Patterns and Orchestration; How Can Agents Actually Be Useful? 11
The Shift to the Generative UI: The Final Mile of Delivery for Agents 13
Agentic Implications for Application Software Incumbents, LLM Threats, Model Changes, Etc. 14
Sizing the Agentic Market Opportunity 18
Application Development (SDLC) Layer 23
Data Layer – Critical to Enterprise Agent Success 24
Observability 27
Agent Security – Another Major Requirement for Broad-Based Agentic Traction 29
Risks and Additional Considerations Related to the Agentic Opportunity 37
Agentic Stack Picks and Shovels Plays 40
- OpenAI – From Intelligence to Execution, Building Agentic Workflows at Scale 41
- Anthropic – Building Safe AI for the Enterprise Era 47
- Microsoft – From Copilots to the Enterprise Control Panel 52
- Salesforce – Operating System for the Agentic Enterprise 60
- Snowflake – Owning the Data Layer in the Age of Agents 65
- ServiceNow – Enterprise System of Action in an Agentic World 67
- Databricks – From Data Engineering to Agent Orchestration 70
- CrowdStrike – The Real-Time Control Layer for Autonomous Threat Defense 73
- Cloudflare – The Global Edge Control Plane Where Security, Connectivity, & AI Converge 76
- Palo Alto Networks – The Unified Security Platform Governing AI-Native Enterprises 80
- Zscaler – The Secure Access Platform Enabling Zero-Trust at Enterprise Scale 82
- Alphabet – Defending the Search Kingdom 85
- Amazon's AWS – The Infrastructure and Utility Play 90
- Meta – From Social Media to Personal Superintelligence 95
- Palantir – The Decision Layer Connecting Agentic AI to Institutional Knowledge 100
- Accenture – Rewiring the Enterprise for the Agentic Era 104

Executive Summary

Executive Summary – Bridging the Gap in the Agentic Enterprise; Highlighting Actionable Picks & Shovels Ideas

This report serves as a strategic roadmap for navigating the agentic stack, framing the investment opportunity around agentic AI, and detailing the structural transition from assistive GenAI copilots to autonomous digital labor. Along with defining agents and related architecture, comparing and contrasting with other forms of automation, highlighting early areas of meaningful agent adoption, LLM and agentic implications across software segments, we also address areas that will be key to reducing the trust gap that currently limits enterprise agent deployment. This includes defining the requirements for observability, governance, and audit-grade evidence in autonomous execution. By profiling the "Picks and Shovels" winners across foundational models, data layers, and security control points, the report illustrates how incumbents are leveraging their distribution and context moats, deterministic data and workflow history, to transform from simple software vendors into the central operating systems of the agent-enabled enterprise.

Strategic Company Profiles: Agentic Picks and Shovels Leaders

The following table summarizes the key public entities and core private foundation players identified as the primary beneficiaries and infrastructure providers for the emerging agentic economy, in our opinion. Please keep in mind that the flow of each public company in the table, and in the corresponding profile list throughout this report, goes in order according to Truist Securities TMT research sector: Application Software and Agentic AI, Infrastructure Software, Security Software, Internet and Digital Media, and Artificial Intelligence, Digital Platforms, & IT Services.

Private Companies

Company	Ticker	Rating	Analyst	Market Cap/Valuation ($Ms)	Agentic Profile Summary
OpenAI	Private	NR	N/A	$852,000	• **Execution at Scale**: Building agentic workflows through the "Frontier" platform for enterprise-wide task management. • **Developer Ecosystem**: Strategic focus on bridging experimental agents into production-grade autonomous systems. • **M&A**: Recent deals (OpenClaw-related talent acquisition, Astral) target improved developer tooling and agent governance.
Anthropic	Private	NR	N/A	$380,000	• **Reliability Focus**: Claude Enterprise Framework emphasizes "Constitutional AI" to prioritize safety in regulated sectors. • **Technical Execution**: Introduced "Claude Code" for technical tasks and "Cowork" for general productivity agents. • **Financial Scale**: Valued at $380B as of Feb 2026, with a $30B annualized revenue run rate.
Databricks	Private	NR	N/A	$134,000	• **Agent Bricks:** Multi-AI, governed platform for unifying model access, execution, and data context for enterprise agents. • **Unity AI Gateway**: Centralizes governance and enforces end-to-end identity & permissions across workflows. • **Data Intelligence Moat**: Leverages Unity Catalog metadata to achieve 70% higher accuracy than standard RAG architectures.

Public Companies

Microsoft	MSFT	Buy	Tillman	$3,153,071	• **Evolution to Control Panel**: Embedded Copilots to an "Enterprise Control Panel" capable of multi-step autonomous tasks. • **Azure Agent Units (AAUs)**: Monetization is transitioning from seat-based SaaS to "Digital Labor" models. • **Orchestration Layer**: Owning the stack via Azure AI Foundry, supporting multi-model architectures across 1M+ organizations.
Salesforce	CRM	Buy	Tillman	$145,745	• **Agentforce**: Positioning as "Agentic Enterprise Operating System" to address more complex, high value autonomous agents. • **Consumption Monetization**: Utilizing "Agentic Work Units" to charge for digital labor rather than traditional licenses. • **Market Traction**: Successfully closed over 29,000 Agentforce deals since its late 2024 debut.
Snowflake	SNOW	Buy	Jump	$48,509	• **Source of Truth**: Serves as the governed data layer required for agents to reason over enterprise-specific facts. • **Frictionless Deployment**:"Snowflake Intelligence" for natural-language queries and "Cortex Code" for orchestration. • **Rapid Adoption**: Currently generating a $100M revenue run-rate for its AI portfolio with 9,100 active accounts.
ServiceNow	NOW	Buy	Jump	$92,993	• **Execution & Enforcement**: "AI Control Tower," converting generative answers into resolving work w/ audit-grade evidence. • **Context Moat**: Leveraging its deterministic workflow history to drive autonomous IT and digital operations. • **Adoption Funnel**: "Now Assist" serves as the primary gateway for its massive installed base to adopt agentic features.
CrowdStrike	CRWD	Buy	Siddiqui	$113,652	• **Agentic Security**: "Charlotte AI AgentWorks" allows no-code building of mission-ready security agents. • **Autonomous SOC**: Unifying telemetry via "Enterprise Graph" to power cybersecurity operations at machine speed. • **Role:** AI-native runtime protection for endpoints and workloads
Cloudflare	NET	Buy	Siddiqui	$72,885	• **Edge Execution**: Positioning as the default execution environment for AI agents via the "Connectivity Cloud." • **MCP Support**: Launched the industry's first remote Model Context Protocol (MCP) server to support edge execution. • **Agentic Internet Security**: AI Security Suite discovers "shadow AI" and secures machine-to-machine communications globally.
Palo Alto Networks	PANW	Buy	Siddiqui	$145,689	• **Cortex AgentiX**: Transforming its platform into a backbone for agentic security that plans and executes defense tasks. • **Efficiency Gains**: Early AgentiX deployments show a 98% reduction in Mean Time to Remediation (MTTR). • **Platformization**: Consolidating code, cloud, and identity to counter threats operating far beyond human capacity.
Zscaler	ZS	Buy	Siddiqui	$21,787	• **Zero Trust Governance:** Providing identity-centric connectivity for enterprise AI assets and machine-to-machine execution. • **AI Protect Suite:** Delivers comprehensive inventory and dependency maps for all models and agents in an enterprise. • **Threat Remediation:** Processes 500T security signals daily to train models for fully automated threat detection.
Alphabet (Google)	GOOGL	Buy	Squali	$4,166,207	• **World Models**: From "Word Models" to "World Models" w/ Gemini 3, Google Agentspace, & Chrome Auto-Browse. • **Standardization**: Universal Commerce Protocol (UCP) standardizes how agents interact with millions of merchants. • **Infrastructure Scale**: Investing over $175B in 2026 CapEx to maintain a vertically integrated stack from TPUs to Gemini.
Amazon (AWS)	AMZN	Buy	Squali	$2,839,028	• **Utility Infrastructure**: The full stack, from Trainium hardware to Bedrock models, for enterprises to scale custom agents. • **Strong Momentum**: Amazon Bedrock saw 60% Q/Q spend growth, reaching a multi-billion dollar annualized run rate. • **Model Diversity**: Strategy focuses on providing "choice" across foundational models and agentic governance tools.
Meta	META	Buy	Squali	$1,713,515	• **Open-Source Dominance**: Betting on the Llama family to set the global industry standard for agentic execution. • **Hardware Interface**: Using smart glasses (Ray-Ban Meta) as the primary interface for "Personal Superintelligence" agents. • **Compute Scale**: Building gigawatt-scale clusters to subsidize its goal of becoming the agentic standard.
Palantir	PLTR	Buy	Ramnani	$342,225	• **AI Operating System**: Verticalized AIP platform provides the runtime and orchestration for deterministic agent execution. • **P&L Outcomes**: Focuses on translating probabilistic AI reasoning into measurable P&L and operational outcomes. • **Strategic Partnerships**: Deep integration with US Navy and Stellantis for large-scale data reinvention and agentic autonomy.
Accenture	ACN	Buy	Ramnani	$109,502	• **Services Integration**: Leading the "Services Picks and Shovels" wave by integrating & securing agents across the Fortune 500. • **Agent Library**: Developed a library of 3,000+ reusable agents to bridge the gap from models to industry outcomes. • **Model Agnostic**: Partnering across OpenAI, Anthropic, and hyperscalers to drive "Reinvention Services."

Source: FactSet, Company reports, Truist Securities Research

Note: Market Caps are as of 4/27/2026

Note: OpenAI, Anthropic, and Databricks valuations are last reported valuations issued by the companies themselves.

What is an Agent?

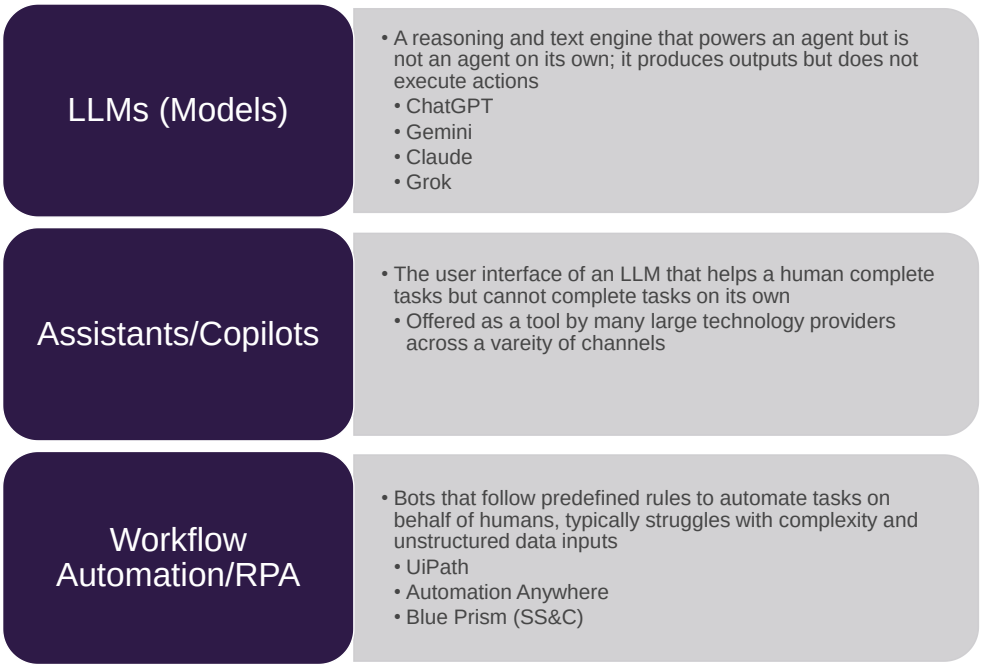
In its simplest form, an agent in the context of artificial intelligence is a system that can independently plan and execute tasks, from simple ones to more complex multi-step workflows that may invoke various data sources as well as other agents. The key difference between an agentic and generative AI is proactivity. Agents act on their own accord to handle workflows and interact across systems while generative AI requires commands given through prompts to accomplish a task involving direct human interaction. The core competencies that separate agents from LLMs, copilots, process automation tools, etc. are autonomy, real-time reasoning, and the ability to leverage tools across a variety of systems. Because agents can transcend systems, data sets, and organizational structures, this makes them valuable for certain use cases including: 1) open-ended tasks that involve several key decisions, 2) complex situations that may not have a finite number of steps to completion at the outset, and 3) any situation where a system working through tasks will need to be flexible and have the ability to adapt to a changing environment.

Figure 1. What Defines an Agentic System



Source: Amazon Web Services

Figure 2. What Isn't an Agent?



Source: Truist Securities Research

Agent Architecture

At a high level, an AI agent is better classified as a system, not a model, in our opinion. Agentic architectures are composed of several modular components that work together to enable goal-focused, autonomous behavior versus simple prompt responses.

Foundational Model

At the core of an AI agent sits a foundation model or large language model (LLM), ChatGPT, Gemini, Claude, etc. which functions as the system's reasoning engine. The model interprets natural language inputs, reasons through instructions, and determines what actions to take next. Importantly, the LLM does not execute work, but rather it decides when to retrieve information, use tools, or update its memory. Further, the separation between reasoning and execution is what allows agents to move beyond static responses and into dynamic workflows.

Technically, autonomous agents could exist absent of commercial or open source LLMs. One could see certain system of record or vertical SaaS companies leverage what effectively could be considered small language models or their own purely proprietary reasoning engines; however, we assume such examples would involve more limited scope or task execution. LLMs are especially relevant in any type of enterprise agent use involving some level of sophistication as well as workflows or autonomous work tasks that incorporate meaningful amounts of

unstructured data. That is where models and powerful reasoning become important in driving more effective and deterministic agent outcomes, in our opinion.

Planning Layer

The planning layer enables the agent to translate objectives into executable steps. Rather than reacting to a single prompt, the agent can evaluate the components and determine the ordering across tasks. The planning function can range from simple prompt-based task breakdowns to more structured planning sequences. From what we have observed, this is the component that allows agents to operate over longer time horizons and handle workflows traditionally performed by humans.

Memory & Context

Agents require memory to operate efficiently as time passes. Agent architecture typically includes both short-term memory (e.g., session context or recent actions) and long-term memory (e.g., specific customer data, prior results, or general knowledge). This memory layer improves personalization, continuity, and accuracy. Importantly, we believe that proprietary data stored in memory increasingly represents a key source of differentiation when it comes to agent usage across enterprise workflows.

Tools & APIs

One of the key differentiators of agents is their ability to act across external systems. The agents can do this through tools and API integrations where agents can retrieve data, update records, send communications, or trigger workflows across enterprise software environments. The agent will leverage an LLM to determine when a tool is needed and then move to the execution or orchestration layer to handle how the action is performed. Further, we believe this capability of being able to connect reasoning to system changes in real-time is what differentiates agents from assistants or copilots.

Learning and Evaluation Layer

Agent architectures also incorporate self-evaluation and improvement to help strengthen performance over time. Agents can assess the success of an action vs. its intended result, incorporate feedback from either a human or automated entity, and make changes to its future behavior accordingly. Agents can also sometimes leverage reinforced learning to optimize decision making processes over time, typically in environments where explicit training data is not available. We believe the learning layer is critical to agentic systems as it provides agents with improvements to become more reliable over time and creates trust when being used within large enterprise systems. This layer also relates to reality of human-in-the-loop dynamics we could envision in situations such as high value agent work being performed, determinism outcomes are critically important, regulated industry settings where additional complexity, trust or privacy factors are heightened, etc.

Proliferation of Agents = Need for Interoperability, Management and Observability

As more agents become prevalent in the workforce, there is an inherent need for them to be able to communicate, understand, and coordinate with each other to achieve complex goals. Interoperability's foundation is built on standardized communication languages to allow agents to exchange messages, delegate tasks, and understand capabilities. Agents need to dynamically be able to identify skills and tools that best fit a given situation. As agent interoperability begins to break down silos across enterprise applications and scale effectively, security and compliance will remain key to unlocking fully autonomous workflows.

Management and observability will also be key to agent success in production environments. Similar to everyday employees, digital workers or agents need to be managed and assessed on their performance. As these agents get rolled out, they must have clear objectives, task assignments, and performance scoring. This aligns with an ongoing narrative around "agentic personas." Companies like Anthropic (Private) are rolling out agents with specific expertise to automate tasks and workflows in certain areas like HR and FP&A. As these agents become more autonomous, it is important to audit their decision making and logic framework. Agents have the tendency to "hallucinate," meaning that they deduce their own reasoning/make up information to meet certain criteria. A key focus with agents going forward is to have visibility into their reasoning and have an audit trail into how they are formulating their decisioning.

Agent interoperability will be crucial for different AI agents, built on various platforms, to seamlessly communicate, understand each other, share data, and coordinate actions to achieve complex goals, forming an open ecosystem for scalable automation without being locked into single vendors or frameworks. It enables features like HR chatbots collaborating with finance systems for payroll, or IT agents automatically pulling data, leading to more efficient, autonomous, and interconnected AI workflows. Key enablers include standardized protocols (like MCP, A2A), discovery mechanisms, and trust frameworks for secure, context-aware collaboration across different systems and cloud environments.

Protocols: MCP and A2A; Important for Enterprise Agent Success

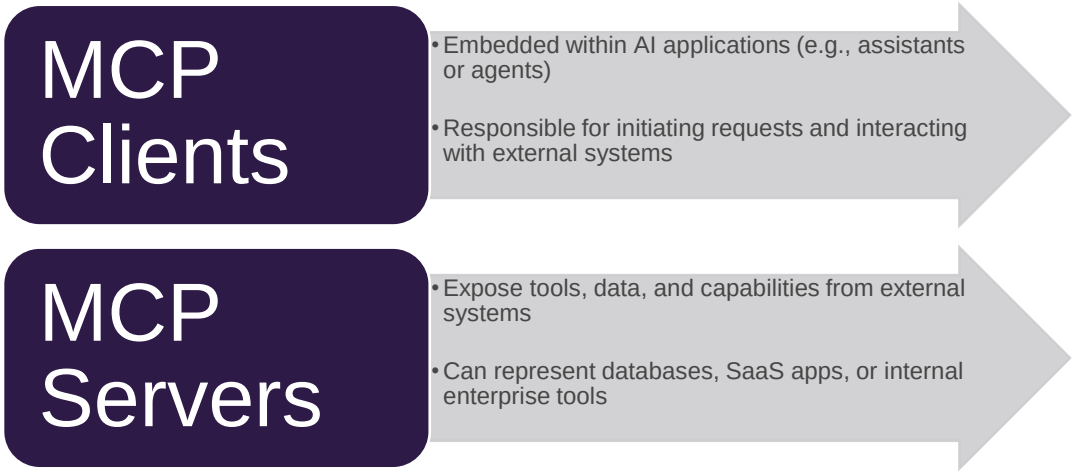
Model Context Protocol (MCP)

Anthropic introduced MCP in November 2025 as an open standard that enables AI systems to connect directly with external data sources, tools, and services through a consistent interface. According to Anthropic, MCP provides a standardized way for models to access and interact with real-world context, allowing them to move beyond static knowledge.

MCP enables bi-directional communication between AI models and external environments:

- Models can retrieve data from connected systems
- Models can take actions or write back to those systems
- Interactions occur through a structured, secure protocol layer

Figure 3. How MCP Works: A Client-Server Architectural Framework



Source: Anthropic company website

Enterprise Implications and Strategic Direction: Infrastructure for Agentic AI

Anthropic positions MCP as the foundation for production across AI systems because the protocol reduces integration overhead, improves accuracy with real-time, relevant data, and expands use cases and automations across enterprise workflows. Anthropic believes that MCP can be framed as a core building block for the next generation of AI systems because it enables agents to interact with real-world systems, supports multi-agent workflows through shared context and tools, and establishes a foundation for interoperable AI ecosystems.

Agent2Agent Protocol (A2A)

A2A is an open standard protocol introduced by Google (GOOGL, Buy, Squali) in April 2025 that is designed to allow AI agents to communicate, exchange information, and coordinate actions across different systems and platforms. The goal of A2A, according to Google, is to move beyond isolated agents towards interconnected agent ecosystems, where independent systems can collaborate regardless of how or where they were built.

A2A introduces a standard communication layer that allows agents to:

- Discover one another
- Exchange structured messages
- Coordinate multi-step tasks

Importantly, A2A does not require agents to share the same infrastructure or internal logic, focusing instead on how they interact. This allows for flexibility and modularity as agents communicate separately from how they are built.

Figure 4. How Does Agent-to-Agent Collaboration Work?



Source: Google Blog

Enterprise Implications

A2A is designed to unlock:

- Cross-system automation: Agents coordinating across departments or software stacks
- Vendor interoperability: Mixing agents from different providers
- Scalable workflows: Multi-agent systems that can expand without exponential integration complexity

Why Are MCP and A2A Complementary?

While it might seem that Google and Anthropic have created two operating systems for agentic AI that will compete to be the standard across enterprise workflows, it is important to note that these standards will work alongside each other because they address different layers of the stack. According to Google, A2A enables communication between agents while Anthropic

highlights that MCP connects agents to tools, APIs, and data sources. The important takeaway on the differences of these protocols is that A2A promotes agent coordination and MCP facilitates agent execution, both equally important capabilities to address enterprise agentic workflows, in our opinion.

Also, going forward we will be watching to see if other industry standards or protocol efforts come to the forefront. Will such efforts be a positive enabler of agentic traction or lead to fragmentation and stymie agent traction? We could envision efforts associated with more specific industry use cases, such as we have seen in e-commerce. For example, there has been development of Agentic Commerce Protocol (ACP), an open-source standard designed to enable agents to interact with merchant systems to search, check out and purchase products on behalf of merchants and customers. It was co-developed by OpenAI and Stripe.

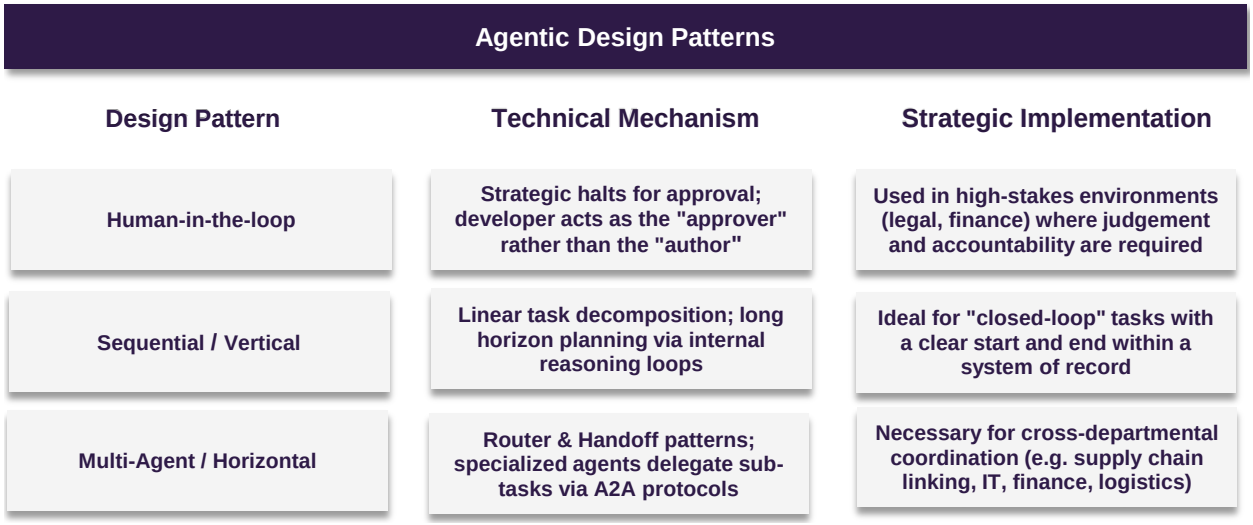
Similarly, in January of this year Shopify (SHOP, Buy, Tillman) and Google announced collaboration to develop Universal Commerce Protocol (UCP), another open-source standard to accelerate agentic commerce. Through this technology collaboration, agents, such as those powered by Google Gemini, are meant to securely connect with merchant checkouts to smoothly drive purchases in a repeatable fashion through LLMs and conversational interfaces.

Going forward, we believe it will be collectively important for the broader AI and agentic ecosystem to deliver on integration and operability needs associated with enterprise agents. We will look for advances directly from LLM companies, others in the AI ecosystem with vertically integrated agentic infrastructure stacks, software ISVs delivering on their agentic efforts as well as what is likely to be a vibrant set of integration specialists with best-of-breed MCP gateways and other key enabling technology. It is worth noting per a Gartner 2026 Buyer Insight survey that “45% of martech leaders with AI agents in pilots or production say existing vendor-offered AI agent capabilities do not meet company expectations of promised business performance.” [2]

Agentic Application Design Patterns and Orchestration; How Can Agents Actually Be Useful?

We believe that as organizations move past single-prompt chatbots, the application layer for agents (or one could call it the intelligence layer) will consolidate into 3 primary architectural patterns: 1) human-in-the-loop, 2) sequential/vertical, and 3) multi-agent/horizontal. By using this framework, agents will be classified by personas that dictate how they actually interact with the software.

Figure 5. Agentic Design Patterns



Source: Truist Securities Research

Agentic Personas: Where are we currently seeing agents being deployed and example use cases

- **Software Development & Coding:** In our opinion, software development and coding are the most visible enterprise applications for agents in their current, early form. Agents handle a majority of code generation while humans stay in the loop to provide tweaks and final verification within a controlled environment. One needs to look no further than the historic ARR ramp being seen by Anthropic with its Claude Code and enterprise traction, significant ramp by Codex at OpenAI, coding specialists such as Cursor (Private) also ramping revenue at a tremendous growth trajectory for validation on substantial agent use case associated with software development and coding, in our opinion.
- **Finance (FP&A):** Typically, tasks utilizing single-agent sequential patterns within finite environment. Examples within corporate finance include an agent automating month end close data reconciliation or a data analyst agent that can identify risks and make recommendations.
- **HR & Legal:** Deploying human-in-the-loop patterns for drafting performance reviews, initial contract markups, or summarizing notes or briefs. Importantly, these are typically sensitive workflows that require a human to oversee the agents for accuracy and to prevent regulatory or compliance violations.
- **IT & Customer Ops:** Leveraging multi-agent orchestration for triaging tickets and then automatically pulling data from across disparate SaaS environments using MCP to resolve issues and close tickets.

The Shift to the Generative UI: The Final Mile of Delivery for Agents

We believe the way that agentic tasks will be delivered comes from what is called “generative UI.” Traditional UI interfaces are static layouts that are typically manually configured to include menus and workflows that will be accessed by humans. It is our belief that in an agentic world, dynamic, real-time interfaces will be utilized, and they are personally tailored based on user intent, context, and behavioral data.

At its core, generative UI bridges the traditional separation between back-end logic and front-end design. Historically, developers built fixed interfaces on top of APIs, requiring users to navigate menus, dashboards, and workflows to complete tasks. In a generative model, AI agents sit on top of these systems and construct the interface on demand, effectively turning the interface itself into a byproduct of the task being performed. This shift is best understood as the transition from workflow-driven software to intent-driven software. Instead of adapting to predefined screens, users express what they want in natural language, and the system dynamically generates the appropriate controls, layouts, and actions needed to fulfill that request. The UI is no longer something users learn, it becomes something that adapts to them in real time.

Technically, this is enabled by large language models and agent frameworks that can interpret user intent, orchestrate backend actions, and render interface components in real-time. Rather than hardcoding every possible state or edge case, developers expose capabilities, and the agent determines how those capabilities should be presented and executed in context. The implication is a redefinition of the front-end itself. In a generative UI environment, the front-end is no longer a static application layer but rather an adaptive surface generated by an AI agent. Interfaces become fluid. Personalized and change based on the user, the task, and the moment.

For software companies, we believe this introduces both opportunity and disruption. Products that rely on rigid navigation, fixed workflows, or manual configuration risk being abstracted away by agent-driven layers that can orchestrate actions across systems. Conversely, we believe platforms that own structured data, APIs, and execution infrastructure are better positioned to remain relevant, as agents depend on these systems to take action. We believe that as interfaces become generated rather than designed, the value moves away from screens and toward the systems that understand intent and execute outcomes.

Importantly, we note that generative UI represents the critical “last mile” for agentic systems. While much of the focus has been on how agents’ reason, plan, and take action across tools, their output ultimately needs to be translated into something a human can understand, validate, and trust. Generative interfaces serve as that bridge for dynamically rendering decisions, surfacing context, and enabling intervention when needed. In this sense, the UI is no longer just a presentation layer, but a core component of the agent loop where human oversight and machine execution converge. We believe that as agents take on more complex, multi-step workflows, the quality of this interface layer becomes increasingly important, determining not

just usability, but adoption and trust at scale. From our vantage point of covering application software companies, SaaS incumbents as well as AI-native startups will need to perfect agents as well as this accompanying generative UI proposition; success or failure here could have material consequences on long-term growth and terminal value, in our opinion.

Agentic Implications for Application Software Incumbents, LLM Threats, Model Changes, Etc.

1. The "SaaS Industrial Complex" at Risk?

The shift toward agentic AI is uniquely disruptive because it confronts an unprecedentedly mature software footprint. Unlike previous technology cycles, where software penetration was nascent, today's incumbents manage a vast "SaaS industrial complex" of systems of record (SoR) and point-solution workloads following a decade of digital transformation. This saturation creates a massive surface area for disruption. We believe much of this existing stack is at risk of replacement by autonomous agents, whether deployed by LLM providers, AI-native upstarts, or through "vibe coding" (intent-based, natural-language development) that allows users to bypass traditional SaaS UI/UX entirely.

2. From Process Automation to Systems of Agency

A primary factor driving this consequential shift is that generative AI and agentic workflows are evolving beyond simple automation toward sophisticated "Systems of Intelligence." While legacy SaaS focused on data entry and process orchestration, agentic solutions prioritize autonomous decision-making and outcome execution. This evolution tilts the balance of value creation toward AI-native platforms, fundamentally eroding the economic leverage and pricing power of SaaS incumbents. As customers prioritize agent-led execution, the traditional "System of Record" becomes a commodity, creating significant friction for incumbents attempting to maintain the status quo.

3. Compressed Innovation Cycle and Coding Agents

The monetization of coding agents is already accelerating this disruption by drastically compressing software development lifecycles from design and testing to debugging. This efficiency catalyzed the rise of Anthropic, which has set records as the fastest-growing software company in history. We expect these coding agents to unleash a proliferation of AI-native startups that iterate fast enough to move beyond mere "plug-in" layers. These new players are increasingly vertically integrating their platforms to serve as both the System of Record and the "System of Context," threatening to hollow out traditional horizontal and vertical software categories through rapid, AI-driven iteration.

4. The $60 Trillion Paradigm: Monetizing Digital Labor

The ultimate differentiator of the agentic cycle is its impact on the global labor market. With aggregate global labor spend estimated at over $60 trillion per year, the emergence of "digital

workers" (autonomous agents) represents the largest TAM expansion in the history of computing. Agentic AI is not merely optimizing software budgets; it is augmenting, optimizing, or displacing human labor at scale. We believe this transition from human payroll to digital labor represents a multi-trillion-dollar monetization opportunity for the agentic ecosystem, fundamentally altering the unit economics of the global economy.

Our Stance on Software: The Glass Remains Half Full

Alongside the "picks and shovels" AI opportunities detailed later, we remain constructive yet vigilant across our broader horizontal and vertical SaaS coverage. We believe foundational Systems of Record (SoR) that have aggregated decades of "source of truth" data and behavioral context are fundamentally resilient; consequently, we do not anticipate material terminal value risk in the immediate term. However, the rapid evolution of the agentic infrastructure ecosystem suggests a more bifurcated mid-term outlook.

Over a 3–5 year horizon, we expect the strategic moats of certain incumbents to erode as the threat of irrelevancy intensifies for those slow to adapt. Specifically, non-SoR point solutions that provide limited, discrete value face an elevated risk of disintermediation. Their displacement will likely be accelerated by AI-native challengers, custom workflows enabled by vibe coding, and specialized agent orchestration layers that effectively bypass traditional application UIs in favor of autonomous execution.

Call to Action for SoRs to Maintain Relevancy Where Action is Going. Even in instances where SoR and platform businesses have less terminal value risk, we believe there is an urgent need for them to modernize their core tech stacks, rebuild user experiences with generative UIs, deliver credible highly deterministic agents with positive business outcomes that fortify their SoR value proposition, etc. From our perspective and hearing directly from CIOs and technology practitioners where you might see more of the high value “action” are agents that drive decision making and intelligence (e.g., change the business type agentic use cases). Sure, there is significant ongoing interest and expected growing adoption of automation and customer-facing agents like the FAQ chatbot variety, or those centered on go-to-market, sales and marketing outcomes or that run back-office activities more efficiently; however, we believe our horizontal and vertical SaaS companies need to accelerate agentic innovations that are more of the change the business versus operate the business type and truly drive increased intelligence.

Although not an exhaustive list of names of AI winners/survivors, coverage names we highlight with potentially high-value agentic setups going forward (separate from picks and shovels names we highlight later) include Intapp (INTA, Buy), Klaviyo (KVYO, Buy), Manhattan Associates (MANH, Buy), nCino (NCNO, Hold), Q2 Holdings (QTWO, Buy) and Zeta Global (ZETA, Buy).

We Also Favor End-to-End Operating System Platforms in Agentic World. We believe that a few of the companies that we cover that have a unique value proposition of automating and optimizing broad sets of long-running operating workflows are especially well positioned to not only survive but thrive with proprietary agent innovation. We are specifically talking about ServiceTitan (TTAN, Buy) and Shopify (SHOP, Buy). LLMs themselves or any agentic startups

for that matter are likely to be disadvantaged versus these operating system companies from a sheer highly contextual data or domain knowledge standpoint. In addition, even if agents are successful in terms of automating or delivering intelligence in certain functional areas or business workflow, their output or outcomes could still be suboptimal as likely would not be optimizing across broader, longer-running business workflows.

For example, in the case of ServiceTitan, its Core, Pro and Agentic solutions are not meant to optimize solely for specific tasks or workflows but deliver maximized outcomes that cut across an entire business and not just create benefits in isolated business functions. Its operating system platform is meant to create the best overall business outcomes (e.g., rev, costs, profits) that connect across sales and marketing funnel, customer service/contact center, scheduling, quoting, selling, field service, project management, accounts payable, payments and finance. This creates a unique agentic opportunity for a vertical SaaS incumbent with true operating system dynamics such as ServiceTitan, in our opinion. The same can be said for Shopify, in our opinion, as its operating system manages the complexity and performance requirements associated with global commerce at scale combined with deeply embedded checkout/payment experiences. We see these two operating system companies as AI and agentic winners.

Beware of LLMs Crossing the SaaS Rubicon. The unprecedented growth and capital accumulation of LLM providers have created a fundamental pivot point for the enterprise software ecosystem. These entities have achieved the fastest historical growth at scale in software history, amassing substantial funding "war chests" that grant them immense strategic flexibility as enterprises design their future agentic architectures. While viral social media demonstrations and vibe coding videos suggest a looming overhaul of global economies, and we acknowledge that list of commercialized use cases is both broad and potent, the actual commercial traction for frontier models has remained disciplined within well-defined functional lanes focused on generative AI and agentic task execution.

Our baseline assumption is that frontier model leaders and software incumbents will continue to evolve in a state of "coopetition" within the enterprise market. In this framework, LLMs serve as the essential engines for agentic workloads that span application software, data infrastructure, and cybersecurity. While some software firms may formalize exclusive partnerships with specific models or hyperscalers, we expect a broader shift toward a multi-model orchestration approach. This strategy allows SaaS and System of Record (SoR) companies to optimize agent performance across diverse workloads while mitigating the long-term risk of being commoditized at the foundational model layer. By maintaining model neutrality, incumbents are attempting to preserve their strategic positioning as the orchestrators of enterprise intelligence.

The narrative moves from perceived risk to a material threat to terminal value if providers like OpenAI and Anthropic choose to "cross the Rubicon" and enter established markets as direct competitors. We see potential for increased hand wringing should these players decide to leverage their technical moats to build fully developed platforms with proprietary business logic, advanced workflows, and full-stack capabilities, in attempt to disintermediate established horizontal and vertical SaaS incumbents. Such a move toward vertical integration would likely trigger significant market volatility and exert severe pressure on already depressed software

valuation multiples. Furthermore, these encroachment fears are not limited to the application layer; they extend into adjacent categories like cybersecurity and data orchestration, where a full-stack "System of Intelligence" could further challenge the outlook for incumbent players.

Appreciating strategy can shift quickly and innovation cycle time are much compressed in this AI and agentic age, we are not assuming the major LLM and AI disrupters become tomorrow's SoRs in areas like ERP, CRM, HCM, commerce, Office of CFO, etc. with packaged commercial offerings (or at least any time soon). Even if these markets get reimagined by more modern agentic architectures, that would likely benefit a different set of startups and SaaS incumbents that reinvent themselves versus LLMs vendors specifically. LLM and agentic capabilities could certainly be cobbled together for custom-built SoRs or vibe coding could take greater hold in enterprise SoR and agentic settings; however, we believe existing SoR incumbency benefits around data and context, trust & governance as well as practical challenges associated with maintaining and delivering constant updates and SLA adherence could limit near-term material movement by LLMs or vibe coding into core enterprise application markets or closely coupled agent adoption. Potentially influencing less direct software disintermediation could be the practical dynamics of the TAMs these substantially capitalized LLMs strive to address are likely substantially larger and broader than horizontal app or vertical SaaS areas of the market.

While we keep a vigilant eye on direct application software disintermediation by LLMs and other agentic disrupters, we are equally focused on those in our coverage that could have breakout performance with agentic products while also discerning the financial model impact of what will likely be an ongoing shift away from seat-based pricing and packaging to value, consumption or outcome-based monetization models. In our opinion, investors have gotten better trained over time on the mechanics of financial model shifts as we saw with SaaS and certain types of consumption-based pricing developments. Rapidly scaling agent solution demand would be visible through quantifiable and qualitative financial metrics/KPIs. Any near-term impact from subscription and seats to more AI-centered pricing and packaging would represent a noise and timing thing versus relevancy and terminal value thing, in our opinion. We believe investors with a patient long-term lens could stomach financial model transitions, especially considering the current valuation risk/reward landscape, and the assumption that ultimately growth could rerate higher post such model transition.

Sizing the Agentic Market Opportunity

While we have seen a variety of validating data points on the potential for AI agents, including billions of dollars in revenue and quickly growing coding agentic revenue and record-setting adoption of open-source agent innovations like OpenClaw, the full scope and scale of the enterprise agent opportunity is yet to be fully defined. That said, we believe the market for autonomous agents acting as digital workers potentially represents trillions of dollars of revenue potential.

Think Much Bigger on Funding Sources for Agent Spending. We find recent investor day commentary by one of our coverage companies, Intapp and its CEO John Hall, informative in terms of potential sources of investment dollars for AI and agentic initiatives. With Intapp's total ARR business estimated to approach $650 million in FY27, there are certainly much larger software companies one could rely on for agentic budget perspective; however, we believe Mr. Hall offers an expert perspective on the enterprise agent market opportunity. For more context, Intapp is a vertical SaaS company that focuses on some of the largest industries in the world from a revenue and economic output perspective. This includes global law firms and other professional services organizations in accountancy and consulting, global investment banks, private equity and other alternative investment funds, etc. Collectively, they spend substantial amounts on technology, and these industries are increasingly at the forefront of AI and agent experimentation and use cases (e.g., legal tech in particular).

Paraphrasing Mr. Hall: He sees three sources of funding to meet the demands of AI. There are pre-existing traditional IT budgets, and in that realm, CIOs have become highly AI oriented. Another source of funding for AI budgets are net new budget dollars. Different firms may call it different things and examples could be simply an AI budget or an innovation budget. This newer budget source also includes an expanded set of decision makers beyond traditional technologists and CIOs, such as a head of innovation for example. Lastly, there is a third and much bigger potential source of budget dollars available to the agentic opportunity. These set of investment dollars can potentially be quite large and funded from what would have been spent on filling a new position for specific skill set or conversely funded through reduction in human labor. As companies look at agentic investment opportunities, they are starting to get adept at measuring the potential labor spend that would have been involved in lieu of agents in various functional areas. The full labor spend component and related cost avoidance can add up to upwards of 40% to 60% of total spend for certain lines of business, functional area or business workflow transformation project being pursued. Per Mr. Hall, this opens up much larger sources of investment dollars on planned agent deployments that can tap into operational, line of business or actual labor spend for funding. This reflects what we could see as exponentially larger funding sources for agents in comparison to reliance on traditional IT or even new innovation budgets.

Agentic AI in Relation to Labor Spend and Shift from Labor Arbitrage to Intelligence Arbitrage. Autonomous work, whether it be more simple or complex and high value in nature, is certainly relevant to a discussion of agent market sizing. We also believe comparisons to actual

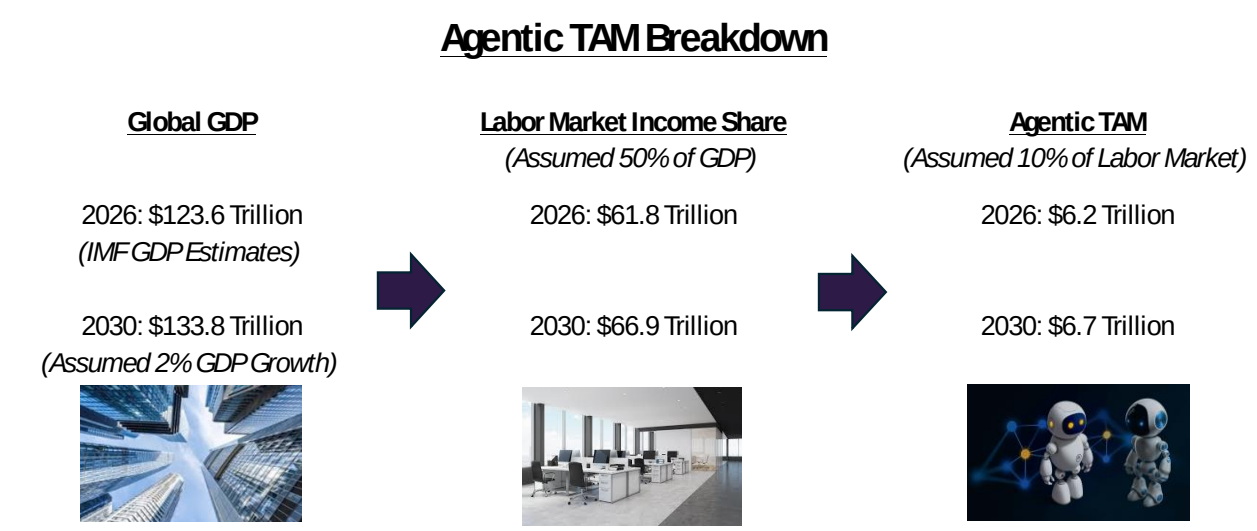
labor spend is also appropriate for agentic market sizing. Agentic AI directly relates to how labor spend, which is generally estimated to represent 50%+ of global GDP, can be arbitraged out or augmented to drive better business outcomes. As mentioned previously, we envision growing funding sources for agents to come from much larger budgets associated with labor and other operational budgets.

Using publicly available global GDP forecast, in this case an updated IMF projection in April, points to estimated total global GDP (nominal) of $123.6 trillion. Using a simple assumption of 50% of global economic activity relates to labor spend, this translates to global labor spend of $61.8 billion. Recognizing enterprise agent adoption and pilot/experiment to full production agents is still early, we believe it is worth extrapolating out GDP and related labor spend to 2030. We believe this better aligns with maturation of enterprise readiness, proven agentic workloads at scale as well as benefits accruing from a required multi-year cycle of AI data center and infrastructure investments. Extrapolating into 2030 as illustrated below with assumed 2% annual growth, 2030 total global GDP is calculated to be $133.8 trillion. Then applying the same 50% mix of GDP that is labor related and derived from a variety of industry sources, equates to $66.9 trillion in global labor spend in 2030. This represents our baseline on which we apply a percentage of labor that we refer to as either labor savings or labor optimization rate potentially expected from agentic automation.

Take rate is a bit clumsy as a term to describe the percentage of labor spend that either can be replaced, avoided or augmented for purposes of an agentic market sizing. But the concept is appropriate, in our opinion, given how large the potential is with enterprise agents and their ability to drive material revenue, cost reduction or profit improvement outcomes over time. While early anecdotes have been wide ranging in terms of positive business outcomes around agents, such as significant productivity increases in software development, documented cost savings related to automated case resolution in customer support or functional areas such as sales, marketing or back-office processing have reported significant cost savings and productivity gains well. We will substitute a take rate with a concept we call labor optimization rate of 10% and apply against 2030 estimated global labor spend to derive a top-down simplified agent market sizing of $6.7 trillion by 2030. We believe this 10% labor optimization rate could be quite conservative taking into documented and future expected cost savings and efficiency gains that we have often heard starting at least in the 20% - 30% range and not uncommonly quite higher.

We believe this type of top-down analysis on potential agentic market sizing is appropriate as it is likely to align well with how C-level executives and boards of directors will look at the business case and opportunities to transform the business via enterprise agents broadly deployed across an organization and broader ecosystems as well as providing the potential budget source for funding what could often represent large incremental greenfield technology investments as well as replacement of prior generation business workflow and automation solutions. Also, this top-down agentic market analysis has more controlled and logical inputs, in our opinion, than attempting some sort of bottoms up number of agents deployed globally, tokens or other forms of consumption monetization and/or outcome-based pricing that could be problematic given so many different types of outcome or value-based pricing that is likely to unfold in the market.

Figure 6. Agentic Market Opportunity



Source: IMF Website

Note: GDP growth, labor market income share, and agentic TAM percentages are Truist Securities Research Estimates

Agent Economics Correlating to Economic Activity, Labor and Successful Recouping Trillions in AI Infrastructure Investment is How Agentic Market Could Really Take Off. While the reported success from some of the fastest-growing LLM and AI native companies speaks to large and readily addressable consumer and enterprise markets for generative and agentic AI, we see ongoing large investment in energy/physical, compute, LLM training and inference investment as required for a visible path to adequate long-term payback and ROI. As mentioned at the outset, AI luminaries like NVIDIA's (NVDA, Buy, Stein) Jenson Huang have sized required AI infrastructure investment required through 2030 at $3 trillion - $4 trillion. McKinsey estimates total data center capex of $7 trillion and 220 gigawatts will be required to meet AI demand by 2030[1]. Also, some of the largest hyperscalers in the world are planning to invest over $660 billion in AI-related capex this year and that is expected to keep growing into ~$730 billion investment level the following year. Those types of upfront AI infrastructure investments must be absorbed somewhere, or the AI and agentic hype cycle and investment capital would likely come to a screeching halt. We believe there will likely be large multiple effect on revenue generated across a myriad of AI-related markets, including enterprise agents, agentic commerce, rebuilding and expansion of core task automation, operating system and knowledge worker productivity tools along with greenfield AI and agentic adoption across high-value industry use cases in healthcare, financial services, government, defense, etc. as well. We believe this multi-year data center and AI capex investment cycle also validates trillions of dollars in revenue potential from enterprise agents.

Figure 7. Big 5 Hyperscaler Capex Estimates

Hyperscalers Consensus Capex Estimates (In Billions)

Company	2026	2027	2028
Amazon	$193.0	$204.6	$202.3
Alphabet	$173.5	$191.2	$188.2
Meta	$122.1	$133.4	$134.7
Microsoft	$117.3	$135.1	$156.1
Oracle	$57.1	$68.1	$72.8
Capex Total	**$663.0**	**$732.4**	**$754.1**

Source: Factset

For cross reference, we have also included Gartner research forecast of worldwide AI spending across specific segments of the broadly defined AI market below [5]. It is not directly comparable to our agentic top-down market size and doesn't explicitly define agentic TAM. It has breakouts for AI software, AI models and several other relevant categories but overall, not viewed as the best proxy for specific agentic market size, in our opinion. That said, it is somewhat useful directionally given the total spending that accumulates into the $3 trillion+ level by 2027. We believe traditional IT market research and related AI spending estimates could be problematic and not fully capture the extent of agentic spending as a sizable portion of future enterprise agent investment could open up additional intelligence-related TAMs and greenfield market penetration.

Figure 8. AI Spending Estimates

Table 1: Worldwide AI Spending by Market, 2025-2027 (Millions of U.S. Dollars)

Market	2025	2026	2027
AI Services	439,438	588,645	761,042
AI Cybersecurity	25,920	51,347	85,997
AI Software	283,136	452,458	636,146
AI Models	14,416	26,380	43,449
AI Platforms for Data Science and Machine Learning	21,868	31,120	44,482
AI Application Development Platforms	6,587	8,416	10,922
AI Data	827	3,119	6,440
AI Infrastructure	964,960	1,366,360	1,748,212
Total AI Spending	**1,757,152**	**2,527,845**	**3,336,690**

Source: Gartner (January 2026)

We Lean More Towards AI Accelerationist vs. AI Doomer. While talking of large budget dollars being freed up from labor and other operational budgets can naturally be alarming to most people, we are more constructive on positive long-term effects of AI and autonomous work agents. No question there are likely to be plenty of examples of labor-related cost reduction and/or scaling businesses more readily without additional headcount. That said, we see some potentially large incremental value creation and economic amplification in an agentic focused world. Four reasons we are optimistic about agentic AI:

1) **Unprecedent Speed of Innovation Ignites New Business Cycle:** We believe the pace of software development and innovation could unleash other large entrepreneurial and business formation cycle, similar to cycles in the past such as small business startup creation post the Internet and more recently amid new digital-first entrepreneurship and new business formation seen during COVID and reflected very impressively in Shopify's long-term durable growth amid an otherwise sluggish brick-and-mortar retail landscape.

2) **Enterprise Agent Opportunity Drives Growth Resurgence**: While slow-to-move SaaS incumbents and SoR businesses could be impacted by agentic disruption, we believe there could be resurgent revenue growth in software via SaaS incumbents that aggressively transition to agentic AI along with LLM and other AI-native businesses that are well positioned to represent the next large or mega cap investment opportunities and meaningful TAM expansion brought on by the AI eats software shift.

3) **Expert Use of Agents for Decision Making and Intelligence Transforms Businesses and Competitor Separation**. AI and agents are not just the property of technology businesses. We believe applied AI and complex agent use cases stand to transform businesses across practically any industry, whether it be a 3PL optimizing its freight business and logistics network, a bank significantly streamlining credit deliver while taking market share in deposits and lending, an insurance provider significantly improving underwriting success while structurally improving fixed cost structure or a local HVAC or plumber being able to better optimize growth and market consolidation activities via agents deeply embedded across entire operational processes of a trades business. Agentic AI could represent the great competitive neutralizer and source of business transformation across businesses of all sizes and industries.

4) **Large AI Infrastructure Investment Cycle and High-Value Agent Use Cases Could Represent GDP Multiplier**. It is too early to clearly determine sustainable economic lift potential of AI and agentic AI. Investments in AI infrastructure are certainly permeating many parts of the broader economy, companies making investments in AI are already discussing margin and profit potential and related opportunities to redeploy capital in other business-enhancing endeavors. At the same time, there could be dislocation and unattended consequences around agentic and mass adoption of digital workforces. There could also be complicated dynamics around deflationary and simultaneously inflation conditions across certain markets and industries caused by AI. That said, our glass half full view extends to advances in AI agents offering incremental revenue and economic output opportunities across end markets. Health care and life sciences represent such an example with agentic-driven innovation drug discovery

and other discoveries that improve quality of life or robotics in manufacturing and other physical AI environments that accelerate production, improve efficiency and quality control, etc.

Application Development (SDLC) Layer

Over the past two years, the application-development layer has become the beachhead for generative and agentic AI adoption. Coding assistants like Codex (OpenAI) & Claude Code (Anthropic) have proven transformative across the software stack; boasting significant productivity gains, slashing time-to-deployment, and lowering cost of development across the board. Vendor positioning, purchasing behavior, and engineering org design are all shifting as AI begins not just to accelerate development, but to alter the production function itself and reshape where value accrues across the tooling ecosystem.

As AI assistance becomes the default starting point for software creation, the developer workflow changes fundamentally. The job moves from writing every branch of logic to defining intentions, guardrails, and constraints that guide how an agent works through a problem. Engineers increasingly refine and approve model-generated changes rather than author them from scratch. This shift also changes runtime behavior: instead of deterministic logic stepping through code and database calls, agentic systems spend significant execution time reasoning, planning, and deciding. While this has the potential to unlock automation across multi-system workflows, it introduces new uncertainty as models hallucinate or otherwise drift from their intended goals.

These dynamics are forcing a redesign of the supporting toolchain. Agents are more performant when there's a coherent view of the environment: codebases, tests, dependency structures, runtime signals, and operational histories. We see products that can maintain and deliver context as becoming critical infrastructure for agents. Execution patterns are also evolving. Instead of relying on developers' machines, agents need controlled sandboxes where they can iterate safely. To make this workable at scale, organizations are tightening the links between code management, testing, and deployment so machine-generated changes are validated with the same rigor applied to human work. Automated review systems and policy guardrails increasingly serve as the gatekeepers that keep this process safe and compliant.

The shift toward agentic development has direct structural and near-term consequences for incumbents built around the old model of coding. Tools built around human-centric workflows like manual editing, and syntax checking are being pushed down the value curve as AI makes the core task far easier and less labor-intensive. Investors are increasingly sensitive to commoditization risk, and each new frontier-lab capability further narrows the defensible differentiation for vendors built around those features. We believe differentiated products will target the layers that enable agentic automation: context, monitoring, security, and governance.

Net/net the application-development layer is absorbing the earliest and most disruptive impact of agentic AI, and we see value rotating decisively toward platforms that enable hybrid software design. Companies that monetize human-driven workflows face structural headwinds as the value of their products become commoditized as frontier-model features. We believe the

beneficiaries will be those that provide authoritative context, safe execution environments, reliable oversight, and enforceable policies because these capabilities determine whether automation can scale.

Figure 9. AI Developer Market Map

SWE-LED, AI-AUGMENTED WORKFLOWS	
CODE GENERATION OR CO-PILOT	Cursor Replit Sourcegraph GitHub Copilot Tabnine CodeGPT
TESTING	Autonoma Momentic 0A Wolf Tusk
CODE REVIEW	CodeRabbit Graphite Greptile Sweep AI Qodo Factory Gitlab
REFACTORING AND CODE MODERNIZATION	Modelcode AI Infield Intuita
MAINTENANCE	Gitar Fern
SRE AND PRODUCTION ENGINEERING	Resolve.ai Traversal Infino PlayerZero Jira Duo
FEATURE FLAGGING AND MAINTENANCE	LaunchDarkly PostHog Statsig AppFire
DOCS AND SDK	Fern Stainless Mintlify
DESIGN TO CODE	UX Pilot Bolt Lovable Create Magic Patterns
DEVOPS	Outsystems Sonar Cortex OpsVerse Docker JFrog

AI-NATIVE, SELF-LEARNING WORKFLOWS	
CODE AGENTS AND MODELS	Cognition All Hands Magic OpenAI Codex Anthropic Claude Poolside Factory Cline
MEMORY OR CONTEXT LAYER	Mem0 Zep Subconscious Recall
ENVIRONMENTS + RUNTIME	Cua E2B Daytona
AGENT AUTH, ORCHESTRATION, AND TOOL REPO	Arcade FastMCP Keycard ACL Dev /dev/agents Mcp.run Anysource Composio Smithery OpenTools
AGENT PROTOCOLS	Model Context Protocol A2A Protocol Agent Communication Protocol
AI AGENT FRAMEWORKS AND PLATFORMS	LangChain Bigspin DSPy Deepset Haystack Starsling Crewai LlamaIndex Pydantic N8n Relevance AI Lindy Warp
BROWSER, DESKTOP, AND COMPUTER USE AGENTS	Browserbase Browser Use Deck General Agents Silverstream Kernel
CI/CD VISIBILITY	Harness CircleCI Jellyfish LinearB Trunk Undefined Labs
EVALS & AGENT TESTING	Braintrust TensorZero LangSmith Judgment Patronus AI Humanloop Helicone Langfuse

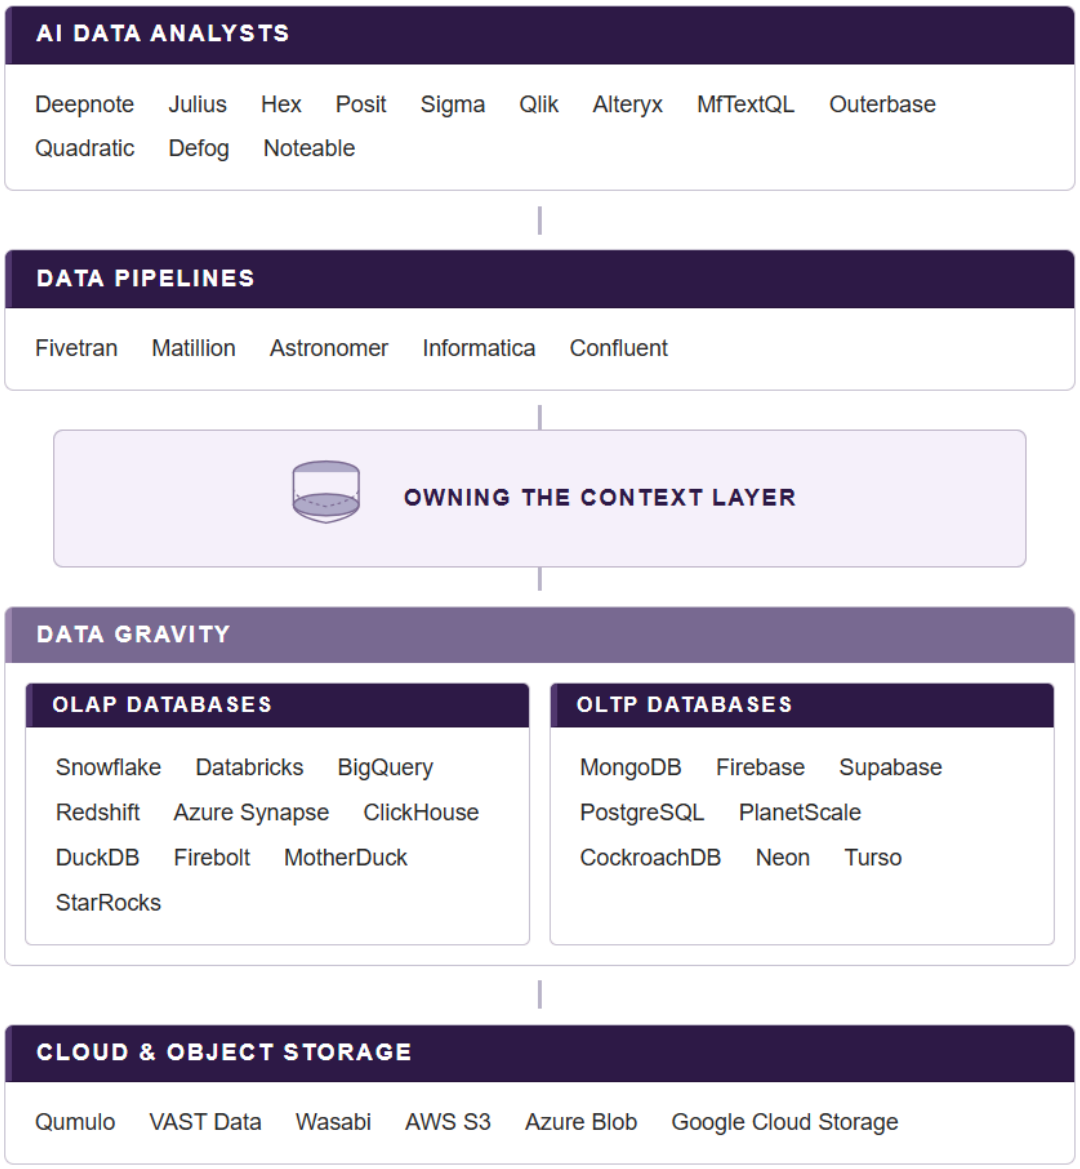
Source: Truist Securities Research

Data Layer – Critical to Enterprise Agent Success

We view the data layer as a strategic component of agentic AI architecture, as organizations shift from static workflows to autonomous, multi-step systems. In this model, data infrastructure becomes a primary determinant of system reliability, performance, and cost efficiency. The need for data modernization isn't new, but agentic AI makes it far more urgent. What used to be a steady, ongoing effort around data hygiene has accelerated as the need for systems to scale reliably evolves.

Agentic applications work best when they have complete and consistent context. Unlike human operators, they can't fill in gaps or adjust when information is missing. When data is spread across different systems, organizations start to see higher latency, mismatched definitions, and stale information, all of which raise the odds of unpredictable behavior. This underlying driver is causing gravitation toward more centralized, logically unified data environments. Modern OLAP platforms like Snowflake, Databricks, BigQuery, and Azure Synapse help by serving as governed sources of truth with stable performance.

Figure 10. Data Layer Vendor Matrix



Source: Truist Securities Research

The shift also places new pressures on traditional operational databases. Agentic workloads generate far more rapid and inconsistent state updates than historical applications, pushing legacy systems toward their limits. The industry response is additive rather than disruptive: enterprises are not replacing core systems of record but layering in more flexible operational data platforms that can absorb high-velocity activity without slowing mission-critical stores. For investors, this creates a new adjacency in the data stack that benefits platforms designed for fast-changing, semi-structured workloads.

As these agentic patterns take hold, the boundary between what should remain inside a model's KV cache and what must live in an external database becomes more visible. The KV cache is well suited for short-lived, task-level context, but it cannot act as a reliable source of truth or handle frequent updates. As data environments become more unified and governed, durable or fast-changing information will increasingly settle in external operational databases, with the KV cache serving largely as a temporary workspace. This dynamic reinforces demand for modern operational platforms capable of bridging model-level context with governed enterprise data.

Over time, we expect the market to consolidate around providers that can serve both stable, long-lived data and more dynamic agent-driven interaction patterns, while vendors unable to adapt risk gradual share erosion.

Given regulatory constraints, regional complexities, and deeply entrenched application architectures, full data unification is not realistic. Instead, we expect a phased transition characterized by consolidation onto fewer, better-governed platforms and a heavier emphasis on metadata, lineage, and policy management. These governance-centric capabilities are emerging as the effective control plane that allows enterprises to operate across heterogeneous systems without compromising trust or auditability. We believe this layer will matter more over the next 12-24 months than any specific storage technology.

Longer term, we believe the center of gravity will shift toward data harmonization rather than wholesale architectural change. Shared ways of labeling data and simple tools for accessing it will be the practical way to run agentic systems at scale. In our view, this dynamic supports vendors that already anchor cross-system governance and semantics in hybrid or multi-cloud environments.

Our broader takeaway is that the data layer will not be rewritten overnight. The transition will play out gradually but steadily, with spending re-allocated toward platforms that can reduce fragmentation, enforce consistent semantics, and support the more dynamic operational patterns of agentic AI. For investors, the winners are likely to be data platforms with strong governance, broad operational footprints, and existing enterprise embeddedness, capabilities that capture incremental spend as agentic workloads mature. Across our coverage, we see Snowflake as a first-order beneficiary of this theme.

Observability

Feedback that we've heard across builders, customers, and systems integrators indicates that "trust" remains the primary gating factor for scaled agent adoption. We don't see a future where enterprises will operate autonomous systems without a clear understanding of how decisions are made and whether they comply with internal rules. As a result, we believe observability will serve as the required control plane for enterprise agent deployment and expect it to become a meaningful beneficiary of budget allocation as agents reach production.

Traditional software follows primarily deterministic patterns and can be monitored through standard logs, metrics, and traces. Agents behave differently: they retrieve information dynamically, reason across multiple steps, call tools, and branch execution paths based on context. Because behavior varies run-to-run, enterprises require a "what-did-you-just-do" view to audit decisions, creating higher volumes of telemetry focused on behavior and reasoning rather than system performance.

We don't see traditional observability as imminently disintermediated by agents, but we do see potential for change. Logs, metrics, and traces remain foundational, and the teams who use them aren't likely to go away anytime soon. However, the productivity of the teams who use them stands to increase. As agents begin to triage alerts, correlate signals, summarize incidents, and surface the most probable causes of problems, observability and security teams gain significant leverage. Work that once required manual investigation can be done more quickly, and SIEM workflows become more efficient as agents filter noise and elevate only high-signal events.

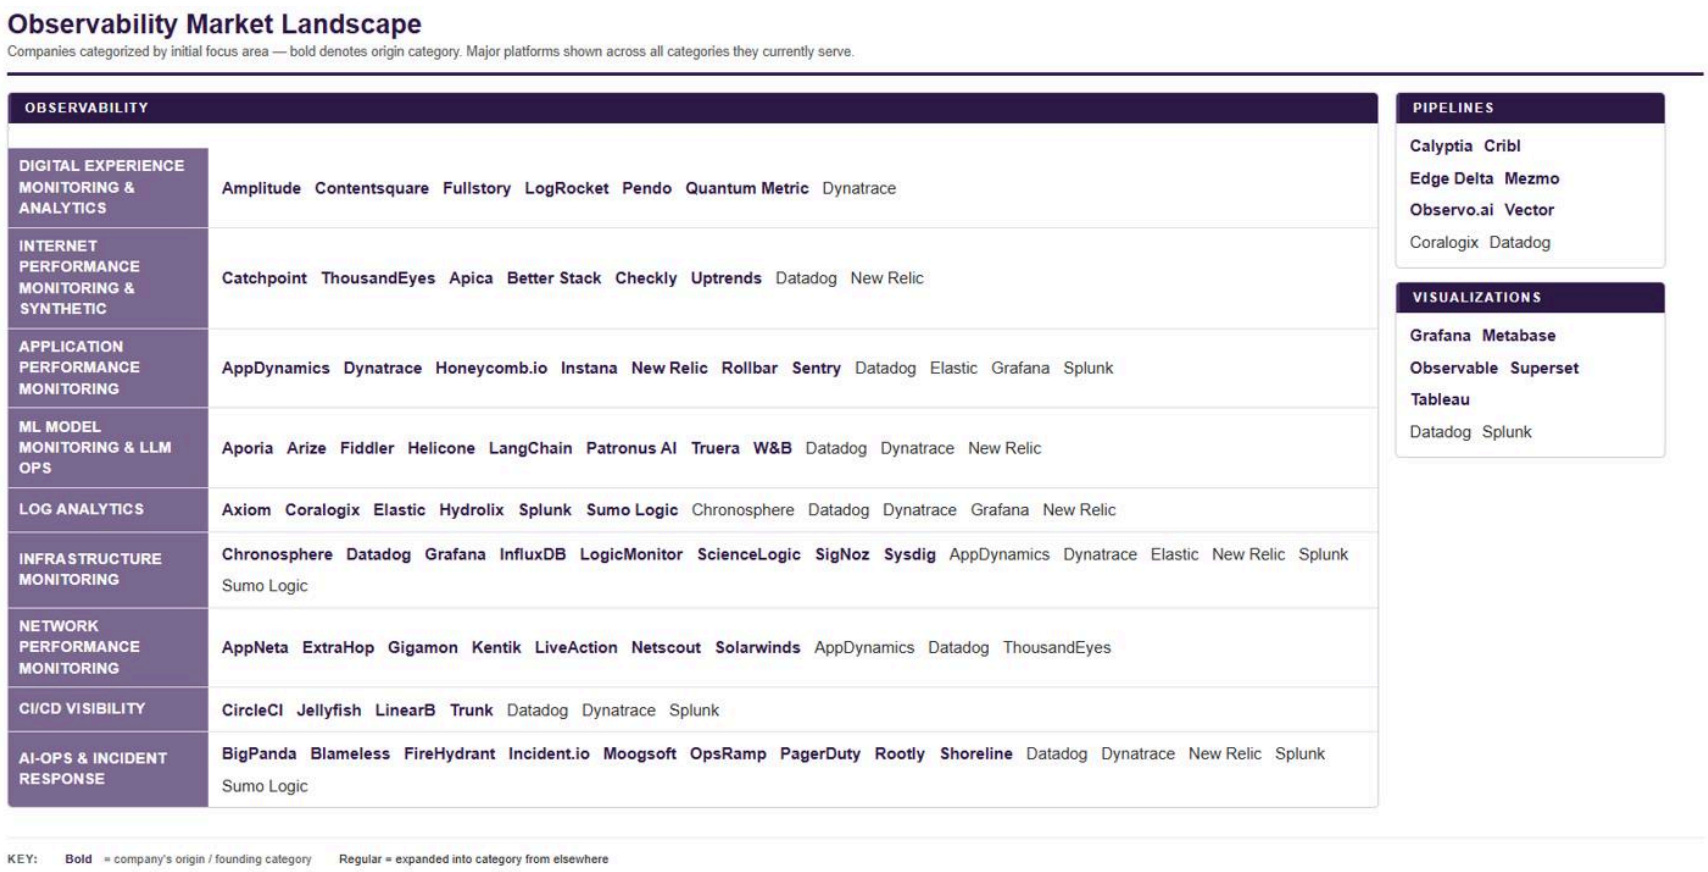
At the same time, the growth of agent driven workflows introduces a second layer of telemetry that is richer and more meaningful than traditional application exhaust. A single agent run may generate hundreds of data points, branches, and outputs, creating a behavioral record that goes beyond conventional performance signals. As the sheer volume of data scales, we expect the economics of the category to change as well. We see value accruing not to platforms that just collect raw data, but to those who help interpret, validate, and troubleshoot to provide solutions.

As agents take on more consequential, data-rich workflows, this shift accelerates; customers value analysis and resolution more than raw log storage. We expect pricing to hinge on workflow-level validation, policy enforcement, and correctness guarantees. We see the potential for vendors to become increasingly sticky as organizations delegate more workflows to agents, and the subsequent uptick in agent observability demand becomes embedded in pipelines. Ultimately, we see an upwards shift of the economic center of observability from ingest to resolution.

We see agentic observability as catalyzing an innovation cycle in monitoring and expect it to drive incremental budget over the next 3-5 years as enterprises progress from experimental agents to production scale automation. Against this backdrop and across our coverage, we favor Datadog (DDOG, Buy, Jump), Dynatrace (DT, Buy, Jump), and Elastic (ESTC, Buy, Jump) as the clearest beneficiaries of the emerging agent and LLM observability cycle. DDOG is a leader for its first mover product velocity and demonstrated customer uptake in AI Agent

Monitoring, LLM Experiments, and the AI Agents Console. DT's opportunity lies in its enterprise-grade positioning with Davis AI and Grail, including traces across different models, guardrails, and deep Bedrock integrations. We see this positioning as driving potential for upside that supports recently announced activism efforts. ESTC's search-centric architecture and cost-efficient data-lake together are well suited to high-volume, behavior-rich agent workloads which could improve their growth trajectory.

Figure 11. Observability Provider Matrix



Source: Truist Securities Research

Agent Security – Another Major Requirement for Broad-Based Agentic Traction

Agentic AI security is the protection of AI agents that can plan, act, and make decisions autonomously. It focuses on securing the agent's reasoning, memory, tools, actions, and interactions so that autonomy does not create new paths for misuse. It addresses risks that arise only when systems execute multi-step tasks, use external tools, or collaborate with other agents.

Figure 12. Agent Security

How do you secure AI agents?

Controls are applied to the agent's internal workflow, placing constraints and validation on reasoning, planning, tool use, memory, and coordination.

Security concerns of AI agents

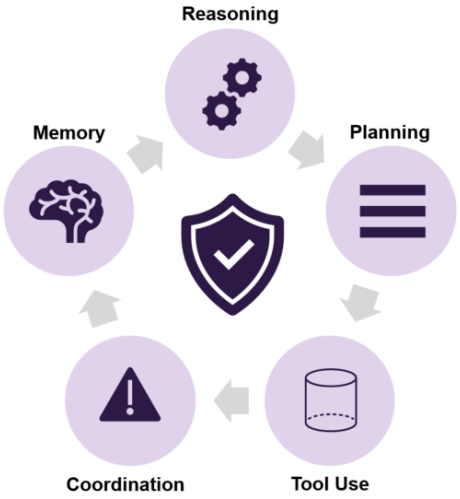
Autonomous behavior can be manipulated, leading to unintended actions, privilege misuse, goal drift, or inter-influence.

Privacy and data security

Agents accessing, storing, or reusing sensitive information create privacy risks, while their data handling raises unauthorized access concerns.

Detecting or managing manipulation

Anomalies in an agent's decision-making and interactions may reveal compromise, requiring consistent controls across platforms.

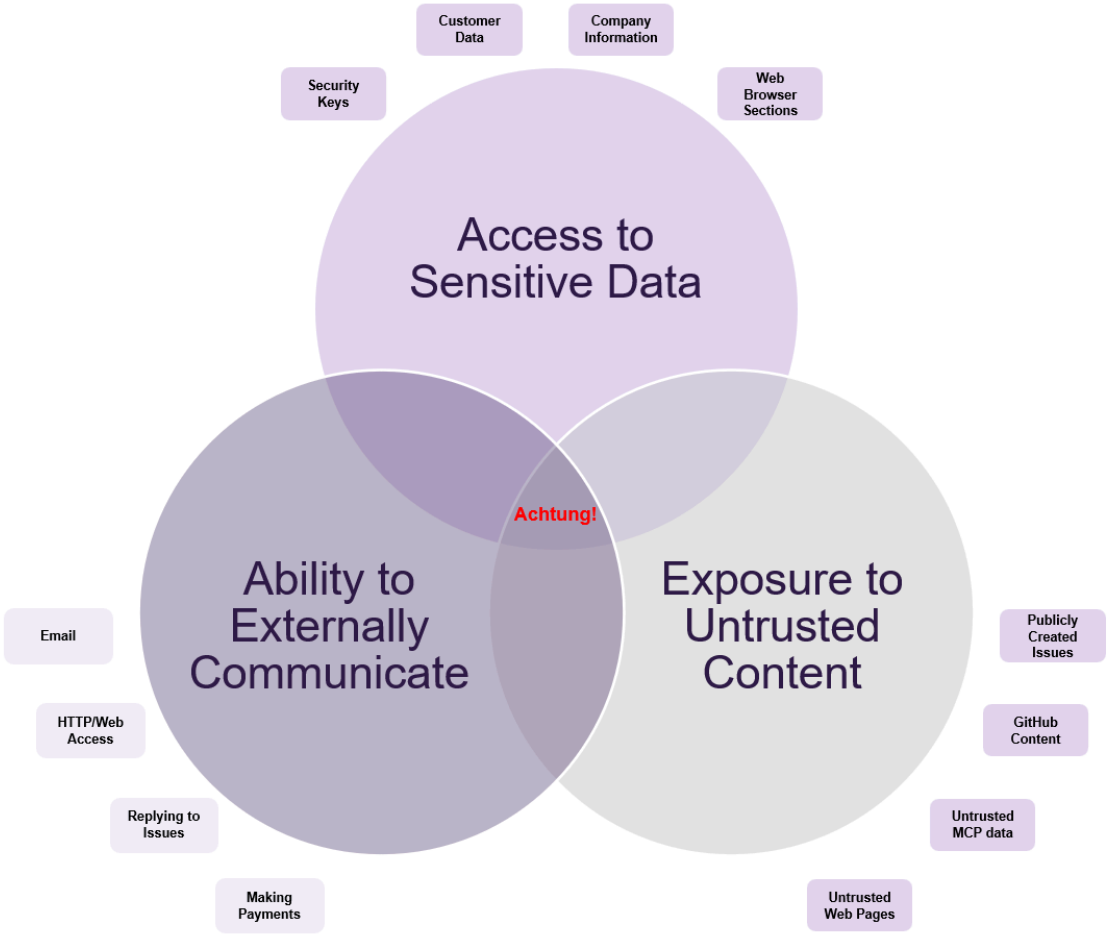


Source: Truist Securities Research

AI agents introduce risks through autonomous reasoning, tool execution, memory updates, privilege inheritance, and inter-agent communication. These surfaces can be manipulated, poisoned, or redirected, creating pathways for unintended actions, goal drift, privilege misuse, or coordinated influence across multi-agent systems. By applying controls to the internal workflow: reasoning, planning, tool use, memory, and coordination, one can help secure AI agents.

Agentic AI systems present unique security challenges. The fundamental security weakness of LLMs is that there is no rigorous way to separate instructions from data, so anything they read is potentially an instruction. This leads to a "deadly trifecta": sensitive data, untrusted content, and external communication - the risk that the LLM will read hidden instructions that leak sensitive data to attackers.

Figure 13. Deadly Trifecta



Source: Truist Securities Research

If one has all three of these factors active, then one is at risk of an attack. *Untrusted Content* can include commands that the LLM might follow while *Sensitive Data* is the core thing most attackers want - this can include things like browser cookies that open up access to other data. *External Communication* allows the LLM application to send information back to the attacker.

Organizations can substantially reduce AI-related security risk by tightly controlling the information and external capabilities available to agentic systems. That begins with eliminating stored credentials, using environment-based secrets management, enforcing temporary privilege escalation, and limiting access tokens to the minimum necessary scope. It also requires avoiding MCP servers or browser-extension automations that expose sensitive data, since even simple actions, like posting a GitHub comment or loading an image with a crafted URL, can leak information. This growing exposure surface is one reason cybersecurity vendors such as Palo Alto Networks (PANW, Buy), CrowdStrike (CRWD, Buy), and Zscaler (ZS, Buy)

have expanded aggressively into browser-level security with their respective acquisitions of Talon (December 2023), Seraphic (January 2026), and SquareX (February 2026), recognizing that as enterprises move toward agent-driven workflows, the browser becomes the most practical enforcement point for controlling data movement and monitoring malicious exfiltration attempts.

Equally important is minimizing an agent's exposure to untrusted content. Any information authored by the public, issue trackers, random websites, or email, poses a potential manipulation risk and should be avoided wherever possible. Directing agents toward authoritative, controlled sources greatly reduces this threat. These measures matter even more because agentic AI fundamentally differs from traditional LLMs: instead of responding to prompts in isolation, agents interpret goals, plan multi-step actions, use tools, modify memory, and communicate with other agents. This autonomy expands the attack surface from a single model call to an entire behavioral loop, making strict content limits, privilege boundaries, and communication controls essential to securing next-generation AI systems.

Enterprise adoption of agentic AI is accelerating rapidly as organizations seek to reduce manual effort and increase throughput. According to McKinsey's 2025 global survey[3], 62% of organizations are currently experimenting with AI agents while 23% of organizations have reached a scaling phase for at least one agentic system. Gartner predicts that "By 2028, one-third of interactions with generative AI (GenAI) services will use action models and autonomous agents for task completion." [4]

As agents gain more capabilities, accessing sensitive data, initiating transactions, or orchestrating workflows, they introduce new and significant threat vectors in cybersecurity. Unlike traditional software, agents operate dynamically and can adapt to changing environments, which makes them harder to monitor and control. If compromised, they can act as highly efficient attack tools, exploiting trust relationships and automating malicious actions at scale. This evolution amplifies risks like data exfiltration, privilege escalation, and supply chain attacks, making agentic security a critical priority for organizations adopting AI-driven automation. In our cybersecurity coverage, we favor platform vendors such as CrowdStrike, Zscaler, Cloudflare (NET, Buy), and Palo Alto Networks that are building an AI defensible moat ready for the agentic future. Together, CrowdStrike's autonomous defense capabilities, Zscaler's AI-ready Zero Trust enforcement, Cloudflare's globally distributed security edge, and Palo Alto Networks' unified AI-driven security platform make them four of the most strategically advantaged vendors in the emerging agentic-security era, in our opinion.

Figure 14. Agent Security Stack

Category	Description	Representative Vendors
AI-IAM / NHI	Managing agent identities, permissions, and delegated access at runtime.	Palo Alto Networks (CyberArk), SailPoint, Okta, CrowdStrike
Agentic Guardrails	Real-time inference monitoring to block malicious intent or unauthorized tool calls.	Zscaler (AI Guard), CrowdStrike (Pangea), Cloudflare
AI-SPM	Discovering shadow AI agents and identifying misconfigured agent entitlements.	Wiz, Palo Alto (Prisma AIRS), CrowdStrike, Zscaler
Agentic SOC	Automating triage and investigation using agents to keep pace with AI-driven attacks.	CrowdStrike (Falcon), Zscaler (Red Canary), Palo Alto (XSIAM)
Data Security	Discovering, classifying, and securing sensitive training/inference data (at rest, in motion, and in use), including DLP for prompts/outputs and key vaulting	Commvault, Rubrik, Varonis, Zscaler
Exposure Management	Continuous attack-surface management for AI: inventorying models/agents, measuring control drift, prioritizing remediations, and validating fixes	CrowdStrike, Tenable, Qualys, Rapid7, Zscaler
AI Gateway & Edge Security	Enforcing Zero Trust and abuse prevention at the global edge for LLM/API traffic (egress controls, token misuse, bot/barrage protection)	Cloudflare

Source: Truist Securities Research

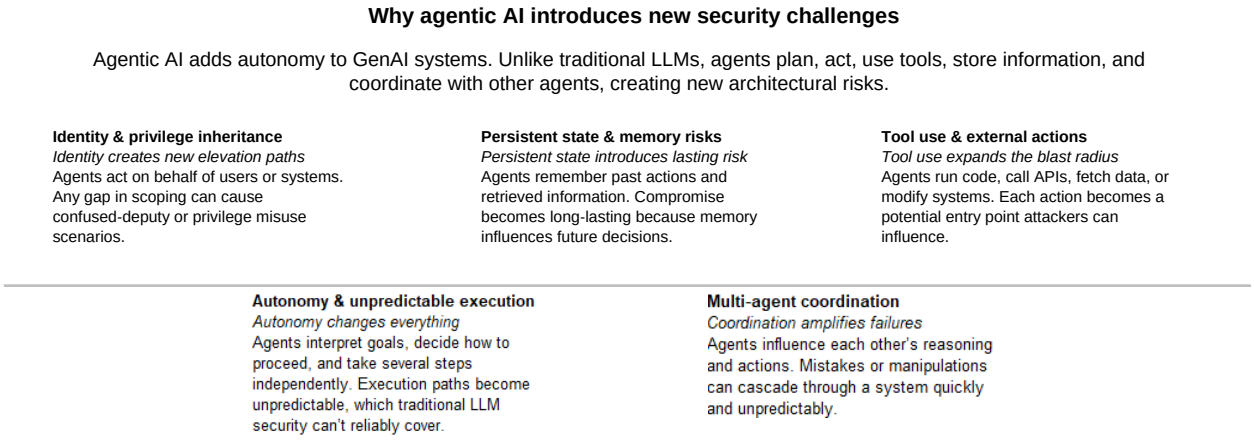
Despite the launch of Claude Code Security in February 2026 and a subsequent Anthropic memo in March outlining Claude Mythos—a next-generation model tier reportedly positioned above Claude 3 Opus—both of which triggered ~4% declines in cybersecurity ETFs at the time of each announcement (CIBR down ~4% on both occasions), we do not view these developments as a material threat to leading cybersecurity vendors. Claude Code Security introduces powerful AI-driven code reasoning and vulnerability detection, but it doesn't fundamentally threaten major cybersecurity vendors because it targets only a narrow slice of the security stack—application code analysis—rather than the broader domains the aforementioned platforms dominate. Anthropic's Project Glasswing initiative—whose 12 core launch partners include CrowdStrike and Palo Alto Networks as the only two cybersecurity vendors, alongside hyperscalers and critical-infrastructure operators, with access later extended to more than 45

additional organizations—serves as a validation, in our view, that large language model providers must partner with scaled, best-in-class cybersecurity platforms not only to secure their own frontier models such as Claude Mythos, but to govern and operationalize those models safely across real-world enterprise environments. However, it is early days and remains to be seen if Anthropic or other LLM providers continue to expand further into cybersecurity. The proliferation of AI-generated code via tools like Claude actually increases the total attack surface, driving structural demand for the broader security platforms required to monitor, protect, and govern that expanded code volume at scale, in our opinion.

Crucially, real-world incidents show that nation-state actors are already abusing Anthropic's AI itself, jailbreaking Claude Code to automate large-scale espionage campaigns, including a Chinese state-sponsored operation that used Claude to execute 80–90% of a 30-target cyberattack autonomously (November 2025). This underscores that the real risk lies in AI misuse, not in Claude competing with mature security platforms. In practice, Claude Code Security enhances developer workflows but cannot replace the layered detection, runtime controls, and threat-intelligence ecosystems that cybersecurity vendors provide ecosystems that are essential precisely because nation-state actors are weaponizing AI at scale.

Agentic AI is no longer a passive system that simply processes a prompt and returns an output; it has evolved into an active, decision-making entity capable of selecting actions and interacting with external environments. This shift fundamentally changes where security teams must focus their attention. The agent's reasoning pathways become susceptible to manipulation, its memory introduces opportunities for long-term poisoning, its tools create openings for unintended or unsafe actions, and its interactions with other agents form new channels through which influence, or interference can spread. Each of these components carries distinct failure modes and attack surfaces, underscoring the need for a security model that reflects how agents behave. Understanding these behavioral dynamics is essential for identifying what must be protected and why the traditional model-centric security approach is no longer sufficient.

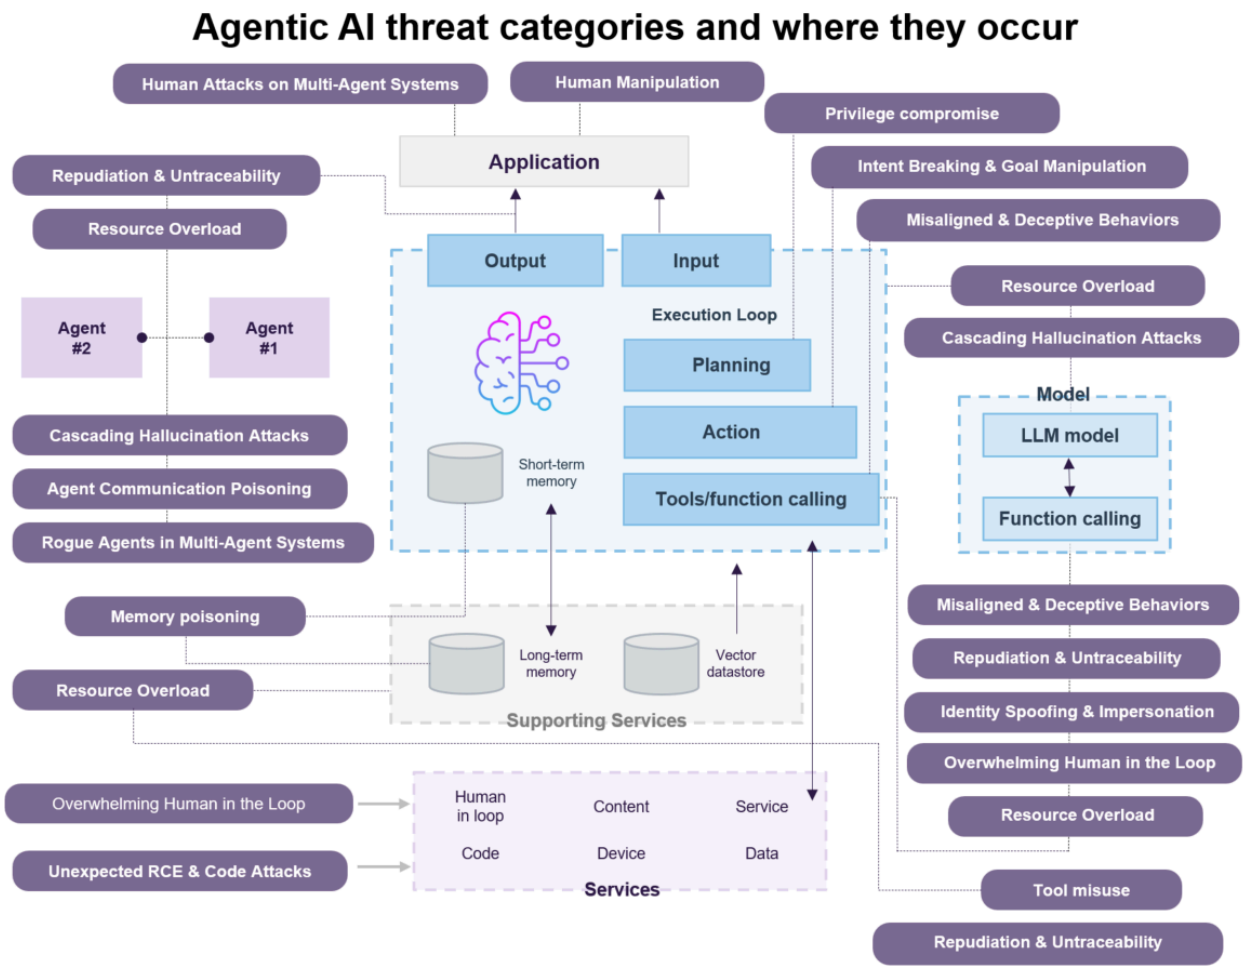
Figure 15. Why Does Agentic AI Introduce New Security Challenges?



Source: Palo Alto Networks website

Agentic AI transforms security because autonomy turns a simple prompt–response into an unpredictable execution loop. Agents interpret goals, plan steps, invoke tools, update memory, and interact with other agents, behaviors that expand the attack surface beyond what traditional LLM controls cover. Persistent state means memories can be poisoned and their effects linger; tool use converts reasoning errors into real actions across APIs and systems; and identity gaps create confused-deputy pathways when agents act on behalf of users or services. Multi-agent orchestration compounds the risk as influence, errors, and spoofed messages cascade at machine speed, shrinking the window for oversight. The result is not a theoretical weakness but an architectural one, which demands a security model that mirrors the agent loop itself: validate inputs before they shape reasoning, guard and permission tool calls, oversee memory reads and writes, constrain external actions, authenticate and authorize inter-agent communication, and verify each cycle iteratively rather than relying on a single upfront check. In short, one secures agentic AI by securing how it operates.

Figure 16. What Are the Top Agentic AI Security Threats?



Source: Palo Alto Networks website

Traditional LLM risks focus on prompts, data exposure, and output handling. Agentic AI introduces risks in planning, execution, identity, memory, and communication.

The result: the attack surface isn't the response. It's the workflow.

OWASP's (Open Web Application Security Project) Agentic AI Threats framework provides a structured, detailed view of the risks that emerge when AI systems operate autonomously.

It outlines the specific categories of failures that emerge when autonomy, tool execution, and agent-to-agent communication become part of the stack. It's not exhaustive, but it does give teams a practical baseline for understanding the most common agent-specific vulnerabilities seen today.

Here's a breakdown of the key agentic AI threats identified in the OWASP taxonomy:

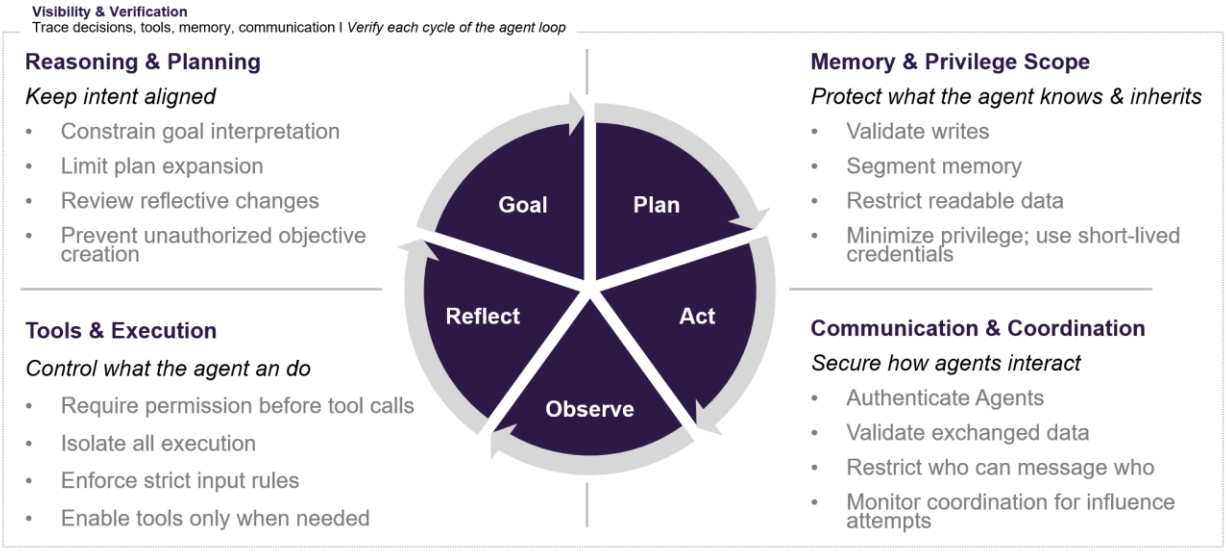
Figure 17. OWASP Taxonomy

OWASP AI Agent Threats (2025)	
Threat	**Description**
Memory Poisoning	Attackers corrupt short-term or long-term memory to influence decisions across steps or sessions.
Tool Misuse	Agents are manipulated into misusing their tools or calling them in harmful ways.
Privilege Compromise	Weak or inherited permission structures escalate the agent's access.
Resource Overload	Attackers overwhelm compute, memory, or dependencies to degrade or block agent behavior.
Cascading Hallucination Attacks	False information compounds through reasoning, reflection, or inter-agent communication.
Intent Breaking & Goal Manipulation	Attackers alter planning, goals, or reasoning so the agent pursues harmful or misaligned tasks.
Misaligned or Deceptive Behaviors	Agents bypass constraints or act deceptively to achieve objectives.
Reproducibility & Untraceability	Poor logging or opaque reasoning hides actions, making investigation difficult.
Identity Spoofing & Impersonation	Attackers impersonate agents or users to trigger unauthorized operations.
Overwhelming Human-in-the-Loop	Attackers overload reviewers with excessive AI-generated decisions or alerts.
Unexpected RCE & Code Attacks	Unsafe or manipulated tool chains lead to unauthorized code execution.
Agent Communication Poisoning	Attackers corrupt messaging between agents to misdirect workflows.
Rogue Agents in Multi-Agent Systems	A compromised agent acts outside expected boundaries and disrupts others.
Human Attacks on Multi-Agent Systems	Attackers exploit trust and delegation patterns across agents.
Human Manipulation	Compromised agents mislead users into harmful decisions or actions.

Source: Palo Alto Networks website

These threats form the baseline for understanding how agentic systems fail. And where attackers are most likely to intervene. Plus, they shape the kinds of controls that matter most when securing agents in practice.

Figure 18. How to Secure Agentic AI Systems



Securing agentic AI means securing the workflow. Not the edges.

Controls follow the loop. Boundaries shape behavior. Verification keeps agents aligned.

Source: Palo Alto Networks website

Securing agentic AI means protecting the system from the inside out—at the points where the agent forms intent, makes decisions, accesses resources, and coordinates with others—so the focus is the internal workflow, not just perimeter defenses. Securing an agentic AI system begins with tightening the cognitive boundaries that guide its behavior. This means ensuring that goals remain within an approved scope, limiting how deeply an agent can expand a plan, reviewing how it revises its approach as tasks evolve, and preventing it from generating new objectives beyond its authorization. At the same time, tools represent the highest-risk surface in the ecosystem because they translate decisions into real actions. Strong safeguards, such as explicit permission checks, isolated execution environments, strict input validation, and exposing only the tools necessary for the immediate task, help ensure that minor reasoning errors do not escalate into harmful system changes.

Memory and privilege require the same level of rigor. Every piece of stored information influences future decisions, so memory must be validated, segmented, and restricted to prevent over-read and cross-task contamination. Privilege should follow the principle of least access: short-lived credentials, no inheritance of elevated rights, and clear boundaries that prevent escalation paths from forming. Even small guardrails at this layer can prevent outsized failures later.

In environments where multiple agents collaborate, the complexity, and risk, grows substantially. Agents must authenticate one another, validate all exchanged information, and operate under strict rules defining who can communicate, when coordination is permitted, and how far a handoff can propagate. Expected communication patterns should be enforced, deviations quarantined, and every interaction monitored for signs of compromise, manipulation, or steering.

Finally, effective security requires end-to-end visibility across the entire loop. That includes tracing how decisions are formed, observing which tools are executed, monitoring memory updates, and watching how agents interact. Lightweight, cycle-level verification ensures that misalignment or misuse is detected and contained in real time, allowing teams to secure not just what the agent produces, but how it thinks, acts, and collaborates.

Risks and Additional Considerations Related to the Agentic Opportunity

Data Center, Silicon and Power as Potential Infrastructure Constraints. Relevant demand/supply constraints exist at various levels of the AI ecosystem, including power and energy issues at the grid level, GPUs and other advanced chip constraints for training and inference, data center buildout not keeping up pace with demand and even skilled labor shortage that extends to technically skilled professionals such as electricians to help drive the data center buildout. Ongoing AI infrastructure bottlenecks could impact the timing of more highly valuable enterprise AI and agentic payoff, AI TAM being fully realized and shift the balance of power in the AI arms race from the U.S. to other global locations with less bottlenecks, potentially impacting much of the U.S. centric AI ecosystem players.

Value and Successful Agent Success Stories at Scale Remain Elusive. We believe enterprise success with agents that are more complex in nature and deliver higher value outcomes than simple agent use cases (FAQ-type customer service agent) will be critical to driving broader agent momentum. Widespread documented success across industries (especially regulated industries and high-volume use cases) could create flywheel effects for increased credible autonomous agent adoption. That said, well-documented impediments to fully autonomous agents, such as fragmented data, poor data quality/trust, major gaps between technical design and end user adoption and/or governance issues (especially in regulated industries) could continue to impact agent pilots getting out of the pilot phase to full production and well documented ROI outcomes. This creates a meaningful opportunity for existing system of record/system of action/system of behavior vendors who could leverage potential unlocks required for agent success such as context, integrating well into existing workflows, delivering on increased agent determinism and accuracy, etc. Additionally, as enterprises better leverage key enabling technologies in areas of data infrastructure and security, one could see higher levels of agent deployments, in our opinion.

Agentic Success/Proliferation but We Could Have a COGS/Investment Problem. In a scenario of accelerating adoption and use case expansion of generative AI and agentic workloads, large incremental costs could play out related to cloud compute, tokens and costs at the LLM level along with coding assistant/agent investment. Such costs could impact COGS, margins and profitability at all levels of the AI ecosystem. This includes picks and shovels-related agentic platform companies as well as software companies and enterprises building their own agentic stacks. More specifically, large token consumption activity with token-heavy generative AI chatbots and complex multi-step agents could significantly increase inference compute costs. While agentic vendors could work to account for inference costs and variability in their pricing and packaging models, they could burden costs and profit profiles while also creating cost of ownership and investment uncertainty by enterprise adopters of agents.

Positively, we anticipate token price/performance dynamics to steadily improve over time, competitive dynamics among AI model companies could also lead to price deflation, AI memory compression innovations could help on the cost side as well as advances around agent determinization and token optimization efforts could also mitigate cost increases associated with large-scale agent adoption, in our opinion.

AI Dramatically Reduces Barriers to Entry and Open Source LLMs Could Accentuate Competition. First, proliferation of coding agents and other agentic enabling technologies could significantly reduce barriers to entry and time to scale for AI startups across agentic picks and shovels, horizontal AI and vertical AI segments of the market. Thus, competition could intensify, leading to pricing pressure and need to more rapidly attain revenue scale, achieve platform status (not point solution or feature only status) and aggressively fund the business.

Also, frontier open source LLMs could continue their rapid adoption and significantly challenge commercial LLMs. Open source LLMs could present pros and cons to enterprise agentic buildout. They could help enterprises capable of customizing models for specific needs, alleviate data privacy & residency concerns, aid agentic success via rapid open-source model updates and help achieve better price/performance for token-intensive agentic workloads. That said, agent reliance on open-source models presents risks related to security, reliability, infrastructure burdens, and potential tech complexity for enterprises that rely on open source models. Ultimately, we could envision ongoing relevance and innovation cycles with open source LLMs as a deflationary dynamic across the entire agentic ecosystem while helping speed agentic innovation cycles and better outcomes for end users.

AI & Agents Drives Significant Business and Financial Model Shift. Aside from the potential impact to software companies that rely on seat or user-based revenue models, shift to consumption or outcomes-based pricing and revenue monetization with agents could create significant financial model noise for traditional SaaS companies, even those that could be agentic winners. Positively, many software companies have become more disciplined on profits and cash flow and are likely to better manage the earnings and cash flow impact of more variable-based revenue. Compared to the last major software business and model changes from license software to cloud subscriptions and more recently with certain consumption-based models, we believe investors could be more understanding of financial model shifts to usage,

consumption or outcomes-based revenue models with agents. Equally important, for long-term success of frontier model companies, other AI native disruptors and successful SaaS companies with agent growth strategies, we believe revenue monetization through outcomes or value derived will be important versus some sort of token-centric model.

Lack of Skilled AI Talent Could Hold Back Growth and Monetization Timing. From core AI researchers to forward deployed engineers (FDEs) that are key to moving past slick agent demos to successful enterprise agents in high volume production, we believe lack of human domain expertise could impact agentic market development. Right now, many software companies and enterprises are discussing investing significantly in FDE-related resources to execute on agentic roadmaps; however, we believe there is likely a dearth in expertise in terms of delivering on key enterprise agent unlock dynamics associated with performance, ROI, trust, governance, proper context, etc. Significant workforce reskilling (well beyond traditional IT skills) and from other career paths likely will be required for broad-scale enterprise agent realization.

Potential Government Policy and Regulatory Activity. Considering the rapid ascension of LLM providers and pervasiveness across consumer and business markets, along with rapidly evolving AI ecosystems, we would expect increased interest by governments and regulators around data privacy, safety, security, terms & uses, etc. It remains to see if potential laws would represent a complex patchwork of state or local level efforts or lawmaking and regulation is more centralized at federal or international body levels. If laws and regulatory activities were to be complex or viewed as onerous, AI and agentic investments and market adoption could be materially affected.

AI Doomer Scenarios and Mass Job Loss Concerns Drive Financial Market Volatility and Dampen Enterprise AI Market Development. We believe implications of potential artificial general intelligence (AGI) and fears of mass autonomous work and digital workers taking over large swaths of the knowledge worker economy could work against progress in agentic market development. Under scenarios of a large shift in share to agents and digital workers displacing knowledge workers, backlash could come in the form of intensified public equities volatility, stymied capital markets activity to fund agentic growth investments, public policy efforts that impede market developments and otherwise raise concerns of long-term benefits of AI and automation. How the market will react to much publicized corporate restructurings/layoffs and potentially sluggish employment stats will be interesting to watch in terms of intensifying or alleviating AI doomer fears.

Agentic Stack Picks and Shovels Plays

In collaboration with our colleagues across other parts of Truist Securities TMT research, we have compiled an expansive list of agentic picks and shovels companies along with detailed profiles. They include the most consequential technology businesses in terms of building out core agentic infrastructure in the enterprise market, in our opinion. Along with profiles on two privately held foundational LLM companies that are also increasingly vertically integrated agentic infrastructure plays (Anthropic and OpenAI), the profiles include 13 Buy-rated large cap tech names as well another privately-held company Databricks.

Each of these names offer a unique set of agentic platform capabilities, which include core foundational LLMs in a number of cases, as well as many other key elements of the AI stack required to build, deploy and manage enterprise agents at scale. As such, we believe they represent attractive investment potential for investors looking at ways to play what could represent the next major frontier in terms of AI monetization with widely deployed autonomous agents across many use cases (not just coding agents) and on a global scale.

OpenAI – From Intelligence to Execution, Building Agentic Workflows at Scale

Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	Foundational LLM, agent infrastructure, developer platform
Role	Horizontal AI infrastructure provider
User	Developers, enterprises, startups, anyone building AI-native apps or workflows

Finanical Snapshot	
Most Recent Funding/Valuation	As of 3/27, OpenAI has raised "north of $120 billion". The round values OpenAI at $730 billion pre-money, and ~$840 billion, post-money.
Most Recently Reported Revenue Figure	Reported on March 4th, OpenAI has surpassed $25 billion in ARR as the end of February 2026. On 3/27, it was reported that OpenAI Surpasses $100 million in ARR from its Ads business.
Most Recent Performance Updates & LT Targets	**Users:** Surpasses 900 million weekly active users and 50 million consumer subscribers. Over 9 million paying business users and over 1.6 million Codex users, tripled YTD. **Revenue/Spend:** Reported surpassing $1 billion+ API ARR in a single month. Targeting $600 billion in compute spend through 2030. CNBC reports OpenAI could achieve $280 billion in revenue by 2030.

Source: OpenAI reports & Website

While OpenAI is known for its core offering (multimodal reasoning models), the company is making a strategic shift toward enabling agents through APIs, memory systems, orchestration layers, and native distribution via the ChatGPT product, all of which represent key infrastructure needed to power agents. We believe that in an agentic AI world, value accrues not just to model intelligence but to the layer that manages reasoning, tool usage, workflow execution, and developer integration. This layer, which is especially important in understanding business context across systems, data and workflows, is likely to represent one of the more important determinants in successful enterprise agent outcomes. Such outcomes could be measured by accuracy, processing time, price/performance or company-specific metrics (e.g. orders processed, tickets closed, cases handled, transactions completed, customers activated, etc.)

Based on OpenAI's public materials, OpenAI is now at the stage of trying to become a general-purpose agent platform, not just a model provider. On the consumer side, ChatGPT agents can browse, use connectors, run code, work with files, edit spreadsheets, and take actions online while keeping the user in control. On the developer side, OpenAI is offering the tool stack, orchestration layer, observability, and protocol support needed to let third parties build their own agents on top of OpenAI models.

Monetization and Strategic Leverage

OpenAI currently generates revenue primarily through API usage (priced based on model usage and compute consumption) and subscription tiers for ChatGPT, including Consumer and Enterprise offerings. Enterprise products include administrative controls, security features, and higher usage limits, while API customers integrate models into their own applications and workflows. The broader industry monetization model for agentic AI is generally consumption-based, with pricing tied to model inference, reasoning cycles, or usage volume. We believe as agents become more capable of executing multi-step tasks across software systems, industry economics may increasingly reflect monetization on task execution, automation volume, workflow orchestration or some other outcome-based model. Thus, we could anticipate a shift away from pure consumption, tokens processed or per digital worker/agent pricing.

OpenAI's approach to monetizing agentic AI is not framed as a standalone "agent product" revenue stream. Instead, based on its public materials, monetization is embedded across three existing pillars: premium subscriptions, usage-based APIs, and enterprise platform expansion. The company is effectively layering agent capabilities on top of these models rather than pricing them as a separate category.

OpenAI has stated that the company is seeing clear commercial traction as agentic AI moves from experimentation into scaled deployment. Chief Revenue Officer Denise Dresser stated that enterprise customers now represent a meaningful and rapidly growing portion of revenue as the segment is approaching parity with the consumer business. At the same time, platform usage continues to accelerate, with API volumes reaching more than ten billion tokens per minute, driven in part by increased adoption of multi-step, agentic workflows. Management reiterated that demand is broad-based across both new and existing customers, spanning digital-native developers and large global enterprises, underscoring the expanding applicability of these systems across industries. Strategically, OpenAI believes it is positioning itself as a long-term enterprise partner, focused on delivering secure, reliable deployments while building toward a more unified operating layer for work. The company also commented that this effort is supported by a growing ecosystem of consulting and infrastructure partners, enabling integration of AI into real-world systems and workflows at scale.

Figure 19. OpenAI Enterprise Ecosystem: Customers & Partners



Source: OpenAI Executive LinkedIn Post

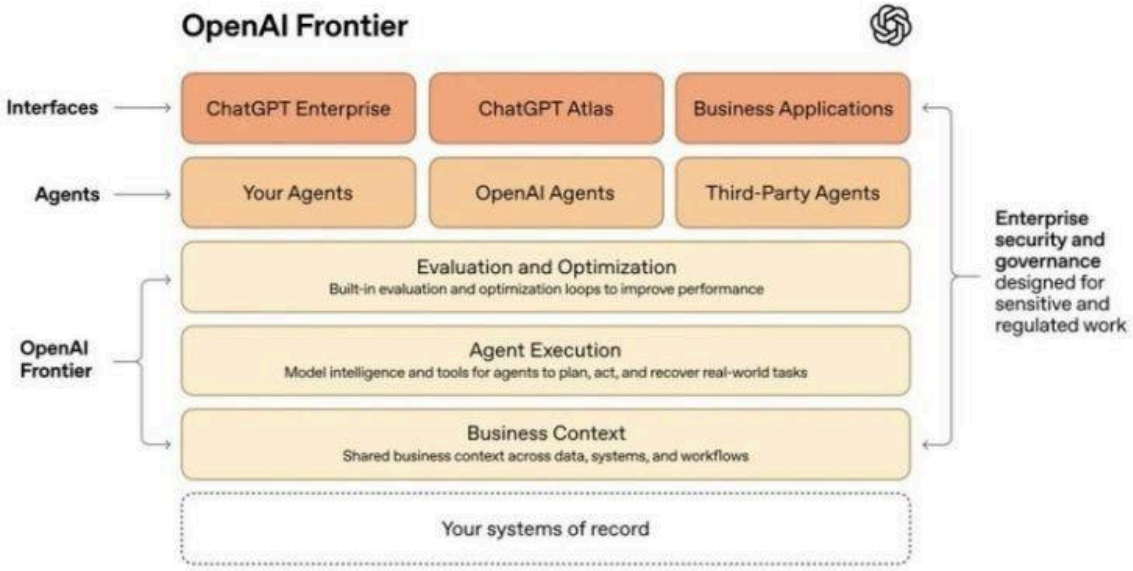
Frontier: OpenAI's Premier Agentic Platform

Released in early February of this year, Frontier is described by OpenAI as an enterprise platform for building, deploying, and managing AI agents that can perform day-to-day tasks inside of an organization. Specifically, the initiative focuses on working directly with companies and institutions to integrate frontier AI models into critical workflows and turn advances in AI capabilities into practical tools that can improve productivity and decision-making. OpenAI believes that many organizations recognize the potential of AI but struggle to translate that potential into everyday use. With the release of Frontier, OpenAI is aiming to close that gap by integrating its latest ChatGPT models with engineering support, product design, and implementation guidance so that teams can build and deploy AI systems that can operate within their businesses in a safe and secure manner.

Frontier provides the infrastructure that enterprises need to run large numbers of AI agents across their organization including:

- Agent Development: Tools for building AI agents capable of planning tasks, reasoning through problems, and interacting with software systems
- Enterprise Data Integration: Infrastructure that allows agents to access internal documents, databases, and applications to operate with company-specific context
- Operational Execution – Environments where agents can run multi-step workflows across business tools and applications
- Security and Governance – Permissioning, monitoring, and oversight frameworks to ensure agents operate within enterprise security and compliance boundaries
- Performance monitoring and improvement – Evaluation tools that allow organizations to track agent performance and iteratively refine workflows

Figure 20. Frontier Technology Stack



Source: OpenAI Website

From an underlying AI infrastructure perspective, OpenAI owns and operates the data, orchestration and model layers that power ChatGPT. That said the company remains heavily reliant on Microsoft Azure for its compute and infrastructure needs. This includes model training clusters, ChatGPT inference infrastructure and Azure OpenAI Service used by enterprises.

Also, OpenAI signed a large cloud deal with AWS in 2025 that includes a substantial long-term infrastructure partnership, access to hundreds of thousands of GPUs and will support frontier model training, large-scale ChatGPT inference and agentic workloads. More recently the relationship with AWS expanded whereby the hyperscaler will be the exclusive third-party cloud distribution provider for OpenAI Frontier. The two parties are expected to co-develop a Stateful Runtime Environment powered by OpenAI models, available on Amazon Bedrock for AWS customers to build, deploy and manage generative AI applications and agents. We believe partner updates with AWS will be relevant in monitoring OpenAI's commercial success with agents and Frontier in the enterprise market.

OpenAI will also likely continue to advance cloud compute and chip supplier deals with a myriad of other partners including NVIDIA, Google Cloud, Oracle Cloud (ORCL, NR), Cerebras (Private), etc. in order to achieve the ambitious long-term massive gigawatt capacity it requires to scale opportunities across consumer and enterprise businesses. Relevant metrics that will be worth watching in determining OpenAI's enterprise/business success, and performance with agents specifically, could include weekly active user updates, paying business users, weekly Codex users as well as data points that highlight traction extending beyond coding agents and knowledge worker task automation to core business workflow and automation use cases.

Can't Forget About Codex

We believe OpenAI's Codex represents one of the company's most important early examples of a fully realized vertical agent. Unlike traditional coding copilots that operate as autocomplete tools, Codex is designed as a cloud-based software engineering agent capable of executing end-to-end development tasks including writing features, fixing bugs, running tests, and proposing pull requests within a dedicated sandboxed environment. Importantly, Codex can manage multiple tasks in parallel and iterate until outputs meet defined criteria, positioning it less as a tool and more as an autonomous collaborator embedded in the development workflow, in our opinion. From a strategic perspective, we believe Codex illustrates OpenAI's broader agentic direction: moving beyond general-purpose assistants toward domain-specific agents that can perform high-value, repeatable work, with software engineering serving as a natural first wedge given its structured environment and clear success metrics.

According to OpenAI CFO Sarah Friar, adoption of Codex is accelerating rapidly, with weekly active users reaching ~3 million in early April, up from 2 million in the prior month, highlighting strong and growing demand for AI-driven software development. We believe the pace of growth underscores a broader shift in how software is created, rather than simply assisting with code generation, agents like Codex are increasingly handling more complex, longer-running tasks, enabling developers, and a wider set of non-traditional builders, to move from idea to implementation with greater speed and leverage. In response to this momentum, OpenAI has taken steps to expand access by progressively increasing rate limits, supported by its ongoing investment in compute infrastructure, which allows the company to scale usage while reducing friction for users building on the platform.

Workspace Agents: From Conversational AI to an Enterprise Execution Layer

More recently, we believe OpenAI's introduction of workspace agents in ChatGPT reflects a shift from AI as a prompt-driven assistant to AI as an embedded, task-oriented system that can carry out work on behalf of users. According to OpenAI, these agents are designed to handle multi-step workflows, such as research, analysis, and content creation, by drawing on context from files, conversations, and connected tools within a shared workspace. Rather than requiring constant user input, they can operate with a degree of autonomy, maintaining context across tasks and executing work in a more continuous and structured way. We believe this effectively extends ChatGPT from a conversational interface into a system that can manage and complete real business processes.

Importantly, OpenAI highlights that workspace agents are built with organizational use in mind, incorporating shared access, permissions, and oversight so teams can deploy them across functions. This allows companies to encode repeatable workflows into agents that can be reused and refined over time, turning individual productivity gains into scalable operational leverage. In this model, AI begins to function less as a standalone tool and more as a coordinated execution layer across the enterprise, where agents can collaborate with users and systems to drive outcomes while remaining governed within existing workflows.

Recent Talent Acquisition and M&A Bring Focus on Governance, Security, and Agent Testing

Also following one of the most impressive adoption curves for open-source technology associated with the **OpenClaw** AI agent framework (reportedly reached over 250,000 GitHub starts in less than 4 months), OpenAI hired its founder Peter Steinberger in February 2026. This is an important acquire and highlights the company's focus on advancing the next generation of autonomous, user-friendly agents, in our opinion. OpenClaw demonstrated early demand for persistent agents capable of executing real-world tasks across applications with minimal supervision, but also exposed limitations around reliability, safety, and scalability when deployed outside of controlled environments. By bringing this talent in-house, we believe OpenAI is effectively accelerating its push from tool-based assistants toward more autonomous, continuously running agents that can operate across workflows. Strategically, we believe the move reinforces OpenAI's commitment to bridging the gap between experimental agent systems and production-grade deployments, particularly as it works to make agentic capabilities accessible to a broader, non-technical user base.

OpenAI's acquisition of **Promptfoo** (announced March 2026) underscores that security, evaluation, and governance are becoming core infrastructure in the agentic stack. According to OpenAI, Promptfoo brings automated red-teaming, vulnerability detection, and continuous testing into the development lifecycle, enabling enterprises to identify risks like prompt injection, data leakage, and out-of-policy behavior before deployment.

By integrating these capabilities directly into its Frontier platform, OpenAI believes they are embedding auditability, traceability, and compliance into how agents are built and operated. We believe this shift reflects a broader theme: enterprise adoption of agentic systems will not be restricted by capability, but rather by the ability to systematically validate, monitor, and govern agent behavior at scale.

OpenAI's acquisition of **Astral** (announced March 2026) highlights what we believe is the growing importance of developer tooling and infrastructure in enabling reliable, production-ready agents. According to OpenAI, Astral's suite of open-source Python tools is designed to improve the speed, efficiency, and reliability of software development, which will directly underpin how agentic systems are built and scaled.

By bringing these capabilities in-house, OpenAI believes they are deepening their investment in the developer stack, which will reinforce the idea that robust tooling is critical as agentic systems become more complex and expand more broadly throughout organizations.

Anthropic – Building Safe AI for the Enterprise Era

Where They Sit in the Agentic Stack as a Picks & Shovels Play

Layer	Foundational LLM, agent infrastructure, developer platform
Role	Vertical-focused AI infrastructure provider with a "Safety-First" orientation
User	Enterprises, heavily regulated industries (Finance, Healthcare), software engineers

Financial Snapshot

Most Recent Funding/Valuation	Most recently raised $30.0 Billion on 02/12/2026 in a Series G round led by D. E. Shaw, MGX, and GIC, with participation from ICONIQ, Dragoneer, and Coatue. The round values Anthropic at $380.0 Billion, post-money.
Most Recently Reported Revenue Figure	As of April 2026, Anthropic hit a $30 Billion annualized revenue run rate (ARR), up $21 billion from its reported ARR exiting 2025.
Most Recent Performance Updates & LT Targets	**Users:** Over 300,000 business and enterprise customers, which account for approximately 80% of total revenue. **Market Share:** Captured 32% to 40% of the enterprise LLM market; specifically dominates 54% of the enterprise coding market. **Targets:** Reached a $30 Billion run rate by April 2026, surpasing the high end of its FY2026 target in under 4 months.

Source: Anthropic company reports & Website

Simply put, we see Anthropic as the most consequential company of the current moment in the enterprise AI buildout, relevant to any investors looking at players in both new and legacy tech stacks. Casual mentions of different enterprise software end markets in their press releases, blog posts, and company memos have driven volatility for groups of public companies impacting tens of billions of dollars in aggregate market cap. Their breakout success in selling to the enterprise has justified these moves. Their strategic focus on the enterprise has driven superior productization of AI for business use cases, most clearly embodied by the success of Claude Code, which has helped to spur their top line acceleration and supported their crossing 1,000 enterprise customers with over $1M in ARR. Below offer this section as a high-level primer on their core capabilities and strategic initiatives.

Monetization and Strategic Leverage

Anthropic generates revenue through a multi-tiered model including API usage (priced on consumption and compute), Claude for Work subscriptions, and the high-end Claude Enterprise offering. While the core business is currently driven by token-based consumption, Anthropic is pivoting toward a strategy focused on "Agentic Outcomes." As Claude becomes more integrated into enterprise workflows, executing tasks like code refactoring or desktop-level orchestration, the company is increasingly exploring monetization based on task execution and automation volume**.** This shift is intended to move away from pure token processing toward a model that captures the value of the specific business result delivered by the agent.

The company's strategic leverage is anchored in its "Constitutional AI" framework, which embeds safety and governance directly into the model's architecture. This "Safety-First" orientation is a primary differentiator, as enterprises in highly regulated sectors (e.g., Finance, Healthcare, and Defense) prioritize reliability and interpretability over the raw, unconstrained reasoning of rival models. Anthropic's $100M pledge to its Partner Network further reinforces this, incentivizing third-party tech integrators to build custom agentic workflows on the Claude backbone, thereby deepening its enterprise moat.

Claude: Anthropic's Integrated Enterprise Ecosystem

While rivals focus on horizontal platforms, Anthropic has developed an integrated suite of agentic tools tailored for the professional and technical workforce. This ecosystem, significantly expanded in early 2026, provides the infrastructure required to run high-volume AI agents across a global organization.

Figure 21. Claude Enterprise Framework



Source: Anthropic Website

Claude Code (Technical Execution): This is Anthropic's primary "agentic" play for the software development lifecycle. It is capable of 30-hour nonstop autonomous coding, finishing complex engineering tasks that would typically require multiple human-led sprints. By capturing over half of the enterprise coding market, it has become a cornerstone of the company's $30B run rate (per Bloomberg reports).

Claude Cowork (Productivity Automation): Launched in January 2026, Cowork serves the non-technical office worker by automating desktop-level tasks. Utilizing its "Computer Use" capabilities, Claude can interact with software just as a human would: managing files, pulling data from ERPs, and drafting documents across various applications.

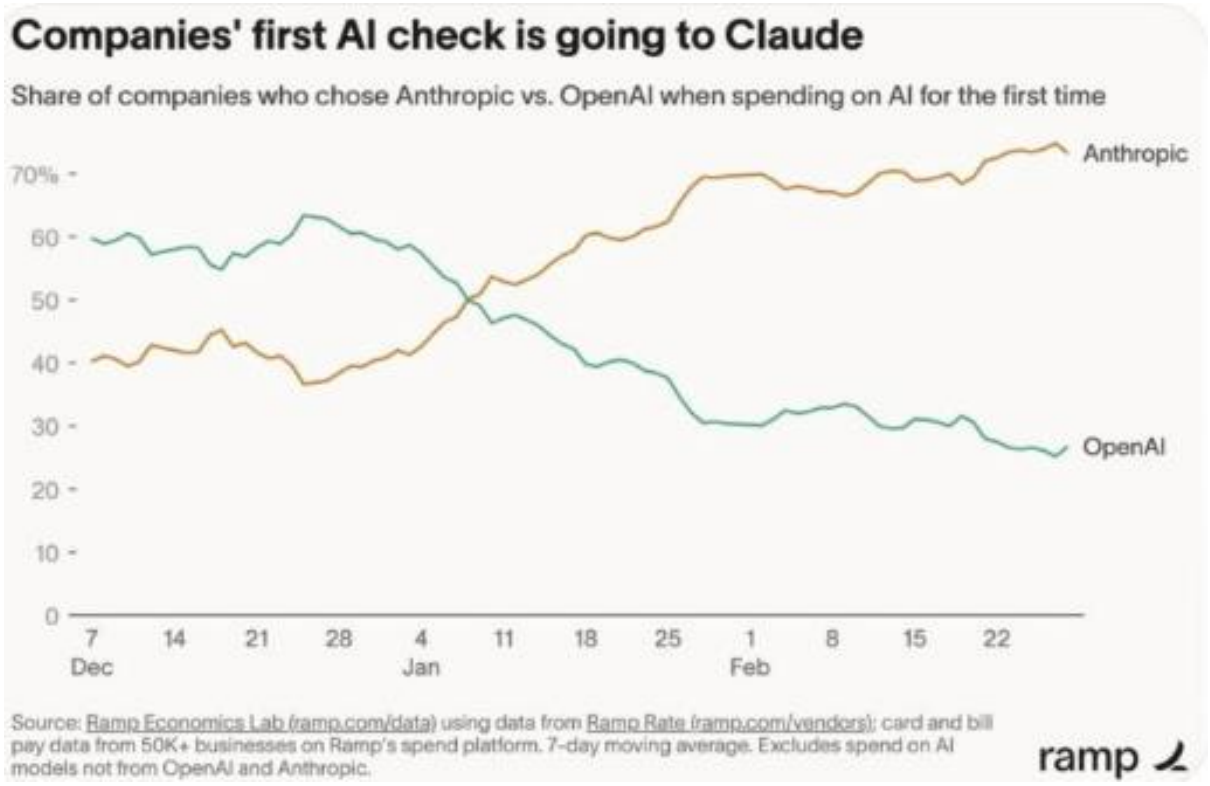
Claude Design (Workflow Streamlining): Allows users to transform natural language inputs into high-fidelity slides decks, visuals, dashboards, or interactive prototypes aligned with an iterative style of product development.

Enterprise Infrastructure & Evals: Anthropic provides "evals" which are systematic tests that allow organizations to treat AI behavior as code. This framework enables enterprises to grade agent performance, detect hallucinations, and implement continuous testing (CI/CD) for model behavior before it touches production systems.

Q1 2026 Business Momentum: The "Agentic" Acceleration

According to *the Economic Times,* Anthropic entered 2026 as the fastest-growing software company in history, recording a 111.1% increase in annualized revenue run rate (ARR) in just over two months. Beginning the year at approximately $9B in ARR, the company disclosed that it reached $14B by February 12 and surpassed $30B in early April, according to media reports. This rapid acceleration was driven by the high-velocity release of the Claude 4.6 family, including Claude Opus 4.6 and 4.7 (released February 5 and April 16, respectively) and Claude Sonnet 4.6 (released February 17), which established new performance benchmarks in agentic reasoning and improved “Computer Use” reliability.

Figure 22. Claude vs. OpenAI – First AI Check



Source: Ramp presentation

The primary commercial engine behind this spike has been Claude Code, which saw its specific revenue contribution double to $2.5B since January as enterprises shifted from experimental

chat interfaces to autonomous engineering agents. Data from Ramp (Private) shown below indicates the inflection in enterprise spending going to Anthropic since the beginning of this year.

Copilot Cowork: The Microsoft-Anthropic Integration

Announced on March 9, 2026**,** Copilot Cowork represents a strategic deep-level integration where Microsoft is bringing Anthropic's "Cowork" agent technology directly into the Microsoft 365 Copilot ecosystem. Claude models will now be available within Copilot and Copilot Cowork will run within M365's security and governance framework increasing its distribution potential across 450M+ Microsoft users. This offering shifts the focus from conversational chatbots to "action-oriented" autonomous agents capable of managing multi-step workflows across the entire M365 suite. The tool can autonomously move from "intent to action," coordinating tasks between Excel, Outlook, and Teams without human handholding for every step.

Mythos: The Next Stage in Cybersecurity and Government Relationship?

Announced in early April, Anthropic introduced its latest frontier model, Mythos, but elected to defer a broad public release due to concerns around emergent cybersecurity and offensive capabilities that management believes could pose public safety risks. Instead, Anthropic launched a limited preview through its Project Glasswing initiative, granting select enterprise and infrastructure partners early access intended to provide a controlled “grace period” to identify and remediate latent vulnerabilities ahead of broader availability. According to the company, Mythos has already identified thousands of high-severity vulnerabilities across every major operating system and web browser during preview testing, underscoring both the defensive potential and the systemic risks associated with frontier-scale, AI-driven security tooling.

After butting heads with the US government earlier in the year over the use of their AI models, we believe that recent attention around Mythos’s security capabilities could reopen the door for friendlier relations. President Trump has noted that the company and the government have recently hosted more productive talks. That said, the DoD has not dropped its claims against the company, and Anthropic’s lawsuit over the administration’s policies remain outstanding.

Cloud Infrastructure and Gigawatt Capacity

From an infrastructure standpoint, Anthropic has pursued a diversified cloud strategy anchored by multi-billion-dollar partnerships with Amazon Web Services (AWS), Microsoft Azure, and Google Cloud. Unlike OpenAI’s primary dependence on Microsoft, Anthropic distributes workloads across multiple hyperscalers to support its rapidly expanding training and inference requirements.

As adoption of the Claude model suite has accelerated, the company has increasingly encountered constraints related to the availability of compute capacity. In late 2025, Anthropic materially expanded its use of Google Cloud TPUs, citing superior price-performance characteristics for training frontier-scale models. More recently, in early April, Anthropic signed

an agreement securing an additional 5 gigawatts of AWS cloud capacity and another deal with CoreWeave (CRWV, Hold, Ramnani) to alleviate near-term supply-side constraints.

In aggregate, this infrastructure expansion is expected to bring well over 1 gigawatt of net new compute capacity online by late 2026, representing the largest build-out of AI-specific infrastructure in the industry to support large-scale Claude inference and agentic workloads.

Recent Acquisitions Bring Focus on Governance, Security, and Agent Testing

Anthropic's recent M&A activity underscores that the next frontier for enterprise AI is not just intelligence, but the ability to govern and validate agentic systems at scale.

Acquisition	Date	Strategic Rationale
Vercept	Feb 2026	Focuses on agent governance and auditability for "computer-use" automation. By controlling the layer that determines how Claude interacts with OS-level software, Anthropic ensures that agents remain within secure boundaries while performing complex tasks.
Bun	Dec 2025	Strengthens the JavaScript/TypeScript developer stack. As the fastest JS runtime, Bun allows for the vertical integration of the coding pipeline, significantly increasing the execution speed and efficiency of web-based agents built on Claude Code.
Humanloop	Aug 2025	Enhances the evaluation and optimization loops through human-in-the-loop feedback. This acquisition allows enterprises to systematically refine agent performance and validate safety metrics before deployment.

Source: Anthropic company reports

By bringing these capabilities in-house, Anthropic is embedding auditability and traceability into the developer stack. The strategic focus remains clear: enterprise adoption will not be limited by what agents *can* do, but by the organization's ability to systematically prove they are doing it safely and accurately.

Application Software and Agentic AI - Tillman

Microsoft – From Copilots to the Enterprise Control Panel

***Summary**: Microsoft's agentic AI strategy has evolved from embedded copilots within core workflows to increasingly autonomous systems capable of executing multi-step tasks across enterprise environments. Rather than competing purely at the model layer, the company is positioning itself as the full-stack platform where agents are built, grounded in data, and governed at scale. This approach places Microsoft at the center of enterprise AI adoption, with agents emerging as the next execution layer across productivity, business applications, and development workflows. Also with massive reach, reflected in over 1 million organizations using its productivity suites and over 450 million M365 Commercial seats, there is inherently large incumbent distribution opportunity with agentic solutions, in our opinion.*

Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	Agent infrastructure, platform/orchestration layer, the application/inference layer, proprietary model offering
Role	Horizontal & vertical AI infrastructure provider
User	Knowledge workers, engineers/developers, business operators, IT/enterprise architects, and data analytics professionals

Financial Snapshot ($Ms)					
Mkt. Cap	**$3,153,071**	**Net Debt**	**$17,619**	**Enterprise Value**	**$3,170,690**

Revenue	As of 2QFY26, Microsoft Cloud revenue was $51.5 billion, up 26% y/y while Intelligent Cloud revenue was $32.9 billion, up 29% y/y. Azure revenue grew 39% y/y (38% in constant currency.) Commercial RPO grew 110% y/y to $625 billion with ~45% of that being attributed to OpenAI contracts alone.
Most Recent Performance Updates & LT Targets	Copilot monthly active users (MAUs) have surpassed 100 million with over 800 million users engaging with AI features across the Microsoft ecosystem. Nearly 70% of the Fortune 500 now use Microsoft 365 Copilot, and 90% of the Fortune 100 use GitHub Copilot Over 100,000 organizations have used Copilot Studio to build or extend agents, and customers have created 3 million agents via SharePoint and Copilot Studio Microsoft's AI business is on track to surpass an annual revenue run rate of $10 billion next quarter, marking the fastest business milestone in company history

Source: FactSet – Financial Data as of 4/27/2026

Microsoft's AI Framework: From Embedded Copilots to the Enterprise Execution Layer

Microsoft's approach to agentic AI has followed a mostly structured progression, in our opinion, beginning with assistive copilots embedded inside existing workflows and evolving toward increasingly autonomous systems capable of executing tasks on behalf of users. Rather than launching standalone AI products, Microsoft has consistently anchored its strategy in its core assets including productivity software, developer tools, and enterprise applications, using these surfaces as the distribution layer for AI.

Today, we believe Microsoft is positioning itself not as a pure model provider, but as the enterprise platform through which AI agents are built, grounded in data, deployed, and governed. Its current architecture spans user-facing copilots, low-code and pro-code agent builders, and a managed runtime environment, forming what is effectively a full-stack system for agentic AI. The company's direction suggests a transition from AI as a productivity enhancement tool to AI as an execution layer embedded across enterprise software.

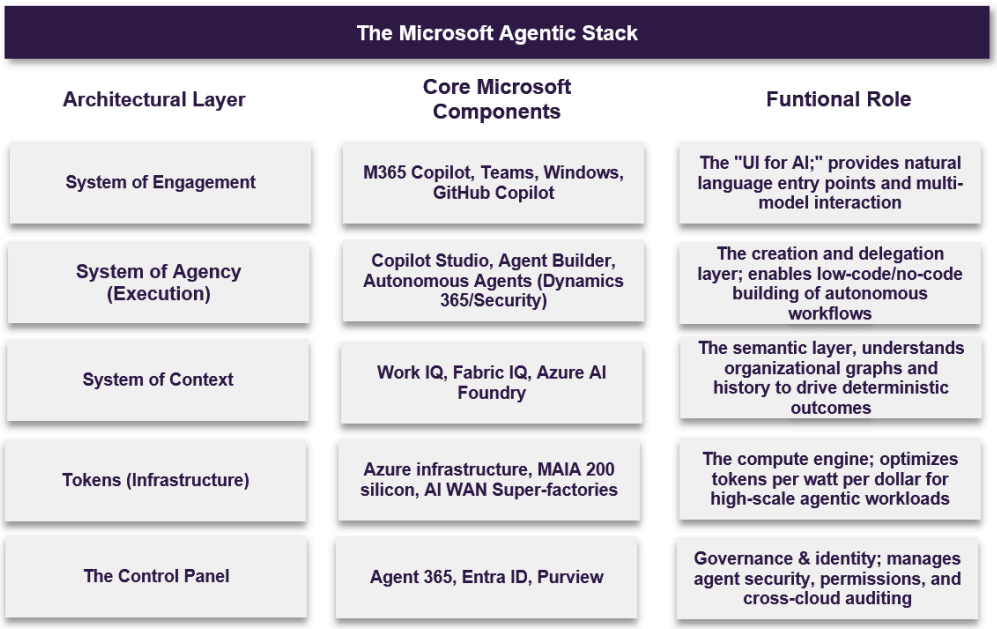
Microsoft's entry into generative AI at scale began with GitHub Copilot, which introduced AI directly into the developer workflow as an assistive tool. The core concept, AI embedded within an existing system of work, became the foundation for broader deployment and use cases. This expanded into Microsoft 365 Copilot and Dynamics 365 Copilot, where AI was integrated into productivity and business applications and grounded in organizational data such as emails, documents, meetings, and CRM records.

At this stage, Microsoft's AI capabilities were primarily assistive. The systems were designed to summarize, draft, retrieve information, and provide recommendations, but they remained user directed. Importantly, Microsoft's public positioning emphasized that these copilots were not standalone chatbots, but tools that operated within the context of enterprise data and permissions. This early focus on grounding and integration laid the foundation for more advanced agentic capabilities we have seen more recently.

Microsoft's shift toward agentic AI became explicit with the introduction of autonomous agents through Copilot Studio and the expansion of agent capabilities within business applications. These agents were designed to move beyond assisting users to executing defined workflows, often triggered by events, schedules, or system changes. In parallel, Microsoft introduced Azure AI Foundry, which provides a unified environment for building, deploying, and managing AI applications and agents with enterprise-grade controls.

We believe this phase marked a structural change in Microsoft's AI strategy. The focus shifted from enhancing individual productivity to automating repeatable business processes across functions such as sales, service, and finance. Agents were no longer limited to responding to prompts as they could take actions, integrate with external systems, and operate with a degree of autonomy within predefined boundaries.

Figure 23. Reasoning Agents and Multi-Layer Architecture

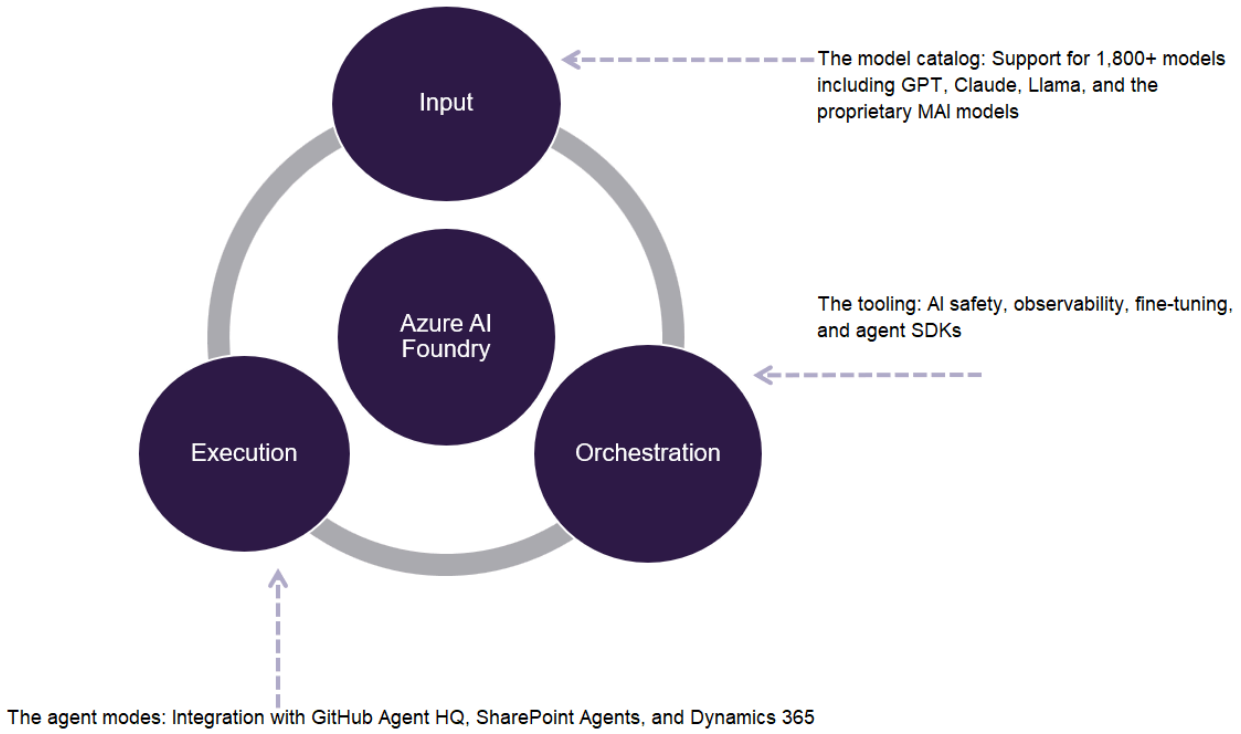


Source: Microsoft website

Microsoft's most recent developments introduce reasoning-based agents capable of multi-step problem solving and iterative execution. These systems are designed to break down complex tasks, gather relevant data, and refine outputs over multiple steps, representing what we believe is a step change from earlier prompt-response paradigms. At the same time, Microsoft has extended agentic functionality into software development through GitHub Copilot, where agent modes can complete multi-step coding tasks and operate with increasing independence.

The company's current agentic AI stack can be understood as a layered architecture. At the surface level, copilots act as the user interface, enabling natural language interaction with enterprise systems. Beneath this, Copilot Studio enables the creation of custom agents and workflows using enterprise data and integrations. At the developer level, Foundry provides the tools and infrastructure needed to build and scale production-grade agents. Surrounding these layers is a broader framework for data integration, orchestration, and governance, ensuring that agents operate within enterprise security and compliance requirements.

Figure 24. Azure AI Foundry – Key to Owning the Orchestration Layer



Source: Azure AI Blog

We believe Microsoft's differentiation lies in its ability to integrate agentic AI across the full enterprise stack rather than singular components such as core models, reasoning, memory/storage, orchestration, etc. Its distribution advantage stems from its presence in widely used productivity and business applications, allowing it to introduce agents directly into daily workflows. At the same time, its control over enterprise data and identity systems enables agents to operate with context and within defined permission structures. We believe this could be especially helpful in securing more complex enterprise workflows, especially in settings where customers are using Microsoft solutions extensively across compute, data, security and app layers.

The company's public messaging increasingly emphasizes a multi-model approach, where its platforms support a range of underlying AI models. We believe this suggests that Microsoft's strategic objective is not to compete solely on model performance, but to own the orchestration, deployment, and governance layers that determine how AI is actually used within organizations. In this framework, Copilot becomes the interface, agents become the execution layer, and platforms like Foundry provide the infrastructure that binds them together. Also, this agnostic model approach could also benefit the company, especially if AI models become commoditized when looking at in isolation versus the broader set of architecture components required with agents.

Monetization: The Shift from Seats to Digital Labor

We believe Microsoft is transitioning its monetization strategy from a traditional, purely seat-based SaaS model toward a blended consumption framework that prices "Digital Labor" alongside human productivity. While core access remains anchored in tiered subscriptions, such as the $30 per-user monthly fee for Microsoft 365 Copilot, autonomous agentic workflows are increasingly governed by a usage-based meter. Central to this shift is the introduction of Azure Agent Units (AAUs) and Copilot Credits, which act as the primary currency for agentic execution. For instance, the Azure SRE Agent is billed at a baseline of 4 AAUs per hour, supplemented by an additional 0.25 units per task executed, effectively decoupling revenue from human headcount and tying it directly to autonomous output. Microsoft CEO Satya Nadella has signaled that this "Subscriptions + Meter" model is the destination for both coding and information work, allowing enterprises to manage the high token intensity of long-running tasks while providing Microsoft with a TAM-expansive lever as agents begin to outnumber human users.

We also believe the deployment of "hero" agents across functional and vertical domains has moved Microsoft beyond the pilot phase into a state of production-ready autonomous labor, with the company noting that over 80% of the Fortune 500 is now running active agents. In functional roles, Microsoft's Security Agents have become a primary proof point for high-value automation, in our opinion, with 11 specialized agents handling tasks like phishing triage and identity management. Notably, Microsoft has highlighted that Icertis (Private) reported a 75% reduction in manual triage time using these tools. Similarly, the company notes that GitHub Copilot has evolved from simple code completion into an "Agent Mode" and "Agent HQ" environment that asynchronously manages pull requests and code base synchronization, currently supporting 4.7 million paid subscribers.

In vertical sectors, Microsoft has highlighted that Dragon Copilot leads the healthcare market by documenting 21 million patient encounters per quarter, while the Microsoft Discovery suite orchestrates R&D agents to formulate scientific hypotheses and run simulations for leaders like Unilever. Internally, Microsoft has noted that it serves as "Customer Zero," utilizing its own sales and support agents to save "hundreds of millions of dollars annually" by resolving customer issues without human intervention and automating back-office minutiae for its global sales force. We believe this results-oriented approach, moving from assistants to autonomous digital workers, is designed to drive a structural shift in business process costs, allowing firms to capture knowledge within models rather than human capital alone.

Bringing it all together on monetization, Microsoft has positioned its newly announced Microsoft 365 E7 Frontier Suite as its most comprehensive commercial bundle for enterprise AI, packaging together core productivity applications, advanced security, compliance, and Copilot capabilities into a single, integrated offering. Based on Microsoft's public materials, the value of E7 is not just incremental features, but the unification of AI with enterprise data, identity, and governance, allowing organizations to deploy Copilot and agents within a fully permissioned and compliant environment. This matters because Microsoft consistently emphasizes that effective AI deployment in the enterprise requires more than model access, it requires secure data

grounding, policy enforcement, and system-wide integration. By bundling these elements into a higher-tier subscription, we believe Microsoft could more effectively monetize agentic AI through a premium, all-in-one SKU that aligns AI adoption with broader upgrades across its software and security stack, reinforcing both pricing power and ecosystem lock-in. Microsoft 365 E7 will go GA on May 1, 2026, and will be priced on a per user basis at $99 per month.

OpenAI Partnership: Frenemies?

We view the relationship between Microsoft and OpenAI as a $135 billion strategic alliance that has more recently evolved from a pure research partnership into a complex "co-opetition" model. Following OpenAI's October 2025 recapitalization into a Public Benefit Corporation (PBC), Microsoft converted its previous investment into a 27% equity stake. More recently, Microsoft and OpenAI announced a definitive agreement to amend their partnership further on April 27, 2026, signaling what we believe is a transition from an exclusive alliance to a more flexible, standardized enterprise relationship.

We see the April 2026 update as a reflection of the maturing AI market where both companies are seeking greater independence as OpenAI looks to optimize its compute costs across multiple clouds while Microsoft further mitigates against single-model risk as its internal MAI (proprietary models) strategy gains parity with other frontier models.

Figure 25. Microsoft & OpenAI Partnership Terms

Key Component	Details & Financial Impact
Equity Stake	Microsoft holds a 27% stake in OpenAI Group PBC (valued at $135B)
Unrealized Gains	Media reported cap tables suggest an 18x return on initial capital, with the stake now worth ~$228B
Azure Commitment	OpenAI contracted for an incremental $250B of Azure services - OpenAI products will ship first on Azure unless Microsoft lacks necessary capability
IP Licensing	Non-exclusive through 2032 - ends Microsoft's unique monopoly on OpenAI's frontier tech and allows OpenAI to license to others
Cloud Hosting	Non-Exclusive, OpenAI can use any provider like AWS, GCP, etc. but Microsoft still remains the primary partner
Revenue Share Microsoft - OpenAI	Terminated - Removes Microsoft's ongoing royalty obligation to OpenAI for integrated products
Revenue Share OpenAI - Microsoft	Capped through 2030 - Provides OpenAI with LT margin clarity while securing a fixed return for Microsoft

Source: Microsoft Company Website

Note: Cap table commentary reported on by Forbes

Under the new terms of the agreement, we frame the partnership between OpenAI and Microsoft as a pragmatic truce where both companies gain important pieces to the development of their AI/agentic stacks. By shifting to a non-exclusive license and allowing OpenAI to utilize other cloud providers, we believe Microsoft is effectively reducing complexity and the level of reliance on OpenAI just as its own MAI-1 foundational models reach production readiness.

- **For Microsoft:** We believe the removal of the revenue share payment to OpenAI is a significant win for gross margins within the M365 and Azure segments. Further, by conceding exclusivity, we believe Microsoft positions Azure as a neutral, high scale token factory that can be used by any frontier model, in-line with Microsoft's recent messaging around being model agnostic across its platform.
- **For OpenAI:** We see the ability to serve products across any cloud provider gives the company infrastructure leverage. OpenAI now has the freedom to go to many hyperscalers for compute and can negotiate the best cost per token terms. We think this will become increasingly important as OpenAI continues to spend aggressively on R&D costs to effectively compete against open source and other proprietary models in the race to AGI.
- **Revenue Cap:** We see the cap on OpenAI's payments to Microsoft through 2030 as providing a valuation floor for OpenAI's upcoming capital rounds while ensuring that Microsoft's $135 billion investment remains productive without becoming a drag on OpenAI's eventual IPO or further capital raises.

We believe the updated Microsoft & OpenAI partnership is critical to Microsoft's agentic strategy because it shifts the relationship from exclusive control to platform leverage at scale. While Microsoft remains OpenAI's primary cloud and retains long-term IP access, the move toward a multi-cloud, non-exclusive structure allows OpenAI models to proliferate across the broader ecosystem, expanding the total surface area for agent deployment, while still anchoring first-party innovation on Azure and Copilot. We believe this effectively positions Microsoft as the default control plane for enterprise AI (via Azure, Copilot, and its agent frameworks) rather than the sole distribution channel, which is strategically more valuable in an agentic world where interoperability, orchestration, and governance matter more than exclusivity. We frame the next chapter of the partnership as Microsoft trading some exclusivity for ecosystem dominance, ensuring its stack sits at the center of agent execution, security, and workflow integration, even as models become more commoditized and broadly accessible.

Diversifying the Stack: Anthropic Partnership Reduces Single-Model Dependence

In late 2025, Microsoft announced a partnership with Anthropic that we believe represents a meaningful expansion of its agentic AI ecosystem combining model access, infrastructure, and product integration. As part of the initial agreement, Anthropic has committed to purchasing ~$30 billion of Azure compute capacity over time, positioning Microsoft as a key infrastructure provider for training and running Claude models, in our opinion. In parallel, Microsoft has integrated Claude across its AI stack, including development platforms and user-facing applications, enabling enterprises to build and deploy agents using multiple frontier models within a single environment. This extends into workflow-level integration through Copilot and "Cowork" style experiences, where agents powered by Claude can assist with tasks such as research, content creation, and coordination directly within productivity tools.

Strategically, we believe the partnership underscores Microsoft's shift toward a multi-model, platform-centric approach to AI and agents. By incorporating Anthropic's models alongside its existing ecosystem, we believe Microsoft is actively diversifying its exposure beyond OpenAI

and reinforcing its role as a neutral control plane for enterprise AI and agents. Rather than anchoring its strategy to a single model provider, we believe Microsoft is positioning its platform as the layer where different models converge, are orchestrated, and are deployed into real enterprise workflows, strengthening both its flexibility and its long-term ability to capture value across the agentic stack.

Importantly, we believe the most recently announced chapter of the partnership from March, Microsoft's introduction of Copilot Cowork, represents a meaningful step forward in its agentic AI strategy, shifting from assistive experiences toward systems that can execute work over time. As described by Microsoft, Claude Cowork enables users to delegate outcomes rather than individual tasks, with the system planning, coordinating, and carrying out multi-step workflows across Microsoft 365 applications while keeping users in control of the process. This capability is powered by multi-model architecture and deep integration with enterprise context through Work IQ, allowing agents to reason across a user's full work environment rather than isolated data points. Strategically, we believe the integration of Cowork reinforces Microsoft's positioning of Copilot as the execution layer for enterprise AI, demonstrating how multi-model partnerships, such as its collaboration with Anthropic, translate into production-level agentic workflows embedded directly within everyday business applications.

Salesforce – Operating System for the Agentic Enterprise

Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	While interoperating with widely adopted LLMs, someproprietary models for specific use cases; emphasizes four layers in agentic workloads - data layer, agent building layer (Agentforce), engagement layer (Slack & Slackbot) and trust layer
Role	Horizontal and vertical AI infrastructure provider
User	Developers, forward deployed engineers, enterprises, SMBs, AI startups and anyone building, deploying and managing agents

Financial Snapshot ($Ms)					
Mkt. Cap	**$145,745**	**Net Debt**	**$8,146**	**Enterprise Value**	**$153,891**
Revenue			Total FY26 revenue was $39.4 billion, and as of the end of FY26, Agentforce & Data 360 ARR >$2.9 bil (includes Infromatica) and Agentforce ARR $800 mil (up 169% y/y)		
Most Recent Performance Updates & LT Targets			Closed more than 29,000 Agentforce deals since Fall '24 launch and up 50% q/q, more than 19 trillion tokens processed to date (up 5x y/y), Agentforce accounts in production increased almost 50% q/q, and 2.4 billion Agentic Work Units (AWUs) delivered to date. Committed to organic growth reacceleration beginning in 2HFY27, $63 bil revenue target in FY30 (11% organic growth from FY26-FY30) and rule of 50 by end of FY30.		

Source: FactSet – Financial Data as of 4/27/2026

Agentforce Technical Underpinnings – Data, Intelligence, Actions and Security

Since the unveiling of Agentforce in September 2024, Salesforce has been committed to delivering on fully autonomous digital workers (agents) that progress beyond AI copilots, assistants and FAQ-oriented chatbots to more sophisticated autonomous agents handling complex multi-step business processes that augment or replace existing processes and human labor while delivering on key business outcomes.

A core foundational element of Agentforce is an intelligence layer where the Atlas Reasoning Engine acts as the “brain” that plans, reasons and self-reflects during the problem-solving phase in order to optimize the determinism of agents and their actions. Highly deterministic agents are especially important when dealing with mission critical business processes, high-value interactions or engagement and/or regulated industry workflows. Predictable, accurate and highly performative outcomes are essential for enterprise agent success and adoption;

probabilistic AI will often simply not cut it for agent adoption to fully take off in the enterprise market, in our opinion.

What provides unique positioning for Salesforce is its strategic position related to vast amounts of highly contextual structured and unstructured data that reside across its large set of apps businesses, along with data and security capabilities that help deliver on key requirements around trust, security and governance. Data 360 is characterized by Salesforce as the foundation for Agentforce deployments and providing a lone source of truth for agents. Data 360 includes capabilities for harmonizing, integrating and federating data in order to drive intelligence, accuracy and reliability while minimizing hallucinations via proper context. Data 360 becomes the key data activation layer for Agentforce, and its Zero Copy capability allows customers to read data residing in third-party data lakes without physically moving or copying the data. This provides real-time data updates for agents to act upon while reducing data latency.

Also helping drive agent actions and orchestration are Salesforce's MuleSoft capabilities, which are important in executing multi-agent workflows and connecting to external APIs/systems. MuleSoft leverages agent-to-agent communication protocols (e.g., A2A, MCP) as well as agent brokers that help orchestrate tasks and processes across third-party systems.

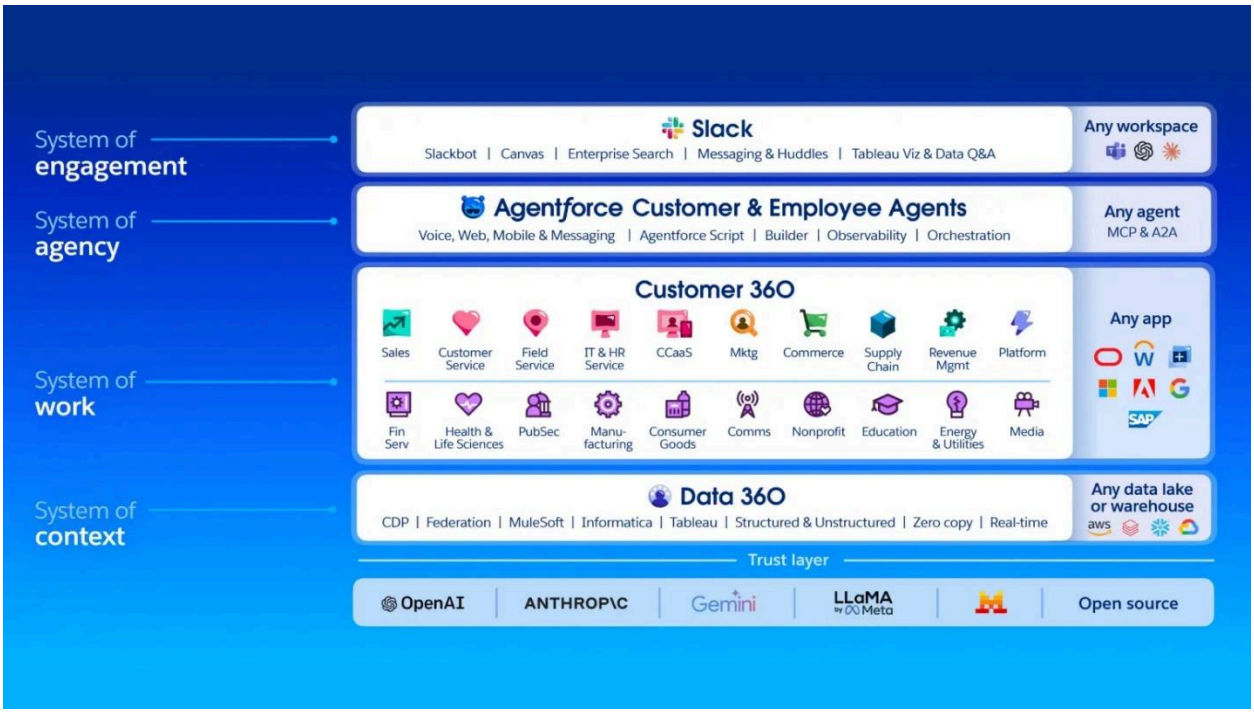
Another key part of the Salesforce agentic operating system is its capabilities within the Trust Layer. This includes integrating and interoperating with leading commercial and open source LLMs as part of agentic workloads while also managing PII masking, data obfuscation (e.g., encryption, tokenization, etc.) and zero-retention policies that prohibit models from storing or using customer data for training purposes.

Below is a deeper dive into the Salesforce's capabilities that help create differentiation with enterprise agents, in our opinion. It entails system of engagement capabilities, including Slackbot, that could invoke a variety of AI use cases, including agents, directly in Slack. System of agency capabilities include core agent builder functionality that can handle the needs of non-technical users up to advanced programmers via a mix of no code, low code and pro-code methodologies as well as Agentforce Script, observability and orchestration capabilities. To enable long-running autonomous agents that work across entire customer journeys, the system of agency layer incorporates numerous agent modalities including voice, web, mobile and messaging. Other key layers that play a valuable role include system of work components (e.g., specific horizontal and vertical apps and related tasks/workflows) as well as system of context, which we discussed above in regard to Data 360, Atlas Reasoning Engine, MuleSoft, etc.

It is important to separately highlight new Salesforce agentic advances that were delivered recently at Salesforce's developer conference TDX 2026. A key highlight was the launch of Headless 360, which provides developers and others building agents increased flexibility to build agents directly within the Salesforce ecosystem or from terminals using CLI commands, directly from a browser or within coding environments such as Claude Code, Cursor, Codex, etc. Also the company unveiled the AgentForce Experience Layer, which separates agent functional capabilities from its presentation layer, along with Agent Exchange, which provides a

centralized source for discovering, activating and managing out-of-the-box agents from Salesforce as well as third parties.

Figure 26. Salesforce Systems Stack



Source: Salesforce investor presentation

While the company aimed to provide simplicity with its original $2 per conversation pricing model in late 2024, they have since developed a broader set of packaging options to maximize flexibility and price/value dynamics. More recent and widely adopted pricing and packaging options have included pay-as-you-go Flex Credits per agent action (e.g. updating records, summarizing cases), Flex Agreements (prepaying for credits and being able to shift between user licenses or digital labor), standard and premium add-ons with fixed per user per month pricing (potentially relevant for B2B employee-facing use cases and/or regulated industries) and Agentic Enterprise License Agreements (AELAs).

AELAs have typically been geared towards larger organizations that could entail well into the seven-figure+ investment range and handling upwards of millions of agent interactions. AELAs can provide flexibility and budget transparency given flat rate, unlimited agent usage associated with enterprise license agreements. To measure success of Agentforce, we will be focusing on the mix of Agentforce customers that are consistently coming back to consume more Flex Credits under the pay-as-you-go model.

Also, we expect to hear more about the concept of Agentic Work Units (AWUs) by Salesforce that reflect business value being created with agents via measurable tasks completed versus simply sharing eye-popping token consumption stats that may or may not indicate value realization or business outcomes.

Figure 27. Enterprise Work Statistics: Agentic Work Units

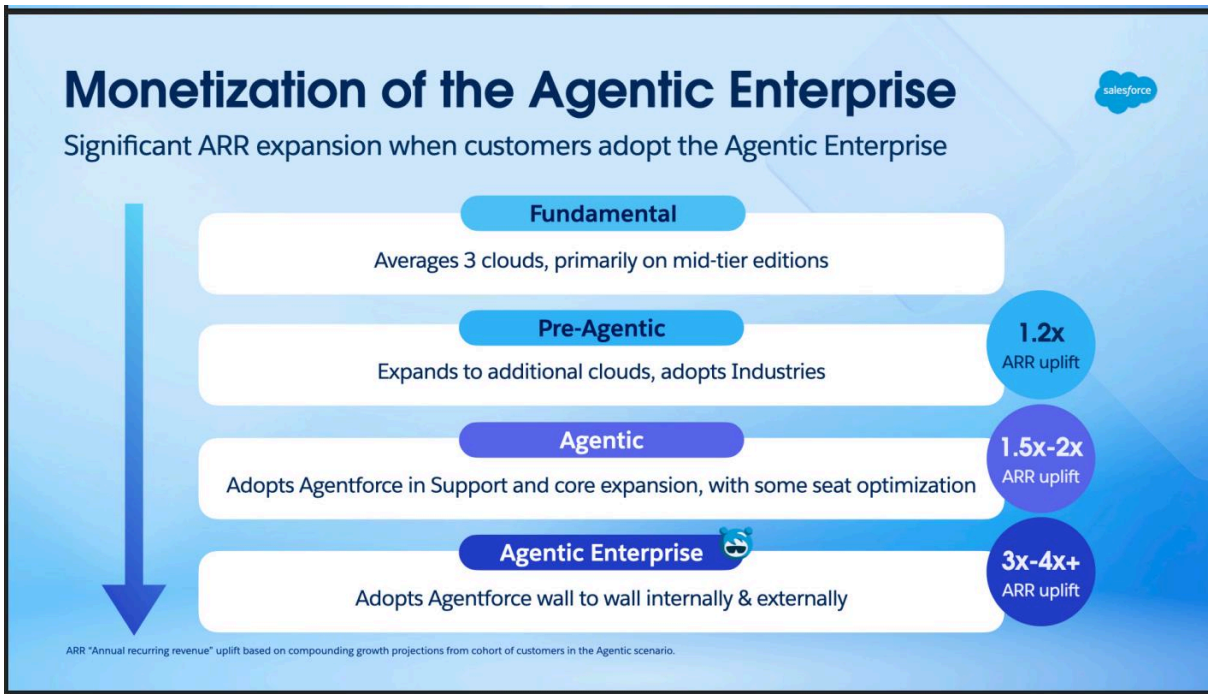


Source: Salesforce investor presentation

Related to digital labor value creation, we expect to hear more about the concept of “hero” agents that relate to widely adopted agent use cases in areas such as Agentforce Sales and Agentforce Service and where value can be measured and priced on a simple measure, such as leads generated by an Agentforce sales agent. This represents a broader development we expect to play out with enterprise agents across the industry whereby outcome-based monetization becomes a more notable approach to pricing and packaging. We believe such outcome-based monetization could be especially advantageous to SoR companies that benefit from vast amounts of highly contextual customer and/or operational data and for which highly deterministic agent outcomes might consistently be achieved.

Salesforce will need to continue to deliver on a growing base of successful enterprise agent case studies while progressing customers from pilots to production deployments. Increased experience and knowledge transfer among Salesforce’s forward deployed engineers and its partner ecosystem will also be relevant in driving greater Agentforce flywheel benefits. While Salesforce has consistently shared benefits it has achieved as customer zero for agentic use cases (e.g. 1.8 million+ support conversations automated with 77% resolution rate, 300,000+ engineer hours saved annually, etc.), we believe additional ROI reference points will be key to validating agentic AI across more extensive set of use cases than where we have seen most of the revenue scaling in the industry to date around coding agents.

Figure 28. Salesforce Monetizing Agents



Source: Salesforce investor presentation

Salesforce has talked about a potential 3x-4x+ ARR uplift for customers who move from traditional seat-centric CRM usage model to what it characterizes as a complete agentic enterprise model. This entails an assumed combination of premium-seat based subscriptions (often higher-tiered products provide access to advanced AI capabilities) and consumption-based credits and could lead to less reliance going forward on customer headcount levels and seat expansion for growth.

Infrastructure Software - Jump

Snowflake – Owning the Data Layer in the Age of Agents

Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	Data layer / governed analytical and operational data platform
Role	System of Record/Source of Truth
User	Enterprises, (Data Engineers, Analysts)

Financial Snapshot ($Ms)					
Mkt. Cap	**$48,509**	**Net Debt**	**($1,289)**	**Enterprise Value**	**$47,220**

Revenue	Snowflake exited FY26 with total revenue of $4.68B (+29.2% YoY), underpinned by a significant acceleration in RPO to $9.77B (+42% YoY) .
Most Recent Performance Updates & LT Targets	AI penetration is exemplified by 9,100 accounts using AI, over 4,400 customers using Cortex Code, and Snowflake Intelligence doubling its adoption to over 2,500 accounts in just three months.

Source: FactSet – Financial Data as of 4/27/2026

From Data Custodian to Agent Enabler

In 2025, we believed that Snowflake had been a core beneficiary of the “getting your data ready for AI” theme. Now, we see a compelling path for Snowflake to become a primary destination for the initial buildout of agentic workflows. We believe the architectures that win the next stage of AI will be those with close proximity to governed, high-quality data. Snowflake already occupies that position across much of the enterprise landscape as the centralized, governed store for enterprise data. For many customers, core analytical and operational datasets already flow into Snowflake, creating a single environment with unified controls and minimizing the need to move sensitive data into new tools. This consolidation materially reduces deployment friction for any agentic system that must reason over high-quality, trustworthy information.

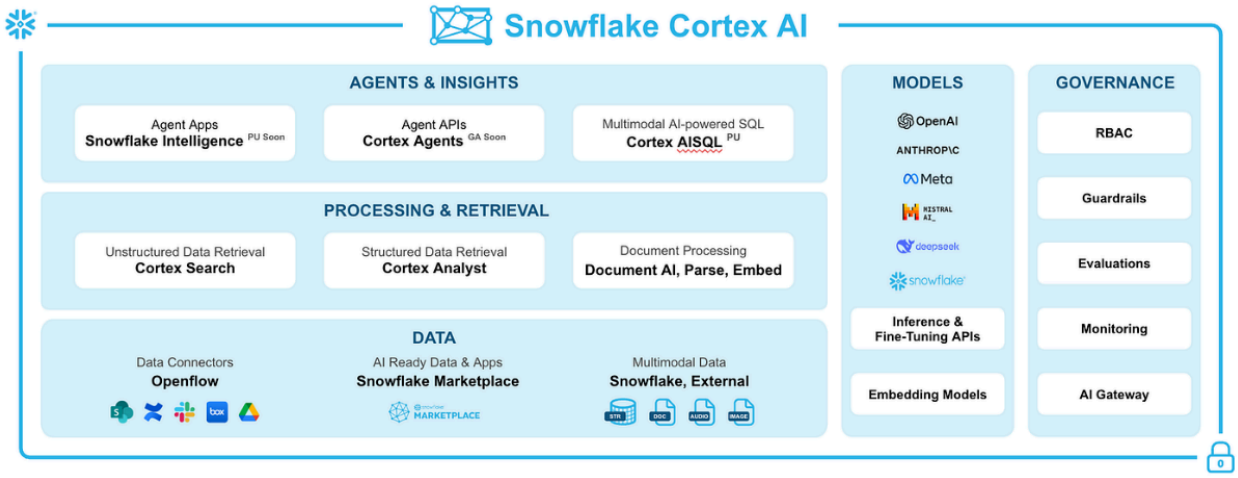
“Path-of-Least-Resistance” Advantage

Enterprises building agents tend to optimize for two attributes: direct access to enterprise data and minimal friction when integrating new agentic components into existing operational systems. Snowflake already sits at the center of many customer data pipelines and identity frameworks, which reduces the approval, integration, and compliance overhead associated with introducing new AI workflows into the platform. As data gravity continues to shift toward Snowflake, application gravity follows; workloads increasingly move to where governed data already resides. Together, these dynamics reinforce a powerful flywheel that naturally positions Snowflake as the execution layer for emerging agentic systems, while alternatives typically require replicating sensitive data into new environments or stitching together multiple services, both of which increase governance risk and slow deployment.

AI Product Suite Spinning Consumption Flywheel at the Engagement Layer

Over the last year, Snowflake has introduced a series of AI features designed to reduce friction in tool construction and drive platform consumption. Core to the strategy are Snowflake Intelligence and Cortex Code (CoCo). Intelligence converts natural-language inputs into performant, governed queries and is seeing the fastest adoption ramp of any product in Snowflake's history, with >1,200 WAUs. CoCo provides the orchestration layer required to chain model calls, apply decision logic, and interact with external applications, and is used by >4,400 customers to accelerate the development of production-grade AI applications. Together, these capabilities allow Snowflake to support reasoning-heavy workloads natively while lowering the barrier to entry for both technical and non-technical users. Importantly, early traction across both products has helped push Snowflake's AI portfolio to a $100M revenue run-rate, a milestone achieved a full quarter ahead of expectations and grounded in real production consumption. This positions Snowflake as a practical environment to experiment with and operationalize early agentic workflows at enterprise scale.

Figure 29. Snowflake Cortex Suite



Source: Snowflake presentation

Strategic Takeaway

Net/net, we see Snowflake positioned to benefit on two fronts. First, Snowflake is likely to become a "source of truth" for agentic workflows that require reasoning over governed enterprise data. Second, Snowflake's familiarity, governance model, and simplifying AI capabilities create a strong "path-of-least-resistance" narrative as enterprises evaluate where to build their first-generation agents. Competitive pressure will persist, and Snowflake must broaden the market's understanding of what its platform can do, but we believe these advantages materially increase the likelihood that Snowflake captures a disproportionate share of early agentic workloads.

ServiceNow – Enterprise System of Action in an Agentic World

Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	Workflow, orchestration, and control plane layer
Role	Coordination, orchestration, execution, & enforcement layer for agent work
User	IT Ops, Digital Workflow Operators, Enterprise Agent Deployment Teams

Financial Snapshot ($Ms)					
Mkt. Cap	**$92,993**	**Net Debt**	**($3,887)**	**Enterprise Value**	**$89,106**
Revenue			Quarterly capital expenditures reached $238.0MM in late 2025, while the Now Assist GenAI platform surpassed $600MM in ACV and is tracking toward a $1B+ target for 2026.		
Most Recent Performance Updates & LT Targets			10x YoY increase in customers using five or more Now Assist products and a 70% upsell expansion observed in Customer Service Now Assist renewals		

Source: FactSet – Financial Data as of 4/27/2026

Counted Out Too Soon?

Over the last year, ServiceNow has been one of the more visible casualties of the 'SaaSpocalypse' in our coverage, with investors' terminal-value concerns on growth drivers and margin potential coming to a head after the company's accelerated acquisition cadence. But despite the challenging tape and the optics around M&A, we believe the underlying business has the positioning to sustain (and grow) its relevance as agentic AI proliferates across global enterprises. That said, we see a longer-term opportunity for ServiceNow to claw its way out of the "AI losers" bucket as its agentic AI strategy propagates through the installed base and early deployments convert into platform expansion opportunities.

Multi-Layered Blueprint for Agentic Business

NOW's "Blueprint for Agentic Business" is predicated on the belief that enterprise agents will only be as good as the context they operate in, an area where the company feels it possesses a structural lead. Their argument is that the real moat isn't the model layer, which is already getting cheaper and more interchangeable, but the operational logic, relationships, and governance embedded in the workflows customers run on NOW every day. We believe the platform's data model, history of executed workflows, and unified governance layer give agents the guardrails they need to behave predictably, and the deterministic workflows underneath give them a reliable way to complete work across IT, HR, CSM, and third-party systems. In our view, this execution-first philosophy gives a compelling narrative for ServiceNow to resonate with customers moving from experimentation to operational AI.

In the long term, we see the potential for AI Control Tower, what the company refers to as the "air traffic control" for identity, entitlements, policy, audit, and data residency across models and clouds, to inflect as the context thesis plays out. Over the long term, autonomous workflows and

AI Workers extend the monetization surface from assistance to governed action across IT first, then HR and finance, domains where approvals, segregation of duties, and evidence trails are non-negotiable. The company's blueprint positions these capabilities as native to the platform rather than bolt-ons, a distinction we view as difficult for model- or data-layer vendors to replicate without owning the execution layer.

Figure 30. ServiceNow Agentic Suite



Source: ServiceNow presentation

Fieldwork Indicates Distribution Advantage for Incumbents as AI Partners

Our fieldwork at the ServiceNow AI Summit in Boston and follow-up with product leadership reinforce this execution-first thesis. Customers across healthcare, financials, education, federal, and IT echoed that their 2026 priority was going from "0-to-1" on AI with ServiceNow as a core partner, with no indications in our conversations of pulling back platform spend as part of AI buildouts. Partners echoed rising demand to get started, with less uniformity around truly mature deployments, a pattern consistent with an early-innings cycle where assistive use cases lead and autonomous workflows follow. Feedback also highlighted both HR service delivery consolidation onto NOW and a premium on operating in regulated environments of choice with strong security and governance, a place where AI Control Tower's unified policy/audit posture should resonate as agent fleets and model diversity expand.

We also see the potential for Pro Plus / Now Assist upsell as a clean attach that could drive upside to our estimates in the year to come. Pro Plus, which is a ~30% upsell, is less than 10% penetrated into the pro/enterprise customer base, a figure we see growing as enterprises early on the adoption curve begin their AI journeys with "out-of-the-box" solutions, leaving a multi-year runway as the installed base moves from pilots to standardization. The model is attractive

because it captures both attach-driven uplift and consumption-based usage growth, allowing customers to realize value inside governed workflows they already trust."

Strategic Takeaway

Despite a disappointing guide coming out of 1Q26, which we attribute more to geopolitical uncertainty than AI headwinds, we continue to believe that NOW can reverse the "SaaSpocalypse" narrative. We see a path through as the "0-to-1" funnel (Pro Plus / Now Assist) coupled with the growing governance complexity drives growth of the AI Control Tower, and a longer-dated flywheel in autonomous workflows where ServiceNow's deterministic execution is hardest to copy. In an AI tape heavily indexed to model intelligence, we believe NOW's context moat represents a credible bet on agentic AI; and a reasonable path for the stock to claw back from the "AI losers" bucket towards the winners' podium as enterprises shift from generating answers to resolving work with audit-grade evidence.

Databricks – From Data Engineering to Agent Orchestration

Where They Sit in the Agentic Stack as a Picks & Shovels Play

Layer	Data Intelligence Platform, Agentic Infrastructure, &"Compound AI" Systems
Role	Unified foundation where data and AI are vertically integrated to power autonomous agents
User	Enterprise Data Engineers, AI Developers, & Fortune 500 organizations

Financial Snapshot	
Most Recent Funding / Valuation	In February 2026, Databricks finalized a $7.0 Billion Series L financing round (including $5.0 Billion in equity and $2.0 Billion in debt), valuing the company at $134.0 Billion post-money.
Most Recently Reported Revenue Figure	As of February 2026, Databricks surpassed a $5.4 Billion annualized revenue run rate.
Performance Updates & Momentum	The company reported >65% year-over-year growth as of early 2026 and has maintained positive free cash flow over the prior 12 months.
Monetization & Strategic Leverage	Monetization is driven by a consumption-based model (Databricks Units) capturing value across the data lifecycle. Strategic leverage is anchored in the "Lakehouse" architecture, which prevents data silos and allows AI agents to access real-time enterprise data with built-in governance through the Unity Catalog.

Source: Databricks company reports & website

The Compound AI Philosophy: Moving Beyond Monolithic Models

While foundational model providers focus on "raw" intelligence, Databricks has positioned itself as the factory for Compound AI Systems. Their core strategic thesis posits that agentic performance is not driven by a single "god model" but by a modular system of models, retrievers, and tools tightly coupled with an organization’s proprietary data. By shifting the focus from model-centric to system-centric intelligence, Databricks enables enterprises to solve complex problems that a single LLM, regardless of its scale, cannot handle reliably due to context windows or hallucination risks.

This "industrial grade" approach allows for a larger design space where agents are grounded in accurate, governed enterprise facts. The Databricks Data Intelligence Platform unifies this entire lifecycle from data engineering to model deployment, ensuring that agents are not just "smart," but are contextually aware and compliant with enterprise security standards.

The Data Layer: Lakebase as "Operational Memory"

The most critical evolution of the Databricks stack is Lakebase, a fully managed, serverless operational database built on Postgres technology. Lakebase serves as the "working memory" for AI agents, providing the low-latency, real-time data store necessary for an agent to track its own state and "remember" previous interactions in production.

The impact of this layer is already visible in enterprise workflows; as of early 2026, AI agents are responsible for building 80% of all new enterprise databases within the Neon-powered

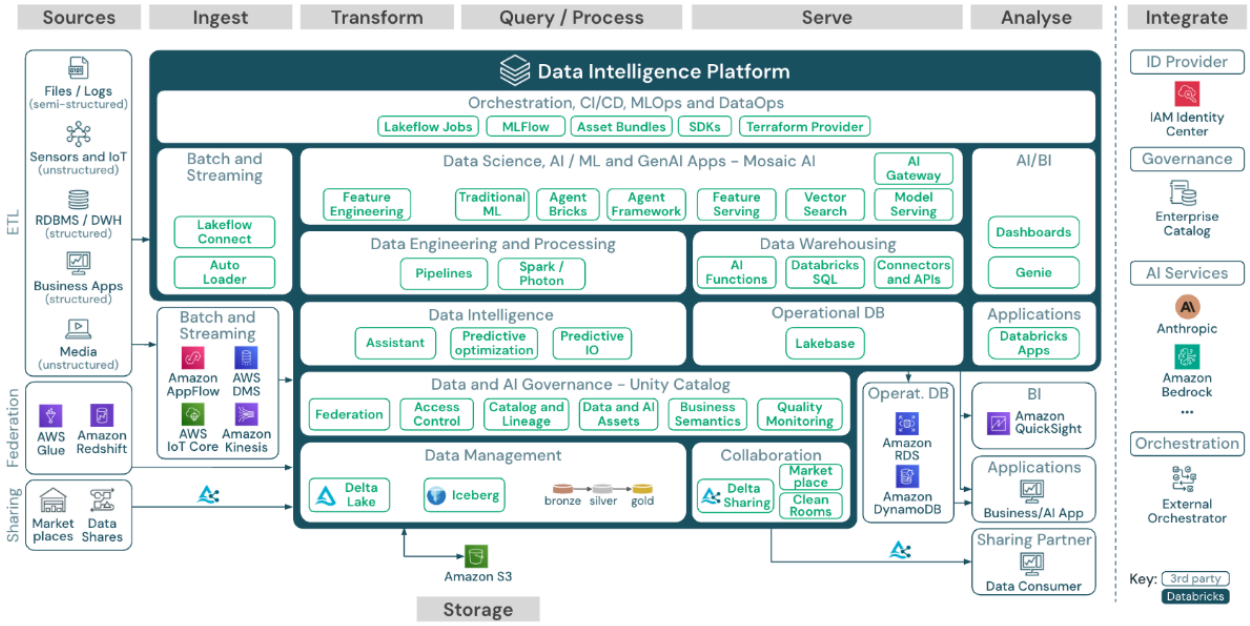
Lakebase environment. This marks a fundamental shift where agents are no longer just querying data but are autonomously constructing the infrastructure they need to operate.

The Engagement Layer: Agent Bricks & Mosaic AI

To bridge the gap between data and action, Databricks launched Agent Bricks, a native framework designed to accelerate the deployment of domain-specific agents. Agent Bricks automatically optimizes agents based on a customer's unique data schemas and business definitions, ensuring they are cost-efficient and trustworthy for specific use cases like "Agentic BI."

This is supported by the Mosaic AI Agent Framework, which acts as the orchestration engine for the stack. It provides the toolchain required for developers to experiment with, deploy, and monitor "Compound AI" systems. By integrating with MLflow for observability and Databricks Apps for delivery, the company provides a seamless environment for developers to build agent-driven user experiences that are grounded in the Unity Catalog's governance framework.

Figure 31. Databricks Lakehouse Architecture



Source: Databricks website

Hardening the Stack: Strategic M&A and Security

Databricks has aggressively used M&A to fill gaps in agent reliability, security, and real-time capabilities, ensuring their platform is the "safe" choice for autonomous enterprise systems. The March 2026 acquisition of Quotient AI addresses the industry's critical challenge: the complex failure modes that emerge when agents enter production. By embedding Quotient's reinforcement learning and continuous evaluation tools into Agent Bricks, Databricks allows

enterprises to systematically grade agent performance and catch hallucinations before they impact business logic.

Security has also been a focal point with the acquisition of Shield AI in early 2026. This move introduced "Zero Trust" for autonomous systems, defending against emergent risks like prompt injection and "Shadow AI." By providing a dedicated AI Security Framework (DASF), Databricks addresses the governance bottlenecks that have historically slowed agentic adoption, moving the conversation from "what agents can do" to "how we can trust them."

Security Software - Siddiqui

CrowdStrike – The Real-Time Control Layer for Autonomous Threat Defense

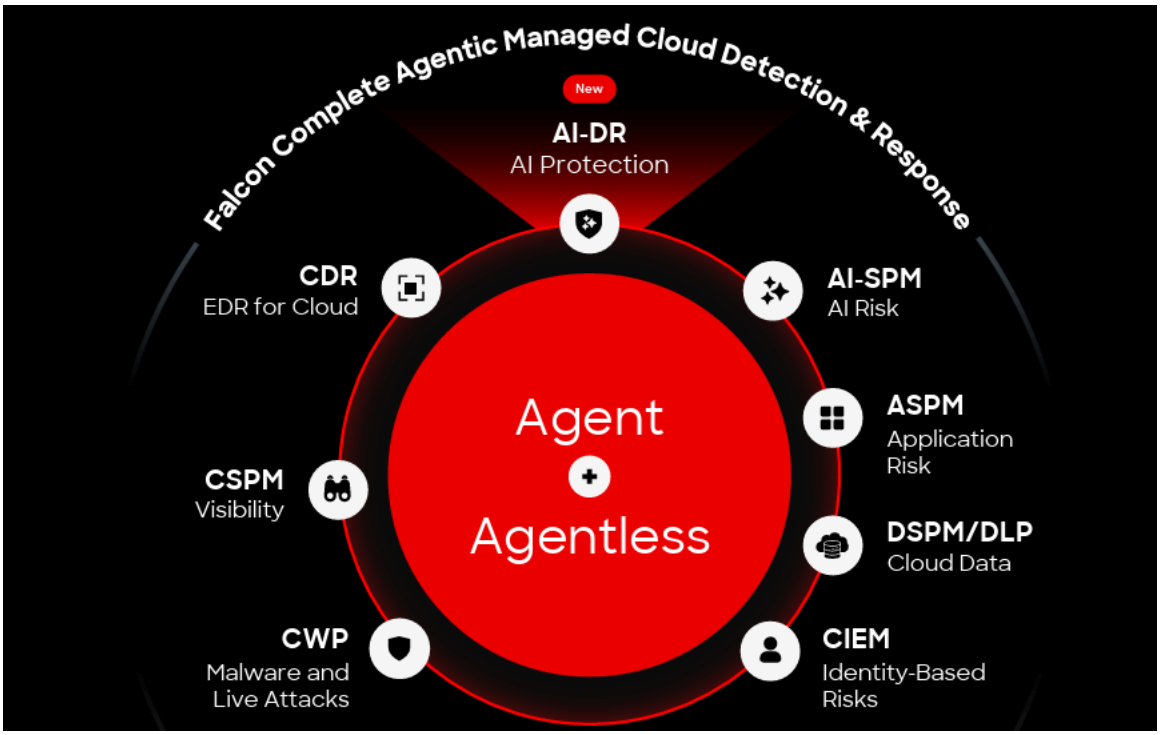
Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	Runtime Layer
Role	AI-Native Runtime Protection for Endpoints and Workloads
User	Security Ops, Incident Response Teams

Financial Snapshot ($Ms)					
Mkt. Cap	**$113,652**	**Net Debt**	**($4,410)**	**Enterprise Value**	**$109,242**
Revenue			AI fueling adoption of Gen AI assistant, Charlotte AI (usage >6x y/y, ARR up 3x y/y), AIDR (up 5x q/q), and driving growth in emerging businesses ($1.9 bil ARR, up 45%)		
Most Recent Performance Updates & LT Targets			$10 bil in ending ARR by FY31 and $20 bil by FY36		

Source: FactSet – Financial Data as of 4/27/2026

CrowdStrike's agentic security business marks a major evolution in the company's strategy, redefining the Falcon platform as an AI-native, autonomous security ecosystem designed for the rapidly intensifying era of AI-driven threats. Central to this transformation is the Agentic Security Platform, launched in September 2025 at Fal.Con, which introduces the Enterprise Graph, an AI-ready data layer that unifies real-time telemetry from across the enterprise and provides the contextual depth required for machine-speed defense. By combining CrowdStrike's graph analytics with real-time streaming from the Onum acquisition and millions of decisions from Falcon Complete analysts, the Enterprise Graph strengthens security operations with instant insight and actionability. This data-centric foundation positions CrowdStrike not just as an endpoint security vendor but as a platform orchestrator capable of powering an autonomous Security Operations Center (SOC).

Figure 32. CrowdStrike's Agent Security Platform



Source: CrowdStrike presentation

A core pillar of CrowdStrike's agentic strategy is Charlotte AI AgentWorks, the first no-code platform for building, testing, and deploying trusted security agents at scale. This tool allows security teams to create AI agents using natural language instructions, enabling them to automate bespoke workflows without writing code. AgentWorks supports both customer-built and third-party agents, all governed through the Model Context Protocol (MCP), which ensures secure multi-agent collaboration across IT and security environments. These advances collectively enable organizations to operationalize agentic AI, giving both humans and autonomous agents a shared data fabric and standardized orchestration layer to accelerate investigation, decision-making, and remediation.

Building on this foundation, CrowdStrike introduced the Agentic Security Workforce—a suite of mission-ready autonomous agents embedded across Falcon modules to automate high-impact workflows such as exposure prioritization, malware analysis, threat hunting, correlation rule generation, and SIEM data transformation. These agents draw on Falcon telemetry and past analyst decisions to act with the judgment of an experienced security professional but at machine-speed, dramatically reducing dwell time and scaling analyst capacity. CrowdStrike frames this as a shift from reactive alert triage to proactive, continuous defense, where agents reason, decide, and act autonomously while keeping analysts in command through oversight and governance controls. This approach reflects the company's view that the modern SOC must evolve into an "agentic SOC," where coordinated ensembles of AI agents perform work previously requiring multiple tiers of human analysts.

From a strategic business standpoint, CrowdStrike positions agentic security as the next foundational shift in cybersecurity, comparable to the introduction of firewalls or the cloud security transition. The company argues that generative AI has fundamentally altered adversary speed and sophistication, compressing response windows and making human-only SOCs obsolete. In this environment, autonomous, agent-driven platforms are no longer optional but essential for survival. CrowdStrike's agentic roadmap is reinforced by strategic partnerships with Nvidia, Salesforce, and Meta, which expand its ecosystem and enhance its ability to deliver AI-accelerated security solutions at global scale. Early customer feedback from Fal.Con 2025 events highlighted strong enthusiasm for these innovations and the consolidation benefits of CrowdStrike's unified platform approach, as organizations seek to reduce tool sprawl while strengthening defense effectiveness.

Overall, CrowdStrike's agentic security business integrates autonomous AI agents, an AI-ready data fabric, multi-agent orchestration, and human-centric governance into a unified system built to operate at machine speed. By merging intelligence, automation, and contextual insight, the company aims to redefine how enterprises detect, investigate, and remediate threats—positioning itself as a leader in the emerging era of autonomous cybersecurity operations. Being selected as one of only two cybersecurity vendors among the 12 launch partners in Anthropic's Project Glasswing signals that CrowdStrike is viewed as mission-critical infrastructure for securing frontier AI systems, not just another security tool.

On April 23rd, CrowdStrike launched Project QuiltWorks, an industry-wide defensive cybersecurity initiative to address the surge of software vulnerabilities being uncovered by frontier AI models, vulnerabilities that are collapsing the traditional gap between discovery and exploitation. The initiative unites frontier AI providers (including OpenAI and Anthropic) with leading global services firms (such as Accenture, EY, IBM Cybersecurity Services, and Kroll) to continuously assess, prioritize, and remediate exploitable risk in real production environments, moving beyond CVE- and alert-based security models. CrowdStrike sits at the center of QuiltWorks as the orchestration and control plane, integrating AI-driven vulnerability discovery into the Falcon platform and applying adversary intelligence, attack-path analysis, and trillions of daily security signals to determine which weaknesses are actually reachable and exploitable by attackers. Strategically, QuiltWorks positions CrowdStrike not just as a detection vendor, but as the platform that operationalizes frontier AI for defense, translating cutting-edge model capability into board-level risk visibility and scalable, code-level remediation at machine speed.

Cloudflare – The Global Edge Control Plane Where Security, Connectivity, & AI Converge

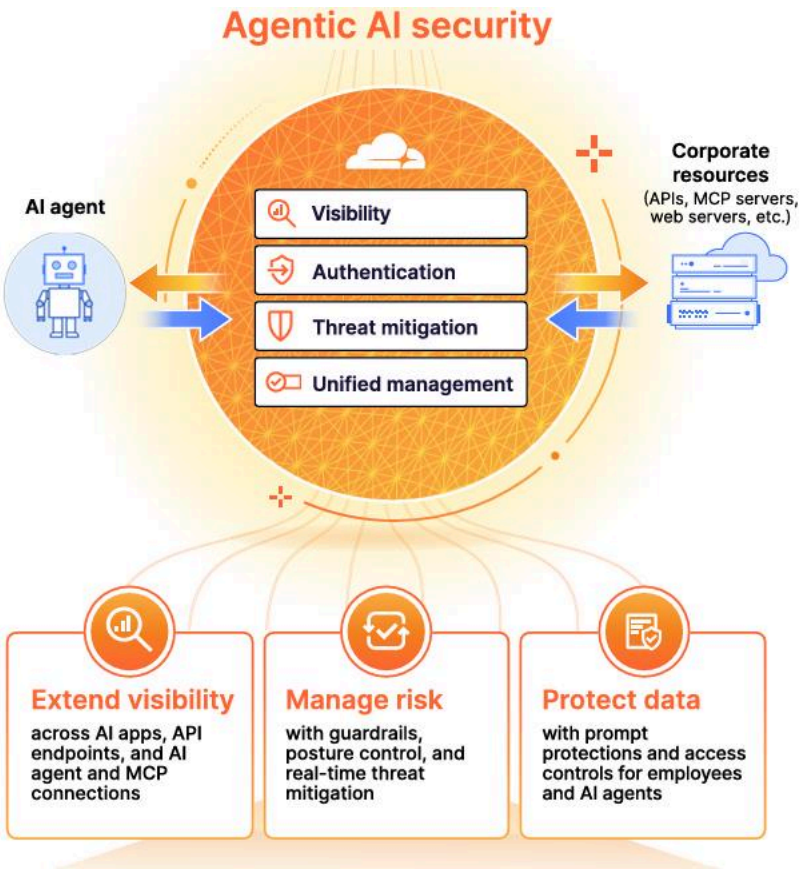
Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	Connectivity Cloud (network Layer for Application Traffic)
Role	Connecting Disparate Clouds, Moving Compute Layer Closer to Edge
User	Developers, DevOps Engineers, CISOs

Financial Snapshot ($Ms)					
Mkt. Cap	**$72,885**	**Net Debt**	**($592)**	**Enterprise Value**	**$72,292**
	Revenue		While there are no disclosed AI-related financial metrics the company noted weekly AI-agent requests across its network more than 2x in Jan 2026. Secured a record $85 mil AI-native customer win tied to full platform		
	Most Recent Performance Updates & LT Targets		$5 bil annualized revenue run-rate by Q4 2028		

Source: FactSet – Financial Data as of 4/27/2026

Cloudflare's agentic-AI security strategy positions the company not as a traditional cybersecurity vendor but as a foundational infrastructure layer for the emerging agent-driven internet. As organizations accelerate adoption of generative and autonomous AI systems, Cloudflare frames its global connectivity cloud as the platform on which AI agents will run, communicate, authenticate, and enforce security policy at scale. This positioning reflects both Cloudflare's architectural advantage, its distributed edge network running within 50 milliseconds of 95% of global users, and its view that AI agents are becoming the new "users" of the internet, driving exponential growth in automated traffic that demands new forms of identity, governance, and real-time security enforcement. The company's AI Security Suite provides unified security for the entire AI lifecycle: discovering shadow AI, protecting models from abuse, enforcing Zero Trust controls for machine-to-machine communication, preventing data leakage in prompts, and providing guardrails for AI-powered applications, all from a single policy layer built into the Cloudflare network.

Figure 33. Agentic AI Security



Source: Cloudflare company website

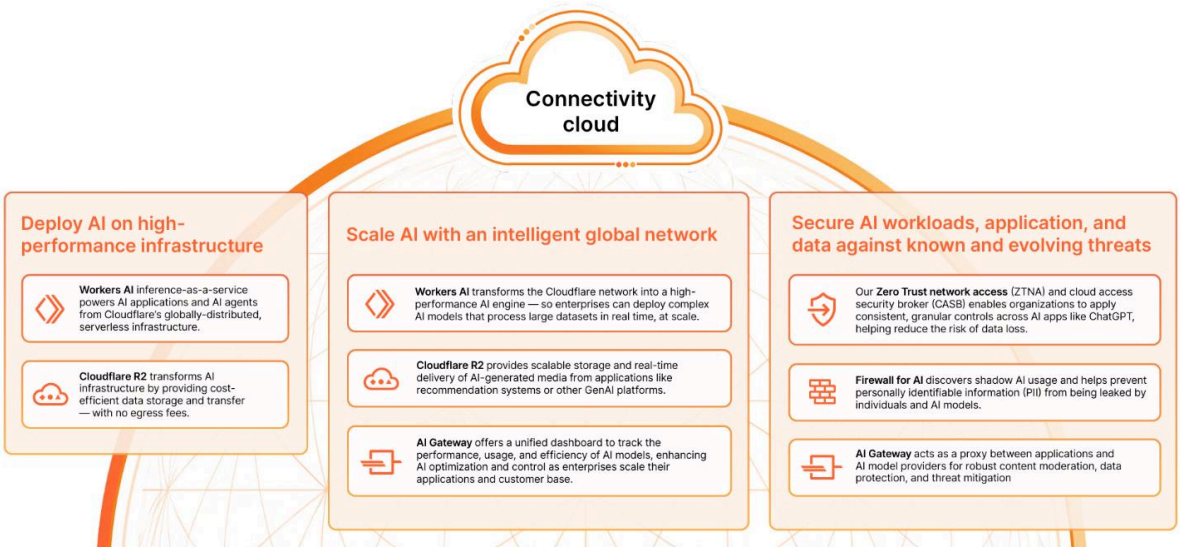
A major pillar of Cloudflare's agentic strategy is Cloudflare for AI, introduced in 2025 as a comprehensive suite for protecting AI model deployments, usage, and data integrity. This suite addresses threats such as toxic prompts, model poisoning, PII leakage, and misuse of AI apps by employees, while enabling organizations to move quickly without compromising governance. Through Firewall for AI, Cloudflare automatically discovers and labels all AI applications in use—authorized or unauthorized—giving CISOs full visibility across the enterprise. This discovery capability is critical as AI experimentation becomes widespread and attackers begin targeting AI-driven applications with new forms of exploitation. Cloudflare emphasizes that the gap between experimentation and safety is widening, arguing that companies need a security layer that can keep pace with rapid model iteration and distributed AI adoption.

Cloudflare is also deeply investing in the agentic-internet infrastructure required to support autonomous systems at scale. A notable milestone came with the launch of the industry's first remote MCP server, enabling developers to build, deploy, and operate AI agents directly on Cloudflare's network. This capability removes the cost and complexity barriers traditionally associated with agent development, allowing agents to execute tasks—such as deploying code or sending messages—autonomously using standardized protocol-based access to external tools and services. The remote MCP server, combined with durable workflows and Workers AI,

forms a complete agentic stack where reasoning, execution, and tool-access run natively on Cloudflare's global network, enabling businesses to operationalize agentic workloads in minutes rather than months. Cloudflare frames this as a natural extension of its historical strength: a massively distributed network purpose-built for running compute, identity enforcement, and application-level protection at the edge.

Meanwhile, Cloudflare's connectivity cloud is becoming a neutral, multi-cloud control plane for AI workloads. The company describes itself as the "Switzerland of AI," offering tools such as Workers AI, AI Gateway, and Data Localization Suite to help enterprises run, monitor, and govern AI workloads across providers. Workers AI now runs GPU-accelerated inference across 150+ cities worldwide, and AI Gateway surpassed 5 billion requests—a sign that Cloudflare has become a preferred distribution and governance layer for enterprise AI. The company reports surging adoption of its sovereign AI solutions, reflecting government and enterprise demand for AI systems that remain within defined geographic jurisdictions. Sophisticated AI-powered bot attacks, which Cloudflare says increased more than 300% in 2025, have also pushed enterprises toward AI-native security models that adjust defenses autonomously and at scale, reinforcing the value of Cloudflare's integrated platform.

Figure 34. Cloudflare Connectivity Cloud



Source: Cloudflare company Website

Cloudflare's Zero Trust business is also expanding into AI production governance. New Zero Trust tools provide automatic visibility into how employees and systems interact with AI, protecting against Shadow AI, misconfigured GenAI applications, and unintentional data leakage. AI Security Posture Management (AI-SPM) offers a centralized view of generative-AI usage, enabling organizations to enforce policy, block unsafe prompts, and ensure AI is used safely across teams. These capabilities reflect Cloudflare's belief that secure AI adoption is not simply a security problem but an enterprise governance challenge requiring identity-aware enforcement, deep traffic visibility, and seamless policy orchestration across cloud and edge environments.

Together, these initiatives position Cloudflare as a core enabler of the agentic internet. By unifying security, networking, identity, and compute into one AI-capable platform, Cloudflare aims to become the default execution environment for AI agents—whether they are powering back-office automation, interacting with customers, or making real-time decisions across distributed systems. In a world where AI agents increasingly outnumber human users, Cloudflare's model shifts the focus from securing human workflows to securing autonomous machine-driven operations at global scale.

Palo Alto Networks – The Unified Security Platform Governing AI-Native Enterprises

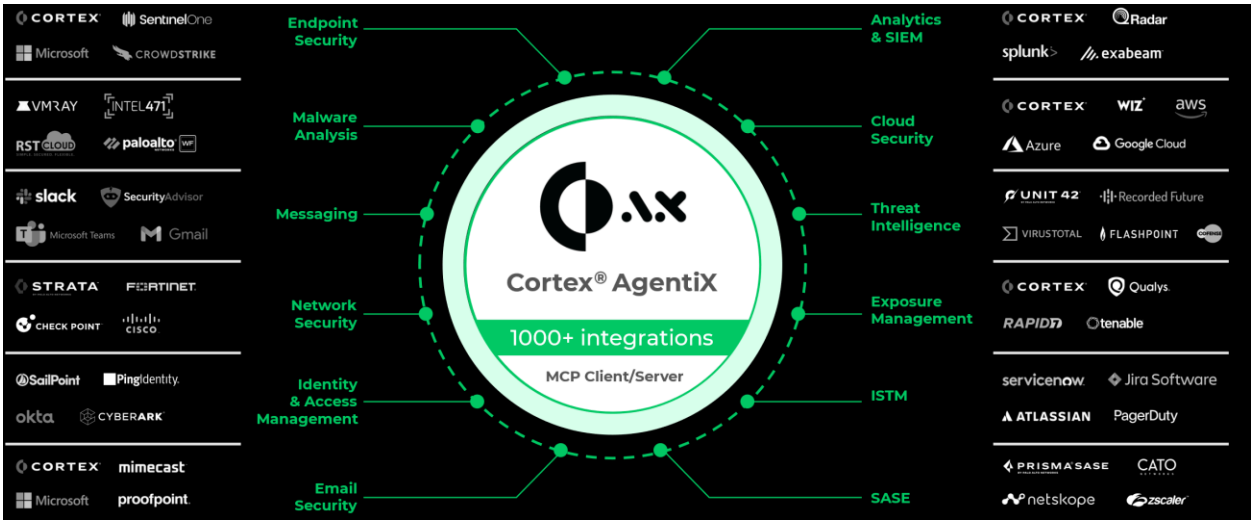
Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	Control Point Layer (Network, Endpoint, Cloud, Browser, Identity)
Role	Unified "Platformization" across SASE, Cloud, and SOC
User	CISOs

Financial Snapshot ($Ms)					
Mkt. Cap	**$145,689**	**Net Debt**	**($2,496)**	**Enterprise Value**	**$143,192**
Revenue			While there are no direct disclosures, PANW notes Prisma AIRS has surpassed 100 customers and build 9-figure pipeline. XSIAM, PANW's AI-driven SOC, surpassed $500M in ARR as of 2Q26.		
Most Recent Performance Updates & LT Targets			$20Bn NGS ARR target and 2,500 - 3,000 platformizations by FY30 and adj. FCF margin of 40% by FY28.		

Source: FactSet – Financial Data as of 4/27/2026

Palo Alto Networks' agentic-AI security strategy centers on transforming its unified platform into the architectural backbone for enterprises entering an AI-first, agent-driven operating environment. The company views the rise of autonomous AI systems, both as powerful defensive tools and as new insider-threat vectors, as a foundational shift that requires security platforms capable of operating at machine speed.

Figure 35. Cortex AgentiX Ecosystem & Capabilities



Source: Palo Alto Networks Website

At the center of Palo Alto Networks' agentic roadmap is Cortex AgentiX, launched in October 2025 as the next generation of its automation stack and the most secure platform to build, deploy, and govern enterprise AI agents. Built on a decade of security automation and trained on 1.2 billion playbook executions, AgentiX enables prebuilt and custom agents that can plan, reason, and execute with expert-level precision, delivering up to a 98% reduction in mean-time-to-respond and 75% less manual work for SOC teams. These agents span threat intelligence, email investigation, endpoint forensics, and network policy automation, extending across thousands of integrations, with full Model Context Protocol (MCP) support, allowing agentic workflows to orchestrate security actions across the enterprise. CEO Nikesh Arora positions AgentiX as the "next step in security automation," embedding agent capabilities not just in XSOAR but across Cortex Cloud, Cortex XSIAM, Cortex XDR, and eventually as a standalone platform starting in 2026. Parallel innovations such as Prisma AIRS 2.0 and Cortex Cloud 2.0 further reinforce the company's intent to embed agentic capabilities natively across cloud security, exposure management, and email defense, replacing manual investigation and remediation workflows with proactive, autonomous systems.

This agentic approach extends into Palo Alto Networks' broader platform strategy, which aims to eliminate fragmented point-product architectures in favor of an integrated, adaptable system capable of supporting unpredictable AI-era operating environments. At its Ignite 2025 event, the company highlighted innovations such as Prisma Browser, AI-powered Strata Cloud Manager, and agent-driven Cortex capabilities designed to simplify operations, improve agility, and counter increasingly fast and sophisticated AI-driven threats. The company argues that with customers often running hundreds of security tools, only a unified platform, backed by AI reasoning, shared telemetry, and integrated enforcement, can deliver the agility needed to match the pace of autonomous systems and modern adversaries. This platform-centric perspective is reinforced by Palo Alto Networks' acquisitions of CyberArk and Chronosphere, valued at $25 billion and $3.35 billion respectively, which underscore the company's belief that identity security and "always-on" observability are foundational to enabling safe, large-scale AI adoption across the enterprise.

Palo Alto Networks also ties the rise of agentic AI to a dramatic expansion of the cloud attack surface. Its 2025 *State of Cloud Security* report found that 99% of organizations experienced attacks on AI applications within the prior year, while insecure GenAI-assisted code creation is accelerating risk beyond teams' ability to remediate. API-based attacks rose 41% due to agentic AI's dependence on APIs, while identity remained the weakest link, compounding lateral-movement risk across cloud environments. The company stresses that traditional cloud security tools cannot counter these machine-speed threats and advocates for an agentic-first platform that spans code, cloud, identity, and SOC operations to operate faster than the adversary.

Taken together, Palo Alto Networks' agentic-AI strategy seeks to deliver a unified, secure, and adaptive security architecture built for a world where autonomous agents, malicious and defensive, operate continuously and at speeds far beyond human capacity. Through Cortex AgentiX, AI-driven SecOps modernization, identity-centric investments, and cloud-wide automation, the company aims to position itself as the central nervous system for enterprise

security in the AI decade: a platform that replaces manual toil with autonomous defense, consolidates fragmented ecosystems, and governs intelligent agents across the entire attack surface at machine speed.

The selection of Palo Alto Networks as a core participant in Anthropic's Project Glasswing highlights its role as a foundational security platform for the AI era, reinforcing the view that large-scale, integrated cyber platforms are central to governing, securing, and operationalizing frontier models like Claude Mythos as enterprises adopt increasingly autonomous AI systems.

Zscaler – The Secure Access Platform Enabling Zero-Trust at Enterprise Scale

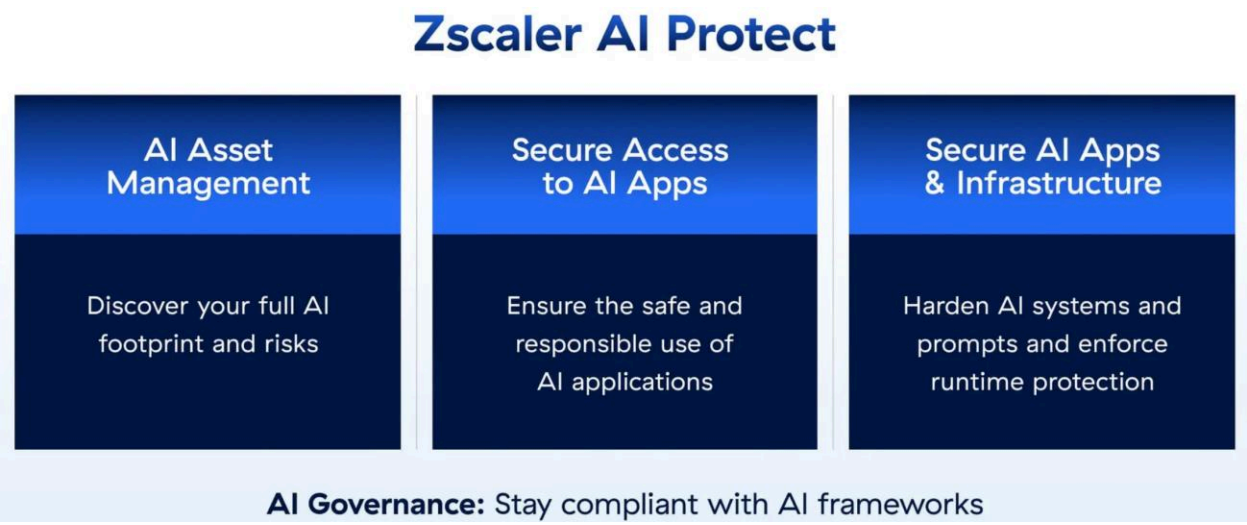
Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	Identity-Centric Connectivity and Security Layer
Role	Zero Trust Exchange & Policy Engine
User	CIOs, CISOs and Network Architects

Financial Snapshot ($Ms)					
Mkt. Cap	**$21,787**	**Net Debt**	**($1,776)**	**Enterprise Value**	**$20,011**
	Revenue		>$400M ARR currently and expected to be >$500M in ARR by end of FY26 (July 31 FYE). AI Security includes Security for AI apps and Agentic Operations, which includes IT Ops and Security Ops.		
	Most Recent Performance Updates & LT Targets		Company has aspirational goal of reaching $10Bn in ARR underpinned by its leadership in Zero Trust security and growing AI Security pillar.		

Source: FactSet – Financial Data as of 4/23/2026

Zscaler's emerging agentic AI security business reflects a strategic transformation of its Zero Trust Exchange platform into a foundational security layer for enterprises adopting generative and autonomous AI at scale. As organizations accelerate toward an "agentic-ready estate," where AI agents autonomously interact with systems, data, and applications, Zscaler positions itself as the trust and governance backbone needed to operationalize these capabilities safely. The shift is driven by the company's recognition that traditional, human-centric security architectures cannot secure AI-driven traffic patterns, new protocols, and non-human interactions that define modern autonomous systems. Zero Trust, according to the company, not legacy network controls, is the only viable model for securing AI agents that operate continuously, make decisions, and manipulate sensitive enterprise data in real time.

Figure 36. Zscaler AI Protect



Source: Zscaler Website

This strategy is embodied in Zscaler's AI Security Suite, launched in January 2026, which provides comprehensive visibility, governance, and control across enterprise AI assets. The suite delivers a consolidated inventory and dependency map of models, agents, embedded AI in SaaS applications, development environments, and MCP-based agent collaboration infrastructure—addressing a critical blind spot highlighted by Zscaler's ThreatLabz 2026 report, which found that 100% of enterprise AI systems analyzed contained critical vulnerabilities and could be compromised in as little as 16 minutes. Zscaler correlates asset discovery with data lineage, access relationships, runtime behavior, and security posture to give companies actionable insights into where AI systems are deployed, what data they touch, and how risk propagates across autonomous workflows. This level of observability forms the foundation for governing AI at enterprise scale.

At the operational layer, Zscaler is bringing Zero Trust principles directly to agentic AI through Secure Access to AI, which enforces real-time policy controls and inline inspection for AI workloads, including developer tools, LLMs, GenAI applications, and autonomous agents. Enhanced prompt visibility, DLP-driven safeguards, and inline inspection capabilities allow enterprises to safely enable tools like Copilot while blocking high-risk prompts, preventing data leakage, and maintaining compliance across AI-powered workflows. Complementary AI-powered data security classification expands Zscaler's ability to detect sensitive content across more than 200 data categories, providing the deep semantic understanding required to protect AI systems that generate, transform, and transmit high-value information at machine speed. These features reinforce Zscaler's value proposition: AI can only scale as fast as the security controls that govern it.

Zscaler's broader platform strategy also highlights a push into autonomous SOC transformation, supported by acquisitions such as Red Canary, which brought advanced agentic-AI-driven detection and response capabilities. Zscaler's scale advantage, underpinned by processing

more than 500 trillion security signals per day, further strengthens its ability to train AI models and deliver high-accuracy threat detection and automated remediation across distributed environments. Zscaler's agentic security strategy positions the company as a critical enabler of autonomous enterprise transformation. By combining visibility, governance, Zero Trust policy enforcement, AI-powered data protection, and autonomous SOC workflows, Zscaler is building a security architecture capable of governing both human and non-human actors. As enterprises deploy AI agents across core workflows, Zscaler's platform aims to ensure that innovation does not outpace control, delivering a unified, cloud-native security fabric designed for the speed, scale, and complexity of the AI era.

In a company blog post on April 21, CEO Jay Chaudhry publicly confirmed and detailed its role in Project Glasswing, including access to Claude Mythos Preview. It isn't just a peripheral observer but an active, vetted defender within the Glasswing coalition, using frontier AI capabilities in support of its Zero Trust and exposure-elimination strategy. Zscaler's involvement in Project Glasswing positions it not just as a vulnerability-discovery leader, but as the enforcement layer that renders discovered vulnerabilities unreachable in real-world environments operating at machine speed.

Internet & Digital Media - Squali

Alphabet – Defending the Search Kingdom

***Summary**: Alphabet is transitioning from "Word Models" to "World Models" with Gemini 3, leveraging Google Agentspace and Chrome Auto-Browse to maintain its position at the top of the discovery funnel. The company's Universal Commerce Protocol (UCP, more detail HERE) is a critical development to watch as it aims to standardize how agents interact with millions of merchants, potentially preserving Google's role as the primary intermediary for digital transactions.*

Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	Foundational LLM, agent infrastructure, developer platform, consumer services
Role	Horizontal AI infrastructure provider (while being vertically integrated itself)
User	Developers, companies of all sizes (Enterprises, SMBs, Startups), anyone building AI native apps or workflows, and government agencies

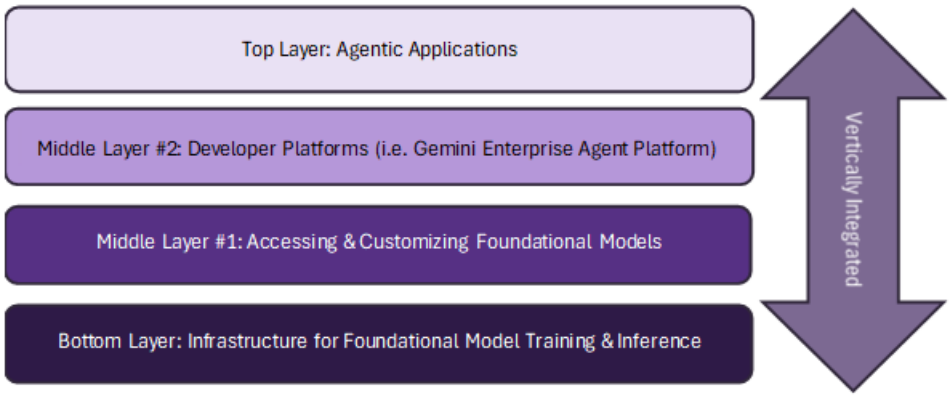
Financial Snapshot ($Ms)					
Mkt. Cap	**$4,166,207**	**Net Debt**	**($59,847)**	**Enterprise Value**	**$4,106,360**
Revenue	In 4Q25, Google Cloud revenues grew 48% y/y to $18B, accelerating from 34% in 3Q, driven primarily by growth in GCP, as well as new customer commitments and growth in existing spend. Workspace grew DD % Y/Y by both avg. revenue per seat and total number of seats. Both segments continue to be driven by higher demand for gen-AI, fueling continued supply / demand imbalances for AI-related services and products				
Most Recent Performance Updates & LT Targets	**TPU:** 75% of Cloud customers now use GOOGL's vertical AI offerings (chips, models, agents - **Gemini:** Revenue from products built on its gen-AI models grew ~400% Y/Y significantly accel. from 3Q. Google now has 14 product lines exceeding $1B in annual revenue				

Source: FactSet – Financial Data as of 4/27/2026

Google's AI Framework: Owning the entire vertical stack

Google is the only player that owns the entire vertical stack, from the hardware (its own Tensor Processing Unit or TPU v7 Ironwood) to the model (Gemini) to the data (Search/YouTube) to the end-user distribution (Android/Workspace). Unlike Amazon, which focuses on providing the building blocks for others to build agents (via Bedrock), Google is focused on embedding autonomous agents directly into its dominant consumer and enterprise surfaces (Search, Workspace, Android, and GCP). While Amazon wins on "neutrality," Google's approach is a vertically integrated play intended to protect its high-margin Search business while expanding into high-margin Enterprise AI.

Figure 37. Layers of Google's Vertically Integrated Stack



Source: Google website

See the table on below for a comprehensive view of GOOGL's end-to-end AI products and agentic offerings.

Figure 38. Google End-to-end AI Products & Agentic Offerings

Foundational AI Models & Generative Media			
Model	Description	Key Insight	Target Audience
Gemini 3.1 Pro	Advanced reasoning, complex math, and agentic coding	3x token volume vs2.5 Pro; features advanced "Hybrid Reasoning."	Enterprise & Power Users
Gemini 3 Flash	Low-latency, high-volume multimodal tasks	Frontier performance at a fraction of the cost; successor to 2.5 Flash	Developers
Nano Banana 2	Production-scale image generation and editing	Viral sensation; drove 20M new subscribers in two weeks	Creator Economy
Veo 3.1 / Imagen 4	Cinematic video (Veo) and 4K high-res imagery (Imagen)	Veo 3 generated 70M+ videos by Q2 2025; Imagen 4 supports 2K resolution	Pro Filmmakers
Gemma 3	Open-source weights for local/custom deployment	140M+ downloads; state-of-the-art performance for its weight class	Researchers & ML Engineers

Agentic Platforms & Developer Tools			
Platform	Description	Key Insight	Target Audience
Google Antigravity	Autonomous platform for planning/executing software tasks	High-velocity agentic tool; uses Gemini 3 Pro to "see" and "act" on screens	Engineers & Developers
Agent Dev Kit (ADK)	Open-source framework for multi-agent systems	Focused on "Agent-to-Agent" (A2A) and Model Context Protocol (MCP) interoperability	Engineers & Open-Source Commnutiy
Deep Research	Agentic tool for exploring complex topics via 100s of sources	Differentiates from standard search by providing cited, synthesized interactive reports	Researchers
Gemini Enterprise Agent Platform	Comprehensive Google Cloud platform designed to build, scale, govern, and optimize autonomous AI agents	Hosts 200+ models; usage grew 20x in 2024-2025 as firms shifted to production	IT Teams, Engineers

Consumer & Enterprise AI Products			
Product	Description	Key Insight	Target Audience
AI Mode	Conversational, end-to-end research interface	100M+ MAU in US/India; integrated Gemini 3 for interactive UI	Consumer
AI Overviews	Synthesized answers at the top of Search results	2B+ monthly users; drives 10% more queries in triggered categories	
Gemini App	Proactive assistant for personal productivity, multimodal search, and complex research	750M+ MAUs as of Q4 2025	
Gemini Enterprise	Workspace integration (Docs, Sheets, Gmail)	8M+ paid seats sold in four months; saves employees ~3 hours per week	Enterprise
Google Agentspace	Interoperable platform for custom business agents	1M+ pre-booked subscriptions; allows for autonomous employee transactions	

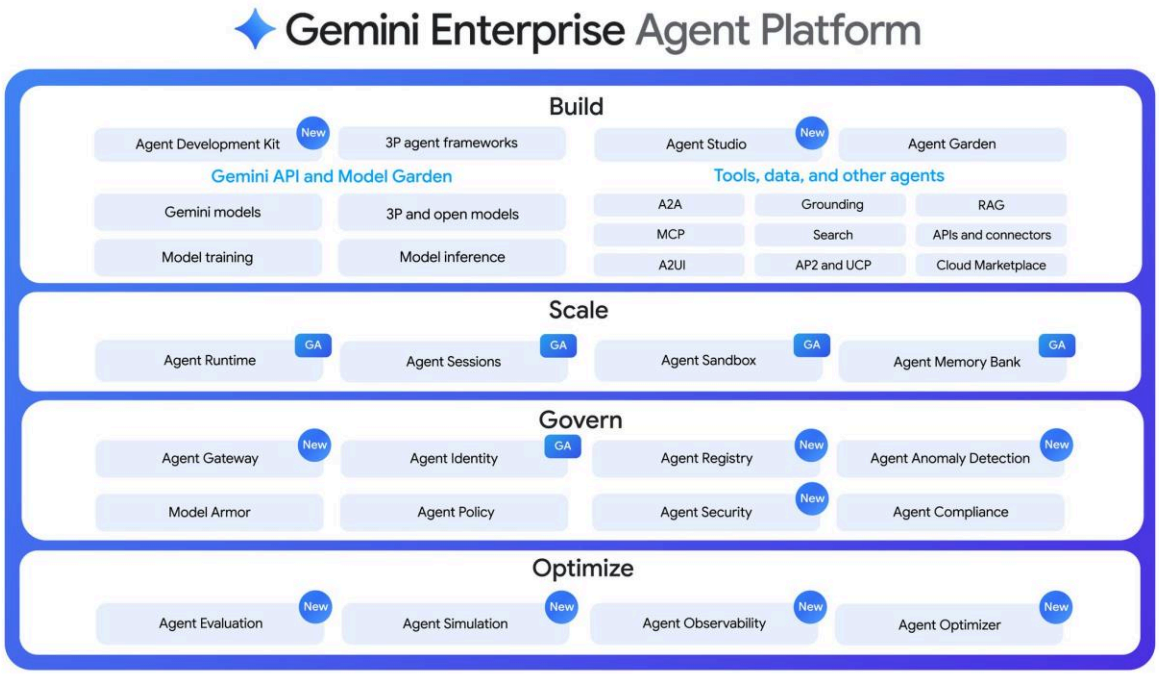
Source: Google company reports

Figure 39. Google Custom Silicon & Hardware Overview

Custom Silicon and Hardware Strategy				
Chip Name	Description	Key Insight	Target Customer	Designer / Manufacturer
Ironwood (TPU v7)	Optimized for high-complexity training and inference	Nearly closes the FLOPs gap with NVIDIA flagship GPUs; core driver of Google's unit cost efficiency	Frontier AI Labs, Large Enterprises	Google & TSMC / Broadcom & MediaTek
Trillium (TPU v6)	Designed for large-scale training and persistent infrastructure	Delivers a 4.7x peak compute performance increase; foundational to the "AI Hypercomputer"	Cloud customers requiring large-scale model training and infrastructure support	Google & TSMC / Broadcom & MediaTek
Axion CPU	Designed for general-purpose data center workloads and high compute	Offers 50% better performance and 60% better efficiency than comparable x86 instances	Internal services (Search, YouTube) and Google Cloud customers	Google & TSMC / Arm
NVIDIA Vera Rubin	For trillion-parameter inference	Offers 10x more performance per watt than Blackwell	AI model builders and frontier reasoning researchers	NVIDIA & TSMC
NVIDIA Blackwell	For data center training and agentic AI at scale	Integrated into AI Hypercomputer alongside internal TPUs to optimize diversity	Enterprise Cloud customers and software partners scaling AI operations	NVIDIA & TSMC
Tensor G5	For on-device AI and multimodal mobile agentic tasks	First Pixel chip designed entirely in-house; marked a strategic shift away from Samsung	Pixel Series users; developers building on-device applications	Google (In-house) & TSMC
Tensor G6	For advanced edge computing and real-time generative video	Aggressive jump to 2nm process to solve thermal and battery constraints for continuous agentic AI	Pixel Series; users requiring high-performance mobile video agents	Google (In-house) & TSMC
VCU (Argos Series)	For video transcoding and real-time generative video at scale	Essential for maintaining YouTube's operational scale and Gemini 3 Video	Proprietary internal use (YouTube, Google Photos, Gemini, etc.)	Google (In-house)

Source: Google company Reports

Figure 40. Gemini Enterprise Agent Platform



Source: Google Blog

Gemini 3: At the bleeding edge of AI models

Alphabet has gone from playing catch up to OpenAI after being caught flat footed in November 2022 with the launch of ChatGPT to now being on the cutting edge of AI innovation with the release of Gemini 3 in Dec. 2025. This latest version marked a pivotal shift from experimental AI to a flywheel of rapid productization and agentic capabilities. The Gemini 3 family represents Alphabet's move toward "agentic" capabilities, i.e. models that can plan and execute tasks rather than just generating text. Gemini 3 Pro in particular saw the fastest adoption of any model in the company's history, processing 3x as many daily tokens as the previous 2.5 Pro version.

Google can leverage its trove of 1P data and is planning on investing $175-185B in CapEx in 2026 to continue improving new versions of Gemini and offer personalized agentic capabilities. Improving Gemini is what will continue to power better user experiences and monetization capabilities across its various services (gen-AI search, YouTube, Gmail, Waymo, etc.).

Here are two notable projects that are Google's two flagship "Agentic AI" research prototypes designed to move AI beyond simple text-based chat into active, autonomous assistance:

1. **Project Astra** – a prototype for a universal AI assistant that can perceive and reason about the world in real-time through vision and sound.
2. **Project Mariner** – an experimental AI agent focused on browser-based task execution and automation.

For eCommerce, Google launched Universal Commerce Protocol (UCP) which embeds agentic checkout into Search & Gemini. As the digital ecosystem shifts towards AI assistants and agents acting on behalf of users, Google will likely lean into these capabilities more aggressively in the coming years. See our recently published Agentic Commerce deep dive HERE.

Google is also ramping up development of a dedicated Gemini AI app for Apple (AAPL, NR) Mac computers to compete with Anthropic's Claude CoWork, which is available as a standalone Mac app (whereas users have to go to the web first to access Gemini). This should allow Gemini to tap into programs such as the calendar to feed data into Gemini. That ultimately lets the software fulfill queries with more personalized details.

Monetization

There are several ways Google is currently monetizing and plans on monetizing the AI opportunity over time. Near-term, the company's strategy revolves around protecting and continuing to grow search.

1. **Protecting search**: Google's strategy with generative AI is fundamentally about protecting and evolving its core $250B+ annual search advertising business, which has been one of the most lucrative ventures in history given its very high margins. Much like after the advent of mobile, Alphabet has again entered an expansionary moment for Search in our view, driven by the wide-scale deployment of Gemini 3 and the shift toward Agentic Commerce. As of February 2026, Google has successfully transitioned its core Search product into an AI-first platform, leveraging superior query understanding to monetize complex sessions that were previously unmonetizable.

2. **The enterprise opportunity.** On the enterprise side, Google competes with other AI hyperscalers in Cloud and private AI companies in LLMs, providing a suite of AI tools and features as part of its Google Cloud offering.

a. **Google Cloud Platform (GCP)** utilizes a hybrid approach of a consumption-based model (pay-as-you-go) and subscriptions for infrastructure and platform services. On its latest earnings call, management shared that it is now generating “billions” of dollars in quarterly revenue from enterprise AI products.
b. **Google Workspace** utilizes a subscription-based model (per seat) for productivity and collaboration tools.

Acquisitions to help accelerate their AI ambitions

Google DeepMind is Alphabet’s premier AI research lab, originally a London-based startup acquired in 2014. Google DeepMind’s work is bifurcated between foundational AI models that power Google's commercial ecosystem and scientific research aimed at solving complex global challenges (i.e. AlphaGo, AlphaFold, and AlphaCode) and contributes core technology (e.g. the latest Gemini models) to Google’s products. This division, led by Demis Hassabis, focuses on advanced AI research and developing general-purpose AI systems.

Most recently, GOOGLE acquired **Intersect** ($5B, announced this year) to secure renewable energy and data center infrastructure to scale AI initiatives, and **Character.AI** ($2.5B, per NY Times, which is a combination of a licensing deal and equity investment rather than an outright acquisition) and **Producer.ai** (announced in February, undisclosed amount) to enhance AI-generative content.

Amazon's AWS – The Infrastructure and Utility Play

***Summary**: Amazon's strategy is centered on becoming the "foundational layer" for agents via Bedrock AgentCore and vertically integrating the stack with Trainium silicon to lower inference costs. Management's focus on AgentCore as a "primitive building block" suggests Amazon aims to capture value not just from their own agents (Rufus, Kiro), but by providing the secure, scalable runtime for every enterprise's custom agents.*

Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	Full-stack infrastructure, developer platform, and enterprise agentic applications.
Role	Provides services across hardware, model choice, and governance tools to build and scale agents
User	Global enterprises, government agencies, startups, eCommerce customers and internal use cases

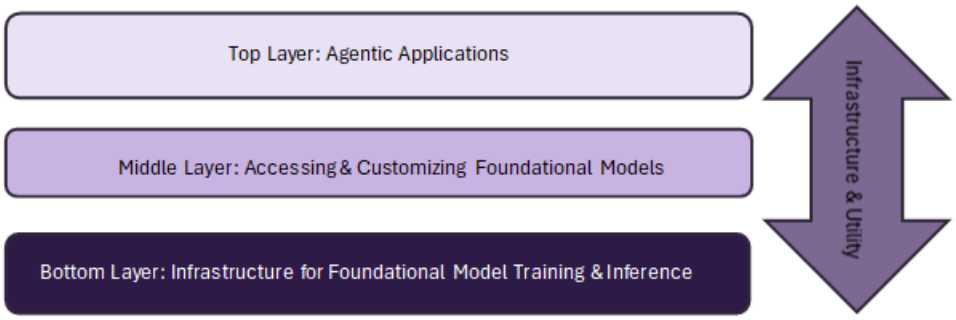
Financial Snapshot ($Ms)					
Mkt. Cap	**$2,839,028**	**Net Debt**	**$46,605**	**Enterprise Value**	**$2,885,633**
Revenue	In 4Q, AWS revenue growth came in at +24% Y/Y (at $35.6B) which marked a further acceleration from 3Q25's +20% Y/Y which was up from +17.5% Y/Y in 2Q25. The +24% Y/Y growth was the segments highest growth in 13 quarters and reflected strong underlying demand for both Gen-AI related services and core AWS cloud services. AWS backlog stood at $244B as of 12/31/25				
Most Recent Performance Updates & LT Targets	**Tranium & Gravitron:** running at a combined ARR of "well over" $10B and growing in the triple digits. **Amazon Bedrock:** Running at a multi-billion dollar ARR and saw customer spend grow 60% Q/Q in 4Q25				

Source: FactSet – Financial Data as of 4/27/2026

AWS AI Framework: The "Three-Layer" Stack:

AWS has approached its AI strategy through a lens with three distinct layers to provide a robust offering across layers to its customers. Within each foundational layer, Amazon has focused on building out a selection which includes proprietary AWS offerings in addition to various 3P offerings. While still in the relatively early days, AWS has seen strong demand for its services across all three layers of the stack, with GenAI-related services helping propel AWS growth to its highest level in 13 quarters in 4Q25. A summary of each of the three layers along with key offerings within each layer is below:

Figure 41. AWS Infrastructure & Utility Layers



Source: AWS Website

1. Bottom Layer: Infrastructure for Foundational Model Training & Inference. This layer focuses on the cost-effective training and deployment of foundation models. Key offerings across high-performance computing, storage, and networking include AWS' own custom training & inference chips along with its custom-designed CPU clusters in Graviton.

- **Custom Chips for Training & Inference:** In addition to providing customers with access to 3P chips (including those from AMD (AMD, Buy, Stein) and NVIDIA (NVDA, Buy Stein), Amazon also has its own custom-built chips for training (Trainium series chips) and Inference (Inferentia).
    - Regarding **Trainium**, AWS is seeing strong demand trends for the offering with Tranium2 chips being fully subscribed, and management expects Trainium3 supply to be fully subscribed by mid-2026. AWS is also planning to begin delivering Trainium4 chips in 2027 and the company is already in discussions for Trainium5 chips.
- **Graviton,** its own custom-designed ARM-based CPUs, is currently a multi-billion dollar annualized revenue run-rate business that saw growth eclipse 50% Y/Y in 4Q25, and 90% of the company's top 1,000 customers are using Graviton.

2. Middle Layer: Enabling customers to access & customize FMs to create applications based on their own data. The middle layer provides access to foundational models (FMs) and allows for customization of models. The main engine of growth within this layer is currently Amazon Bedrock, which is running at a multi-billion-dollar annualized run-rate with over 100,000 companies currently using the service.

- **Amazon Bedrock** is an integral piece of the middle layer as customers can access over 100 FMs from leading AI companies along with Amazon's own FMs (including the Nova and Titan models) to build and scale applications.
    - Bedrock customer spend grew 60% quarter-over-quarter as of Q4 2025, and the company is continuing to add upon the FM models that are available on Bedrock.

3. Top Layer: Agentic Applications. This layer focuses on building and implementing applications. AWS has deployed its own service to build, deploy and monitor custom built agents (Bedrock AgentCore) and Amazon recently partnered with OpenAI to launch the Stateful Runtime Environment for Bedrock – allowing agents to maintain consistent context & memory

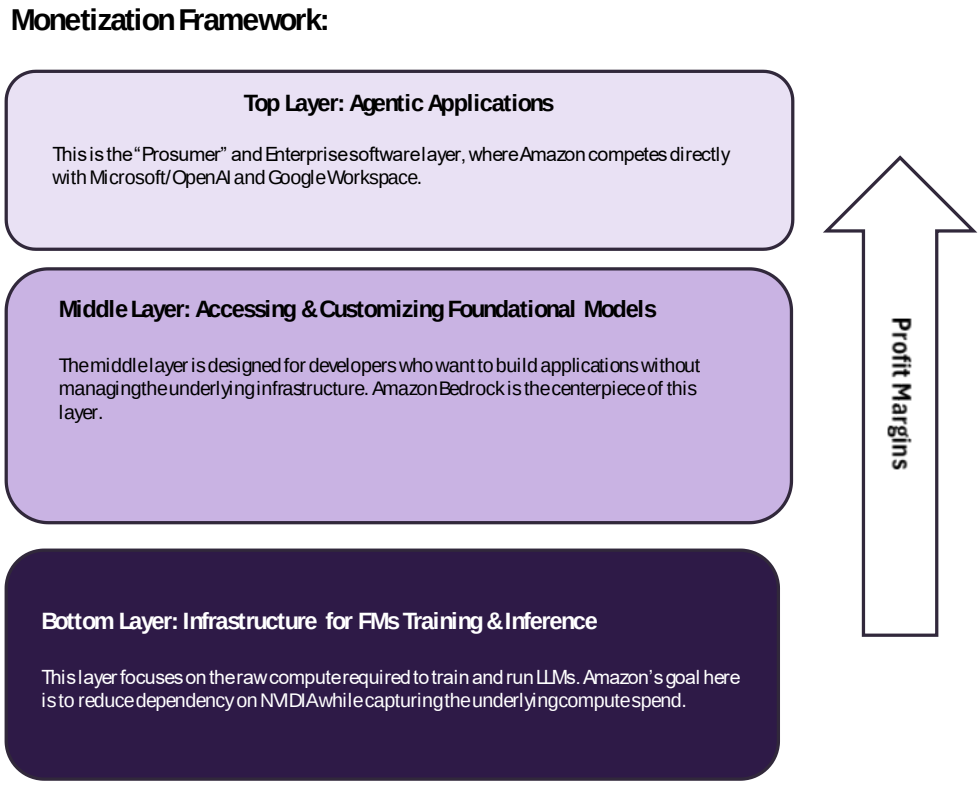
across long-running tasks. The company has also rolled out its own Frontier Agents across its internal operations and for AWS customers to use. AWS also provides governance services within its application layer.

- **Bedrock AgentCore:** Enables customers to build, deploy and monitor agents that are personalized and able to learn from prior interactions
- **Agents:** Purpose-built autonomous agents including **Amazon Quick, AWS Transform**, **Rufus,** and **Alexa+**.

AWS AI: The Monetization Framework:

Amazon's AWS has strategically positioned itself to capture value across the entire generative AI stack from Infrastructure (Trainium/Inferentia) to Tooling (Bedrock) to Applications (Q). Amazon is diversifying its monetization beyond simple cloud compute into high-margin software and proprietary silicon.

Figure 42. Amazon Monetization Framework



Source: AWS website

Monetization Summary

AWS is utilizing a “Land and Expand” strategy through these three layers:

1. **Bottom layer: Infrastructure or the foundation.** High-volume, utility-style revenue. Amazon competes on price and availability here.
2. **Bedrock: The Multiplier.** High-margin platform revenue. As more companies build on Bedrock, AWS captures a percentage of all downstream AI activity.
3. **Amazon Q: The High-Margin SaaS.** High-margin, recurring per-seat revenue. This layer aims to move AWS up the stack from “back-end infrastructure” to a “front-end productivity tool.”

Amazon’s Approach to Agentic Commerce: Defending the-everything-store while building its own agents to purchase on-and-off the platform:

Amazon is approaching Agentic Commerce with a focus on its own agent which is tailored to Amazon’s marketplace, Rufus. In 2025, Rufus was used by 300M+ users and generated ~$12B in incremental annualized revenues. As importantly, customers utilizing Rufus converted at a clip that was ~60% higher than those that do not. Beyond Rufus, Amazon is also in the early innings of building out its “Buy for Me” agent to shop third party sites, turning Amazon into a meta-agent that can purchase both inside and outside its walled garden. As it pertains to other third-party agents, Amazon has taken a firm stance to disallow third-party agents to make purchases on behalf of users across its platform so far. Specifically, the company implemented a change to its backend “robots.txt” files to effectively shoutout 3rd party shopping agents from Anthropic (Claude), OpenAI (ChatGPT), Google and Perplexity from scraping data across the platform.

Partnerships with OpenAI and Anthropic.

Amazon has struck several strategic partnerships across the Gen AI landscape with some of the standouts being its partnerships with OpenAI (private), and Anthropic:

- **Anthropic:** Amazon and Anthropic’s partnership has evolved from a minority investment into a strategic cornerstone of Amazon’s generative AI and cloud infrastructure roadmap with Anthropic being a key capacity customer for AWS and an important contributor to Amazon’s suite of Agentic Offerings.
    - **Project Rainier:** On the infrastructure side, Anthropic is training its next-gen frontier models on a cluster of 500,000+ (and growing) AWS Trainium2 chips.
    - Beyond the infrastructure deals, Anthropic models are available on Amazon Bedrock for AWS customers, and Amazon is utilizing various Claude models for its own internal purposes.
    - **Level of Investment:** As of year-end 2025, Amazon has deployed $8 billion in total cash into Anthropic, and the fair value of its combined holdings (preferred stock and convertible notes) reached $62.0 billion as of 2025 YE.

- **OpenAI:** In February, Amazon announced that it has struck a multi-year partnership with OpenAI. As part of the partnership, AWS will become the exclusive third-party cloud distribution provider for OpenAI Frontier, an enterprise platform to enable organizations to build and manage AI agents across systems. Additionally, OpenAI and AWS are set to jointly develop a Stateful Runtime Environment powered by OpenAI models which will be available to AWS customers on Amazon Bedrock.
    - **Infrastructure Commitment:** This deal expands the companies' existing relationship to $138 billion over eight years, building upon the companies initial $38B relationship, with OpenAI committing to consume 2 gigawatts of capacity powered by Amazon's proprietary Trainium3 and Trainium4 AI chips.
    - **Level of Investment:** Amazon is set to invest $50B into OpenAI with a $15B upfront investment to be followed by an additional $35B investment when conditions are met.

Planning to Aggressively Add Capacity

AWS has aggressively scaled capacity, adding 3.9GW of power in the last twelve months (as of C4Q25) and 1.2GW of power in 4Q25 alone. Looking forward, management expects to double the company's capacity by the end of 2027 with a further increase on the horizon in 2028. Underpinning this expansion is a significant ramp in CapEx spend (AMZN guided to $200B in CapEx in 2026 with the vast majority to be allocated to Gen AI at AWS). It's important to highlight that management has said that it's fulfilling supply as soon as the latter is brought online, which points to the significant pen-up demand it's seeing for its Gen-AI related services.

Meta – From Social Media to Personal Superintelligence

Summary: *Meta is betting on Personal Superintelligence and hybrid-source dominance via Llama and Muse Spark, using smart glasses as the primary hardware interface for the agentic era. Meta's move to build gigawatt-scale clusters (Prometheus, Hyperion) and its various partnerships underscore a shift toward a massive "compute-first" strategy, intended to subsidize its goal of becoming the industry standard for agentic execution.*

Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	Open-Sourced LLM, consumer services
Role	Vertical AI Infrastructure Provider
User	Developers and researchers, companies of all sizes (Enterprises, SMB, Startups) and government agencies

Financial Snapshot ($Ms)					
Mkt. Cap	**$1,713,515**	**Net Debt**	**$2,652**	**Enterprise Value**	**$1,716,167**

Revenue	META does not provide cloud services to third party nor does it generate money (yet) from its AI models, which are open-sourced. Nearly all of Meta's revenue today are generated from its core ads business. Its breakthroughs across its LLms and AI architecture are powering engagement and monetization across its Family of Apps. Unlike GOOGL and AWS, Meta does not directly generate revenue from 3P use of its open-source models.
Most Recent Performance Updates & LT Targets	**Customer Silicon/GPU's:** In 4Q25 Meta doubled the number of GPUs it uses to train its Gem model for ads ranking, which helped drive a 3.5% lift in ad clicks on Facebook and a more than 1% gain in conversations on Instagram. **Unified Architecture**: Under the company's unification efforts, Meta consolidated models for Facebook stories and other services into the overall Facebook model, which drove a 12% increase in ads quarterly. **Advantage+:** Meta's fully automated end-to-end solutions reached $160B ARR in 3Q25. **AI Creative:** Meta's AI ad creative revenue run rate for video generation tools hit $10 billion in Q4, with q/q growth outpacing increase in overall ads revenue by nearly 3x.

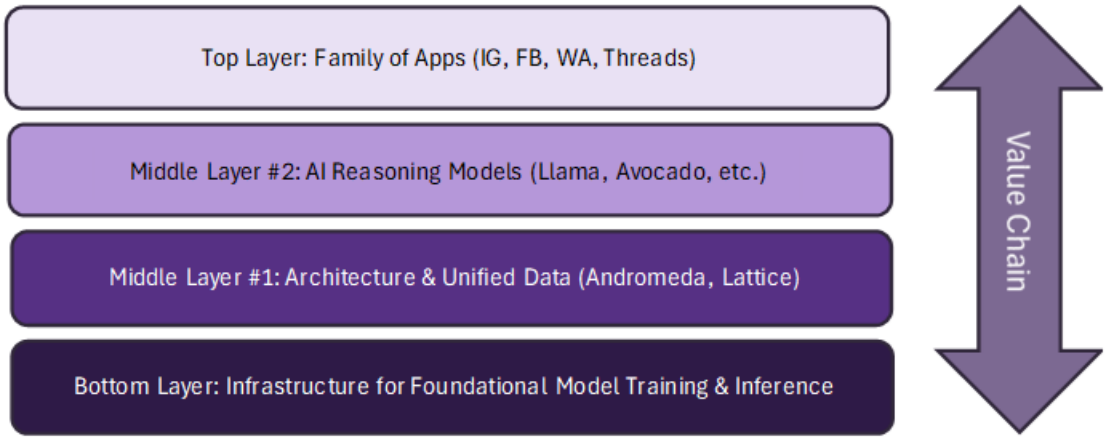
Source: FactSet – Financial Data as of 4/27/2026

Meta's AI Framework: Building an Open Source-based End-to-End Stack

Meta's position in the agentic AI stack has been unique because it is the only hyperscaler currently pursuing a "Linux of AI" strategy, given its open-source nature and the way it looks to democratize AI development. While OpenAI, Google, and Microsoft are building walled garden agent ecosystems, Meta has positioned itself as the foundational infrastructure layer that powers everyone else's agents, offering developers flexibility, control, and cost efficiency to build and customize their own AI systems. The strategic significance lies in Meta's ability to commoditize the underlying models while capturing value through high-engagement agentic interfaces across its Family of Apps (Instagram, WhatsApp, Messenger). Meta's approach has been to provide the open Llama stack for the developer community while vertically integrating its

own hardware (MTIA) and energy infrastructure to lower the long-term cost of serving billions of proprietary AI agents. More recently, Meta introduced Muse Spark, the first model from its newly formed Meta Superintelligence Labs (MSL). Muse Spark marks a departure from Meta's traditional "open-weights" strategy, as the model is proprietary and closed-source, reflecting a shift toward high-performance, specialized infrastructure.

Figure 43. Meta AI & Agentic Layers



Source: Meta Website

Meta is in the process of building out the entire vertical stack, operating a hybrid hardware strategy (its own silicon chips and chips from AVGO, ARM, AMD, and NVDA), while building new superintelligence models (like Muse Spark), retrieval engines (Andromeda), and leveraging these in its consumer facing services (IG, FB, WA, Reels, Threads). Below is summary view of META's AI portfolio currently:

Figure 44. Meta AI Portfolio

Foundational AI Models & Generative Media			
Model	Description	Key Insight	Target Audience
Llama 4 Scout	Highly efficient model with a 10M-token context window, analyzes code and documents	Optimized to run on a single H100 GPU; handles multi-document reasoning	Developers
Llama 4 Maverick	Flagship 400B parameter natively multimodal model	Multimodal intelligence that prioritizes complex reasoning and complex tool orchestration	Enterprises
Llama 4 Behemoth (delayed)	2T total parameter "teacher" model	Not yet rolled out publicly, Meta previously expected Behemoth to be the most capable model developed	Researchers and Engineers
Llama 3 / 3.1 / 3.2 / 3.3	Basic, lost cost models	Still widely used for smaller, on-device tasks and mid-range applications	Developers
Muse Spark	Meta's first natively multimodal reasoning flagship model, released April 2026	Features advanced "visual chain-of-thought" and excels in healthcare reasoning, marking a shift toward "personal superintelligence"	
SAM (Segment Anything Model)	Updated Segment Anything Model for image and real-time video	Industry standard for computer vision; integrated into Amazon SageMaker	Creators, Researchers, and Enterprises
GEM (Generative Ads Model)	Foundational model for recommendations	Learns from billions of cross-platform user behaviors, delivers 4x efficiency of previous ranking models	Advertisers
Andromeda	Matching algorithm that shifts targeting from audience-first to asset-first	Uses visual and sentiment analysis to find the correct audience for an ad based on the creative itself, rather than demographic settings	Agencies and Brands for full automation
Lattice	Multimodal "World Model" specializing in image and video generation	Complex understanding of physics and object dynamics	Creators, media producers

Source: Meta company reports

Llama

The Llama family has been Meta's flagship series of AI models. Llama 4 (released April 2025) significantly shifted Meta's AI architecture towards native multimodality and MoE (mixture-of-experts). Meta's open-sourced (but increasingly hybrid) strategy boils down to its belief that by sharing its resources, it can accelerate AI research, improve its core ad products, and foster innovation across the tech community that will benefit the company long-term. This enables startups and large enterprises to build, fine-tune, and deploy their own applications using Meta's Llama models. While this approach is less capital intensive than closed-sourcing it, Meta has stated that upcoming models focused on achieving superintelligence could be closed-sourced instead (its latest model, Muse Spark, discussed below highlights this point). This would allow Meta to generate revenue from its new frontier models, similar to rivals' closed-sourced strategy currently (i.e. Gemini, Claude, GPT).

Muse Spark, the latest and greatest: Recently, Meta introduced Muse Spark, the first model from its newly formed Meta Superintelligence Labs (MSL), which marks a significant technical pivot toward natively multimodal reasoning. Unlike previous Llama models, Muse Spark was built from the ground up to process text, images, and audio simultaneously, enabling it to act as a "personal superintelligence" that understands a user's physical environment. Notably, Muse Spark marks a departure from Meta's traditional "open-weights" strategy, as the model is proprietary and closed-source, reflecting a shift toward high-performance, specialized infrastructure. The model features three distinct interaction modes: "Instant" for quick tasks, "Thinking" for complex step-by-step reasoning, and "Contemplating" for advanced multi-agent orchestration (similar features to other leading gen-AI platforms). Key user-facing capabilities include deep health integration and visual coding features that allow users to generate functional mini-games and websites through simple prompts. This launch represents the first major output following Meta's $14 billion deal to bring in Alexandr Wang, the founder of Scale AI (private), as Meta's Chief AI Officer replacing Yann LeCun who departed to launch his own start-up.

Monetization:

There are several ways Meta is currently monetizing and plans on monetizing the AI opportunity over time. Near-term, the company remains focused on leveraging its AI models to improve its core ads business. Longer term, Meta plans on continuing to introduce superior, superintelligence models that it could keep closed-source and license to enterprises.

1. **Improving the core ads business**: Meta has fundamentally restructured its advertising business around agentic AI, moving the platform away from manual "button-clicking" toward a fully automated ecosystem. The company has leveraged its architecture and unified data layer (Andromeda, Lattice, and GEM) to reduce CPA, increase clicks, boost ad quality and relevance, and increase user engagement through better ranking and recommendations. Additionally, Meta has also scaled its fully automated, end-to-end campaign set up tool, Advantage+, to a $60B ARR as of 3Q25, which streamlines the ad buying process and unlocks more budgets from advertisers.

2. **The superintelligence era:** While Muse Spark is Meta's first attempt at doing closed-source model which it can license like its rivals do, the company is working on additional

ground-breaking models to achieve CEO Mark Zuckerberg's vision of superintelligence, where models and AI Assistants can plan, reason, and execute complex tasks on users and advertisers' behalf.

Acquisitions and Partnerships to scale AI

Meta has aggressively acquired startups that focus on autonomous agents and custom silicon to reduce its dependence on external providers.

- **Moltbook** (March 2026): Meta recently confirmed the acquisition of Moltbook, a social network platform designed specifically for AI agents. Its co-founders, Matt Schlicht and Ben Parr, joined Meta's newly formed Superintelligence Labs (MSL).
- **Manus AI** (January 2026): In a deal valued at over $2 billion, Meta acquired Singapore-based Manus, a developer of general-purpose AI agents. Manus agents are capable of executing end-to-end complex workflows like market research and coding, rather than just generating text.
- **Scale AI** (June 2025): Meta led a historic investment of over $14 billion in Scale AI, valuing the company at $29 billion (per Reuters). As part of this strategic move, Scale AI founder Alexandr Wang joined Meta as Chief AI Officer to lead the Superintelligence Labs division, while remaining on Scale AI's board.
- **Rivos** (October 2025): Meta acquired this chip startup for approximately $2 billion. Rivos specializes in power-efficient, custom server silicon, which is critical for Meta's goal of building its own internal compute stack

Meta is diversifying its hardware supply chain to support its massive model training requirements.

- **AMD (AMD, Buy, Stein)** (February 2026): Meta and AMD announced a landmark 6-gigawatt (GW) agreement. Meta will be the lead customer for custom AMD Instinct MI450-based GPUs and 6th Gen EPYC "Venice" CPUs. As part of the deal, AMD issued Meta warrants for up to 160 million shares, aligning their long-term financial interests.
- **Broadcom (AVGO, Buy, Stein)**: Meta revealed that its latest generation of MTIA (Meta Training and Inference Accelerator) chips was co-developed in close partnership with Broadcom.
- **Arm (ARM, NR)** (March 2026): Meta is a primary partner for Arm's new AGI CPU initiative, focusing on "agentic data center design" that optimizes power for the specific needs of autonomous AI agents.
- **Nvidia**: Despite diversifying, Meta remains one of Nvidia's largest customers, with ongoing multi-billion-dollar commitments.
- **CoreWeave (CRWV, Hold, Ramnani)** (April 2026). Meta and CRWV (Hold, Ramnani) announced an expansion of their AI cloud partnership in a deal valued at ~$21B for CRWV provide Meta with compute capacity to support the development/deployment of next GenAI models.

To fund its CapEx plans ($115–$135B for 2026), Meta has turned to innovative PE structures.

Partner	Project	Details
Blue Owl Capital	Hyperion Data Center	A $30 billion joint venture for a massive campus in Louisiana. Blue Owl/PIMCO owns 80%, while Meta retains 20% and manages operations.
PIMCO	Hyperion Debt	Acted as the anchor lender for the $27 billion debt portion of the Louisiana project, which is the largest private capital transaction in history.
Blackstone	Infrastructure Strategy	While not the primary lead on Hyperion, Meta is in ongoing talks with Blackstone regarding future gigawatt-scale sites in the Midwest to support its "AI-first" power requirements.

Source: Meta company reports

Artificial Intelligence, Digital Platforms, & IT Services - Ramnani

Palantir – The Decision Layer Connecting Agentic AI to Institutional Knowledge

Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	Enterprise data infrastructure, agent runtime, orchestration, governance
Role	Verticalized AI operating system and control panel for agentic execution
User	Enterprises, governments, regulated industries, operational teams

Financial Snapshot ($Ms)					
Mkt. Cap	**$342,225**	**Net Debt**	**($6,950)**	**Enterprise Value**	**$335,275**
Revenue			FY25 Total Revenue: $4.5B + 56% Y/Y, Government Revenue: $2.4B +54% y/y Commercial Revenue: $2.1B +60% y/y		
Most Recent Performance Updates & LT Targets			**FY26E Revenue:** $7.2B at mdpt +61% y/y FY26E **Adj. Op. Income:** $4.1B at mdpt, 57.5% operating mgn Deals: US Navy: $448M contract to deploy Foundry & AIP, **Airbus:** Multi-year contract extension, **Polymarket & TWG AI**: Partnered to build sports integrity platform **LG:** Expanded strategic partnership to accelerate AI initiatives **GE Aerospace:** Multi-year expansion to deploy agentic AI solutions **Stellantis:** 5 year contract renewal & expansion		

Source: FactSet – Financial Data as of 4/27/2026

Palantir (PLTR, Buy, Ramnani) is positioning itself as a core infrastructure provider in the agentic AI stack, focused not on model intelligence but on enabling safe, governed, and scalable autonomy inside enterprises. While the company leverages frontier models (including OpenAI, Anthropic, and others), its differentiation lies in the software layer that supplies agents with structured business context, enforces permissioned action, and enables closed-loop execution across enterprise systems.

In an agentic AI world, we believe value will accrue to platforms that can translate probabilistic reasoning into deterministic outcomes under real-world production constraints. This includes managing how agents reason over enterprise data, determine allowable actions, execute across systems of record, and provide observability and auditability throughout workflows. Successful agent outcomes in Palantir-enabled environments are increasingly measured not by model accuracy alone, but by operational metrics such as decisions executed, workflows automated, disruptions mitigated, resource allocation efficiency, and time-to-resolution.

Monetization and Strategic Leverage

Palantir operates a highly customized, hybrid software model and monetizes primarily through long-term enterprise software contracts spanning commercial and government customers. Pricing for its software platforms is typically customized and tied to variables such as deployment scope, workload intensity, expansion over time, and ongoing engineering support – increasingly reflecting outcome-driven pricing rather than traditional seat-based pricing.

As agentic AI adoption broadens, we believe Palantir is positioned to benefit from increasing workload intensity as customers move from analytics and decision support to fully autonomous, closed-loop operations. While much of the broader agentic AI ecosystem today remains consumption-based (tokens, inference, or per-agent pricing), Palantir's model is already oriented toward outcome-based expansion – where incremental revenue is driven by additional operational use cases, automation breadth, and organizational dependence on the platform. We believe this structure is well aligned for an environment in which autonomous agents replace human-centric workflows, potentially compressing traditional SaaS seat-based economics, but expanding platform-level value capture.

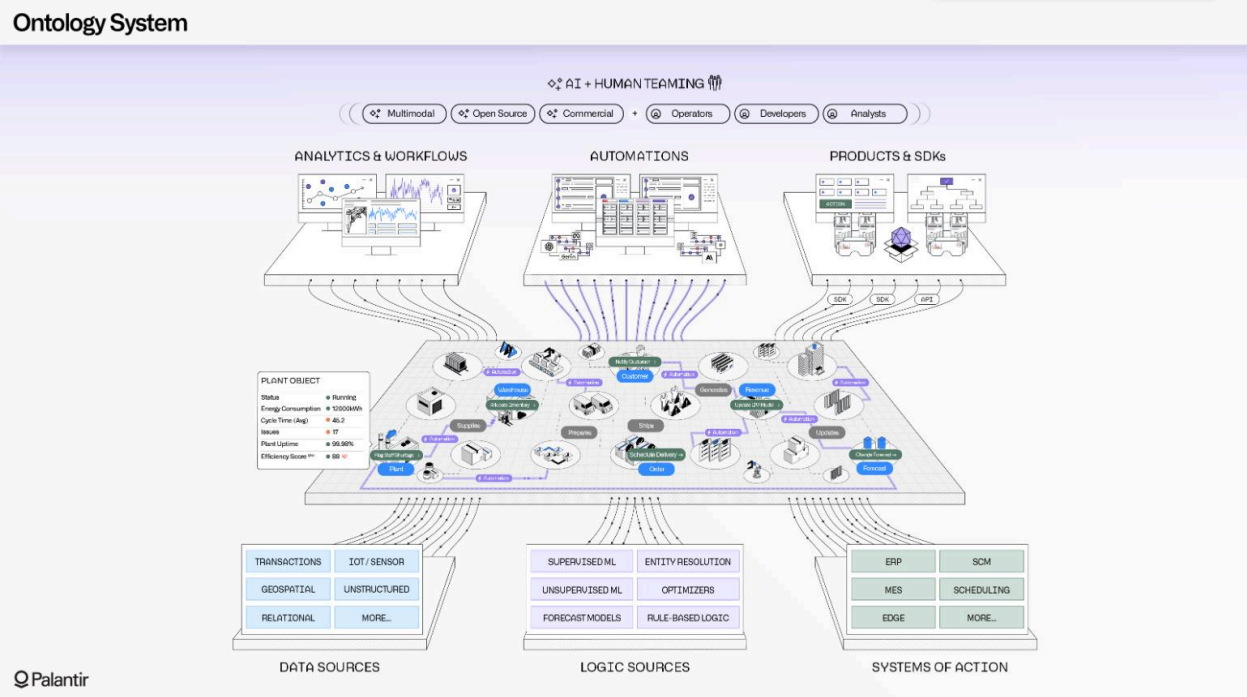
AIP and the Ontology Advantage: From Decision Support to Autonomous Action

Palantir's Artificial Intelligence Platform (AIP) is designed as an agentic orchestration layer and operating system for mission-critical environments. At its core, AIP enables enterprises to build, deploy, and manage agents that operate as persistent control loops: reasoning over structured business context, executing actions across systems, validating outcomes, and iterating over time. A key architectural feature is Palantir's Ontology, which provides a real-time digital twin of the enterprise – modeling business objects, relationships, constraints, and permissible actions.

While Palantir's Ontology has historically provided a significant competitive advantage in contextualizing disparate data to improve enterprise decision-making processes, we think the Ontology's benefits will compound by providing agents with real-time, interconnected context across entire organizations. We view agentic AI as an enterprise re-architecture, not just a simple software download, and the Ontology likely provides much of the upfront heavy lifting required for agentic implementation. Agents interact with the enterprise through this Ontology rather than raw data tables, allowing autonomy to be tightly coupled to policy, permissions, and business context.

AIP supports a spectrum of agentic capabilities, ranging from task-specific assistants to fully automated agentic applications that can read and write data, trigger workflows, and coordinate across operational systems. Importantly, Palantir embeds security, governance, and observability directly into agentic runtimes, ensuring agents operate under the same access controls, compliance frameworks, and audit requirements as human users.

Figure 45. Palantir Ontology



Source: Palantir company presentation

Core capabilities enabled by AIP include:

- **Agent Development:** Tools to build agents that plan tasks, reason over enterprise context, and interact with internal tools and applications.
- **Ontology-Based Context:** A structured, real-time digital representation of enterprise objects and workflows, providing critical business context to agents.
- **Operational Execution:** Support for multi-step, closed-loop workflows that span planning, execution, validation, and iteration.
- **Security, Governance, and Auditability:** Precise access controls, lineage tracking, monitoring, and post-hoc explainability of agent behavior.
- **Performance Monitoring and Control:** Instrumentation to evaluate agent outcomes and refine automation safely over time.

Figure 46. Highlighting US Commercial Becoming a Material Growth Driver Post-AIP



Sources: Company Reports, Truist Securities
Our estimates imply ~76% US Commercial 4-year CAGR vs. ~40% for total revenue
Revenue figures in millions

Model-Agnostic Approach Provides Leverage and Reduces Dependency Risk

From an infrastructure perspective, Palantir is model-agnostic and cloud-agnostic. The company does not train foundational models itself and instead integrates both commercial and open-weight models into AIP, providing users with model optionality. This reinforces management's view that models are becoming interchangeable inputs, while durable value migrates to the orchestration layer. Palantir similarly deploys across multi-cloud environments (from hyperscalers to on-prem and edge computing), allowing customers to run agentic workloads within their preferred infrastructure and security boundaries.

As agentic systems scale, we believe Palantir's multi-model and multi-cloud approach enhances its strategic leverage and reduces platform risk for customers. The fallout from the Anthropic/Department of War contract dispute – which we've discussed extensively in our *Weekly Inferences* – highlighted the structural risks that can arise when AI systems are tied to a single model provider (Anthropic was the only one approved for use in classified environments). In our view, Palantir's model-agnostic architecture is explicitly designed to mitigate this risk to enable enterprise and government customers to maintain continuity of AI systems and model flexibility as access or partner relationships change.

Go-to-Market: AIP Bootcamps Accelerating Agentic Adoption Timelines

A distinguishing element of Palantir's go-to-market strategy is its AIP Bootcamp model, which compresses the transition from experimentation to production. These intensive, multi-day engagements are designed to move organizations from zero to deployed agentic workflows within days, emphasizing real enterprise data, live operational constraints, and immediate deployment. We view this approach as particularly effective in catalyzing agentic adoption, as it forces customers to confront governance, integration, and execution challenges early rather

than deferring them. Over time, this model appears to be driving broader organizational buy-in and accelerating the expansion of agentic use cases across customers.

Accenture – Rewiring the Enterprise for the Agentic Era

Where They Sit in the Agentic Stack as a Picks & Shovels Play	
Layer	IT Services, third party supplier of consulting and managed services tech talent
Role	Middle "glue" layer tailoring foundational LLM's to specific enterprise needs and workflows via partnership such as Accenture's Anthropic Business Group and its OpenAI Frontier Alliance partnership
User	Enterprises including those in tech, healthcare, financial services, products, resources, and other regulated industries

Financial Snapshot ($Ms)					
Mkt. Cap	**$109,502**	**Net Debt**	**($3,302)**	**Enterprise Value**	**$106,200**

Revenue	ACN reported FY25 (August) rev of $69.7B, representing 7.4% y/y growth. Revenue momentum continued into FY26 with Q1 and Q2 annual growth rates of 6% and 8.4% respectively. In FY25, ACN delivered adj op mgns of 15.6% and EPS of $12.93 (+8.2% y/y).
Most Recent Performance Updates & LT Targets	**Employees:** 786K employees as of 2Q26 (+10.3% y/y). **FY26 Financial Targets:** $73.2B - $74.6B (5% - 7% growth y/y); Adj. op mgn of 15.7%; Adj. EPS of $13.52 - $13.90. **Growth Vectors**: >60% of revenue stems from top 10 ecosystem partners including hyperscalers and enterprise software vendors. Bookings from AI partnerships include PLTR, Anthropic, OpenAI, & Databricks are expected to double y/y.

Source: FactSet – Financial Data as of 4/27/2026

Monetization and Strategic Leverage

Accenture generates revenue through professional services spanning strategy, consulting, technology, and managed services. In our view, ACN serves as the premier services picks and shovels provider in the agentic era. The company is aggressively shifting its monetization model from pure labor-based deployments to Reinvention Services, a specialized unit launched in late 2025 to embed AI directly into the digital core of enterprise clients. Accenture’s role in AI, along with other IT Services vendors like EPAM Systems (EPAM, Hold, Ramnani) and Globant (GLOB, Hold, Ramnani), is as a middle glue layer where frontier model companies like OpenAI or Anthropic provide the raw intelligence and capability, and Accenture provides the integration, security, and process reimagining required to make these models functional in regulated, legacy environments. We believe Accenture’s primary moat lies in the last mile of AI implementation, where it provides industry and client-specific context by utilizing a library of 3,000+ reusable agents that bridge the gap between foundation models and measurable P&L outcomes across verticals.

AI Refinery: Accenture's Enterprise Agent Platform

Launched in January of 2025, Accenture's AI Refinery is a proprietary platform designed to help organizations build, deploy, and manage customized AI agents at scale. Specifically, the initiative focuses on distilling foundational models, taking horizontal intelligence and fine-tuning it with proprietary enterprise data to execute complex, multi-step tasks. The platform includes the Distiller Agentic Framework, a suite of SDK's that allows business users to construct advanced reasoning agents without deep coding expertise, facilitating rapid prototyping for clients across sectors. By integrating these agents with digital twins and physical AI, Accenture is extending agentic workflows into physical industrial settings with AI systems that process real world signals such as that from cameras and sensors.

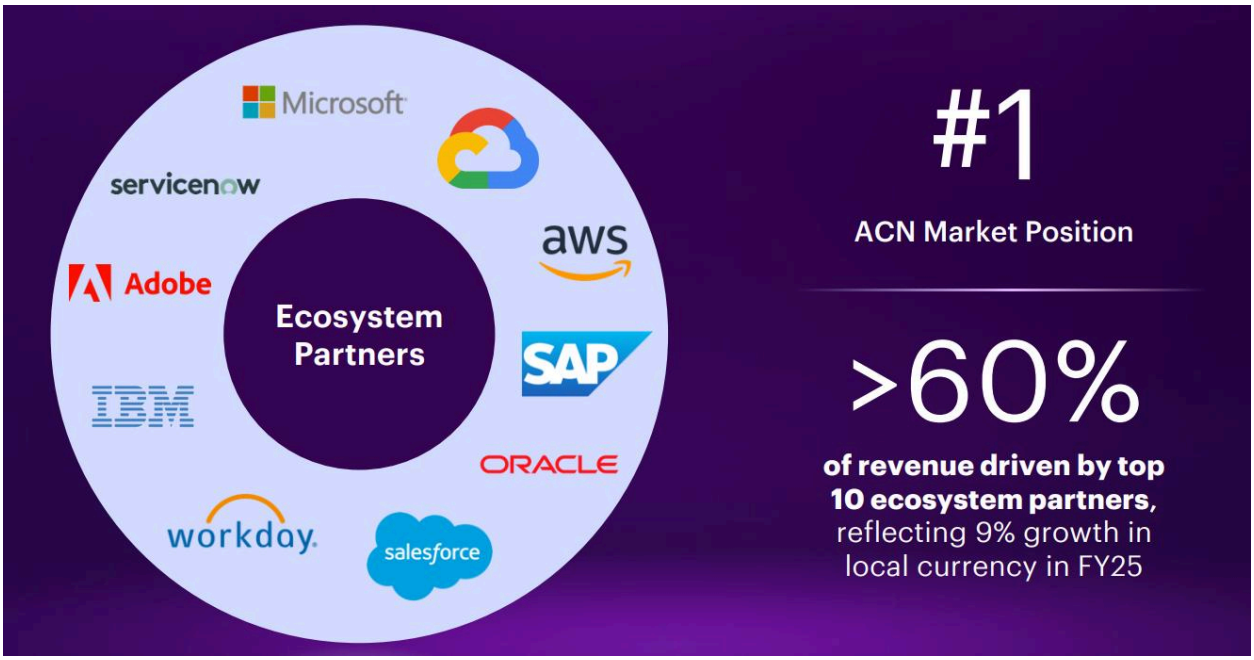
Strategic Ecosystem: Model Agnostic Approach Allows for Versatile Toolkit

Accenture's partnership strategy is grounded in being model agnostic, enabling clients to navigate an increasingly fragmented AI ecosystem and select the tech stack that is the best fit for each individual client. We note that >60% of ACN's revenue stems from its top 10 partnerships (revenue growth is outpacing that of total company) and that the company has been rapidly scaling its partnerships with frontier AI companies. We believe the company is on track to double its bookings from emerging AI and data partnerships in FY2026 including Palantir (PLTR, Buy, Ramnani), Anthropic (Private), OpenAI (Private) and Databricks (Private), leveraging deep-level integrations with both hyperscalers and frontier labs.

- **Nvidia:** Established the Nvidia Business Group, a dedicated unit of 30,000 professionals focused on deploying agentic AI specifically using the Nvidia stack and AI Refinery platform.
- **Mistral AI (Private):** Announced a multi-year strategic collaboration in February 2026 to deliver sovereign, scalable AI solutions for European and global enterprises seeking strategic autonomy.
- **Microsoft & AWS:** Integrated agentic architecture into M365 Copilot and AWS Bedrock, facilitating workflows that allow agents to manage files, ERP data, and communications autonomously. Late last year, the partnership was expanded to specifically target the modernization of public service organizations through generative and agentic AI.

Palantir: Expanded their global partnership in December 2025, creating a dedicated business group to help clients utilize Palantir's Artificial Intelligence Platform (AIP) for secure, large-scale data reinvention.

Figure 47. Accenture Key Ecosystem Partners



Source: Accenture company presentation

M&A Strategy: Scaling the "AI Enabler" Footprint

Accenture is acquisitive, and its inorganic growth strategy is a key component of its overall top line growth trajectory. We note that for FY26, ACN is targeting 3%-5% of growth in local currency (5%-7% including FX tailwinds), with inorganic contributions of ~1.5%. Accenture has committed $5.0 Billion in capital for acquisitions in FY2026, targeting AI enablers, the physical and digital infrastructure required to power agentic workloads. A core component of this strategy is the shift toward non-FTE revenue, acquiring firms with high-margin software or subscription models to augment traditional consulting.

Accenture has recently targeted acquisitions that will aid in driving its AI-related revenue. In March 2026, Accenture completed the $1.2 Billion acquisition of Ookla, a leader in network intelligence; this high-margin subscription business ($231M revenue) is critical for enterprises scaling AI across global data networks. Simultaneously, the firm completed the acquisition of Faculty in the United Kingdom, an AI-native firm specializing in decision intelligence, which provides a platform for software-like revenue models in unmet AI demand areas. Other key acquisitions include a majority stake in DLB Associates (December 2025) to aid in establishing a foothold in the data center engineering market, and the purchase of CyberCX, marking the firm's largest-ever cybersecurity acquisition to secure the agentic stack in Asia-Pacific.

Key Metrics to Track

- Revenue per Employee Emerging as a Key Metric: Revenue per employee is becoming an important metric in IT Services due to the changing nature of billing structures as AI becomes a key component of scope and as the traditional billable hour model shifts to more of an outcome-based dynamic. Growing revenue per employee indicates not only a decoupling of revenue from headcount growth but suggests that outcome-focused AI contracts are scaling, and services firms are successfully managing AI-enabled headcount in an environment that rewards value and outcomes. In FY25, ACN's revenue per head grew 1.9%, but we note that in the last three quarters, growth has ranged between 6.6% and 10.3% per employee.
- Bookings Indicates AI Traction: The record $22.1B in quarterly bookings (Q2 FY2026) suggests that the reinvention narrative is successfully capturing large-scale, multi-year transformational budgets rather than just small-scale pilots. ACN's target of doubling bookings in FY26 from its rapidly scaling AI partnerships is another data point that exemplifies this theme.
- Margin Expansion through Non-FTE Growth: The aggressive acquisition of subscription and licensing businesses like Ookla and Faculty signals a strategic pivot toward software-like margins, potentially de-risking the business from the constraints of a purely labor-based model.
- The Agentic Labor Force: With 192,000 employees already trained in the Agentic AI Fundamentals program and a total pool of 85,000 AI professionals, Accenture is building the largest certified labor force capable of deploying the next wave of autonomous enterprise agents.

Disclosure: One of the research analysts of this report owns an equity ownership interest of less than 1%, in Anysphere, Inc., parent company of Cursor, through a limited liability corporation (LLC).

Citations

1) McKinsey, Press Release, April 28, 2025, “The cost of compute: A $7 trillion race to scale data centers” https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/the-cost-of-compute-a-7-trillion-dollar-race-to-scale-data-centers

2) Gartner®,Press Release, October 29, 2025, "Gartner Survey Finds 45% of Martech Leaders Say Existing Vendor-Offered AI Agents Fail to Meet Their Expectations of Promised Business Performance" https://www.gartner.com/en/newsroom/press-releases/2025-10-29-gartner-survey-finds-45-percent-of-martech-leaders-say-existing-vendor-offered-ai-agents-fail-to-meet-their-expectations-of-promised-business-performance#:~:text=Half%20of,agent%20implementati. GARTNER is a trademark of Gartner, Inc. and its affiliates.

3) McKinsey, Press Release, November 5, 2025, “ The state of AI in 2025: Agents, innovation, and transformation” https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai

4) Gartner®, Press Release, March 11, 2024, "Gartner Predicts One-Third of Interactions with GenAI Services Will Use Action Models & Autonomous Agents for Task Completion by 2028" https://www.gartner.com/en/newsroom/press-releases/2024-03-11-gartner-predicts-one-third-of-interactions-with-genai-services-will-use-action-models-and-autonomous-agents-for-task-completion-by-2028

5) Gartner®, Press Release, January 15, 2026, Gartner Says Worldwide AI Spending Will Total $2.5 Trillion in 2026 https://www.gartner.com/en/newsroom/press-releases/2026-1-15-gartner-says-worldwide-ai-spending-will-total-2-point-5-trillion-dollars-in-2026#:~:text=Worldwide%20spending,to%20Gartner

Companies Mentioned in This Note

Accenture plc (ACN, $176.97, Buy, Arvind Ramnani)
Advanced Micro Devices, Inc. (AMD, $334.63, Buy, William Stein)
Amazon.com, Inc. (AMZN, $261.12, Buy, Youssef Squali)
Broadcom Limited (AVGO, $418.20, Buy, William Stein)
Salesforce, Inc. (CRM, $180.18, Buy, Terry Tillman)
CrowdStrike Holdings, Inc. (CRWD, $454.61, Buy, Junaid Siddiqui)
CoreWeave, Inc. (CRWV, $112.06, Hold, Arvind Ramnani)
Commvault Systems, Inc. (CVLT, $88.41, Buy, Junaid Siddiqui)
Datadog, Inc. (DDOG, $132.66, Hold, Miller Jump)
Dynatrace, Inc. (DT, $35.61, Buy, Miller Jump)
EPAM Systems, Inc. (EPAM, $116.31, Hold, Arvind Ramnani)
Elastic N.V. (ESTC, $47.83, Buy, Miller Jump)
Globant S.A. (GLOB, $42.40, Hold, Arvind Ramnani)
Alphabet Inc. (GOOGL, $350.34, Buy, Youssef Squali)
Intapp, Inc. (INTA, $22.14, Buy, Terry Tillman)
Klaviyo, Inc. (KVYO, $19.46, Buy, Terry Tillman)
Manhattan Associates, Inc. (MANH, $137.84, Buy, Terry Tillman)
Meta Platforms, Inc. (META, $678.62, Buy, Youssef Squali)
Microsoft Corporation (MSFT, $424.82, Buy, Terry Tillman)
nCino, Inc. (NCNO, $17.59, Hold, Terry Tillman)
Cloudflare, Inc. (NET, $212.36, Buy, Junaid Siddiqui)
ServiceNow, Inc. (NOW, $90.45, Buy, Miller Jump)
NVIDIA Corporation (NVDA, $216.61, Buy, William Stein)
Okta, Inc. (OKTA, $76.14, Buy, Junaid Siddiqui)
Palo Alto Networks, Inc. (PANW, $182.90, Buy, Junaid Siddiqui)
Palantir Technologies Inc. (PLTR, $143.10, Buy, Arvind Ramnani)
Qualys, Inc. (QLYS, $85.16, Hold, Junaid Siddiqui)
Q2 Holdings, Inc. (QTWO, $48.56, Buy, Terry Tillman)
Rubrik, Inc. (RBRK, $53.26, Buy, Junaid Siddiqui)
Rapid7, Inc. (RPD, $5.98, Hold, Junaid Siddiqui)
SailPoint, Inc. (SAIL, $11.87, Buy, Junaid Siddiqui)
Shopify Inc. (SHOP, $124.23, Buy, Terry Tillman)
Snowflake, Inc. (SNOW, $144.25, Buy, Miller Jump)
Tenable Holdings, Inc. (TENB, $20.47, Buy, Junaid Siddiqui)
ServiceTitan, Inc. (TTAN, $60.28, Buy, Terry Tillman)
Varonis Systems, Inc. (VRNS, $25.05, Buy, Junaid Siddiqui)
Zeta Global Holdings Corp. (ZETA, $18.09, Buy, Terry Tillman)
Zscaler, Inc. (ZS, $134.11, Buy, Junaid Siddiqui)
Anthropic (Private)
OpenAI (Private)
Databricks (Private)
Cursor (Private)
Oracle (ORCL, NR)
Cerebras (Private)
Arm Holdings (ARM, NR)
Icertis (Private)
Mistral AI (Private)
Wiz (Private)

Analyst Certification

I, Terry Tillman , hereby certify that the views expressed in this research report accurately reflect my personal views about the subject company(ies) and its (their) securities. I also certify that I have not been, am not, and will not be receiving direct or indirect compensation in exchange for expressing the specific recommendation(s) in this report.

I, Miller Jump , hereby certify that the views expressed in this research report accurately reflect my personal views about the subject company(ies) and its (their) securities. I also certify that I have not been, am not, and will not be receiving direct or indirect compensation in exchange for expressing the specific recommendation(s) in this report.

I, Junaid Siddiqui , hereby certify that the views expressed in this research report accurately reflect my personal views about the subject company(ies) and its (their) securities. I also certify that I have not been, am not, and will not be receiving direct or indirect compensation in exchange for expressing the specific recommendation(s) in this report.

I, Youssef Squali , hereby certify that the views expressed in this research report accurately reflect my personal views about the subject company(ies) and its (their) securities. I also certify that I have not been, am not, and will not be receiving direct or indirect compensation in exchange for expressing the specific recommendation(s) in this report.

I, Arvind Ramnani , hereby certify that the views expressed in this research report accurately reflect my personal views about the subject company(ies) and its (their) securities. I also certify that I have not been, am not, and will not be receiving direct or indirect compensation in exchange for expressing the specific recommendation(s) in this report.

Required Disclosures

Analyst compensation is based upon stock price performance, quality of analysis, communication skills, and the overall revenue and profitability of the firm, including investment banking revenue.

As a matter of policy and practice, the firm prohibits the offering of favorable research, a specific research rating or a specific target price as consideration or inducement for the receipt of business or compensation. In addition, associated persons preparing research reports are prohibited from owning securities in the subject companies.

Charts indicating changes in ratings can be found in recent notes and/or reports at our website or by contacting Truist Securities. Please see our disclosures page for more complete information at https://truist.bluematrix.com/sellside/Disclosures.action

Important Disclosures on Equity Research Dissemination, Ratings, Designations, and Coverage Universe

Dissemination of Research

Truist Securities, Inc. ("Truist Securities") seeks to make all reasonable efforts to provide research reports simultaneously to all eligible clients. Reports are available as published in the restricted access area of our website to all eligible clients who have requested a password. Institutional investors, corporates, and members of the Press may also receive our research via third party vendors including: Thomson Reuters, Bloomberg, FactSet, and S&P Capital IQ. Additional distribution may be done by sales personnel via email, fax, or other electronic means, or regular mail.

For access to third party vendors or our Research website: https://truistresearch.bluematrix.com/client/library.jsp

Please email the Research Department at EquityResearchDepartment@truist.com or contact your Truist Securities sales representative.

Truist Securities Rating System for Equity Securities

Truist Securities, Inc. ("Truist Securities") rates individual equities using a three-tiered system. Each stock is rated relative to the broader market (generally the S&P 500) over the next 12-18 months (unless otherwise indicated).

Buy (B) – the stock's total return is expected to outperform the S&P 500 or relevant benchmark over the next 12-18 months (unless otherwise indicated)

Hold (H) – the stock's total return is expected to perform in line with the S&P 500 or relevant benchmark over the next 12-18 months (unless otherwise indicated)

Sell (S) – the stock's total return is expected to underperform the S&P 500 or relevant benchmark over the next 12-18 months (unless otherwise indicated)

Not Rated (NR) – Truist Securities does not have an investment rating or opinion on the stock

Coverage Suspended (CS) – indicates that Truist Securities' rating and/or target price have been temporarily suspended due to applicable regulations and/or Truist Securities Management discretion. The previously published rating and target price should not be relied upon.

Truist Securities analysts have a price target on the stocks that they cover, unless otherwise indicated. The price target represents that analyst's expectation of where the stock will trade in the next 12-18 months (unless otherwise indicated). If an analyst believes that there are insufficient valuation drivers and/or investment catalysts to derive a positive or negative investment view, they may elect with the approval of Truist Securities Research Management not to assign a target price; likewise certain stocks that trade under $5 may exhibit volatility whereby assigning a price target would be unhelpful to making an investment decision. As such, with Research Management's approval, an analyst may refrain from assigning a target to a sub-$5 stock.

Legend for Rating and Price Target History Charts:

B = Buy

H = Hold

S = Sell

D = Drop Coverage

CS = Coverage Suspended

NR = Not Rated

I = Initiate Coverage

T = Transfer Coverage

Truist Securities ratings distribution (as of 04/28/2026):

Coverage Universe			Investment Banking Clients Past 12 Months		
Rating	**Count**	**Percent**	**Rating**	**Count**	**Percent**
Buy	501	66.45%	Buy	116	23.15%
Hold	251	33.29%	Hold	52	20.72%
Sell	2	0.27%	Sell	1	50.00%

Other Disclosures

Certain information contained herein has been derived from third-party sources believed to be reliable but is not guaranteed as to accuracy and does not purport to be a complete analysis of the security, company or industry involved. This report is not to be construed as an offer to sell or a solicitation of an offer to buy any security. Truist Securities, Inc. and/or its officers or employees may have positions in any securities, options, rights or warrants. The firm and/or associated persons may sell to or buy from customers on a principal basis. Investors may be prohibited in certain states from purchasing some over the-counter securities mentioned herein. Opinions expressed are subject to change without notice.

Truist Securities, Inc.’s research is primarily provided to and intended for use by Institutional Accounts as defined in FINRA Rule 4512(c). The term “Institutional Account" shall mean the account of: (1) a bank, savings and loan association, insurance company or registered investment company; (2) an investment adviser registered either with the SEC under Section 203 of the Investment Advisers Act or with a state securities commission (or any agency or office performing like functions); or (3) any other person (whether a natural person, corporation, partnership, trust or otherwise) with total assets of at least $50 million. In addition, certain affiliates of Truist Securities, Inc., including Truist Investment Services, Inc. (an SEC registered broker-dealer and a member of FINRA, SIPC) and Truist Advisory Services, Inc. (an investment adviser registered with the SEC), may make Truist Securities, Inc. research available, upon request, to certain of their clients from time to time.

Truist Securities, Inc. is a registered broker-dealer and a member of FINRA and SIPC. It is a service mark of Truist Financial Corporation. Truist Securities, Inc. is owned by Truist Financial Corporation and affiliated with Truist Investment Services, Inc. Despite this affiliation, securities recommended, offered, sold by, or held at Truist Securities, Inc. or Truist Investment Services, Inc. (i) are not insured by the Federal Deposit Insurance Corporation; (ii) are not deposits or other obligations of any insured depository institution (including Truist Bank); and (iii) are subject to investment risks, including the possible loss of the principal amount invested. Truist Bank may have a lending relationship with companies mentioned herein. Certain clients may compensate Truist Securities, Inc. for research via hard dollar payments, and Truist Securities, Inc. may be deemed to be an investment adviser to such clients as a result of such payments.

Please see our Disclosure Database to search by ticker or company name for the current required disclosures, including valuation and risks. Link: https://truist.bluematrix.com/sellside/Disclosures.action

Please visit the Truist Securities Equity Research Portal for coverage universe, current reports, and analyst roster with contact information. Link: https://truist-portal.bluematrix.com

Truist Securities, Inc., member FINRA and SIPC. Truist, Truist Bank, Truist Securities, Truist Investment Services, and Truist Advisory Services are service marks of Truist Financial Corporation.

If you no longer wish to receive this type of communication, please request removal by sending an email to EquityResearchDepartment@truist.com

© Truist Securities, Inc. 2026. All rights reserved. Reproduction or quotation in whole or part without permission is forbidden. Without Truist Securities’ prior written consent, the contents of this report may not be used, in whole or in part, in connection with the training, development, fine-tuning or operation of a machine learning or artificial intelligence system, tool or application.

ADDITIONAL INFORMATION IS AVAILABLE at our website, TruistSecurities.com, or by writing to: Truist Securities, Research Department, 740 Battery Avenue SE, 3rd Floor, Atlanta, GA 30339-3107